

# Chapter 1: Introduction



# Outline

---

Computer System Architecture

Operating Systems

- Definition

- Goals

- Operations

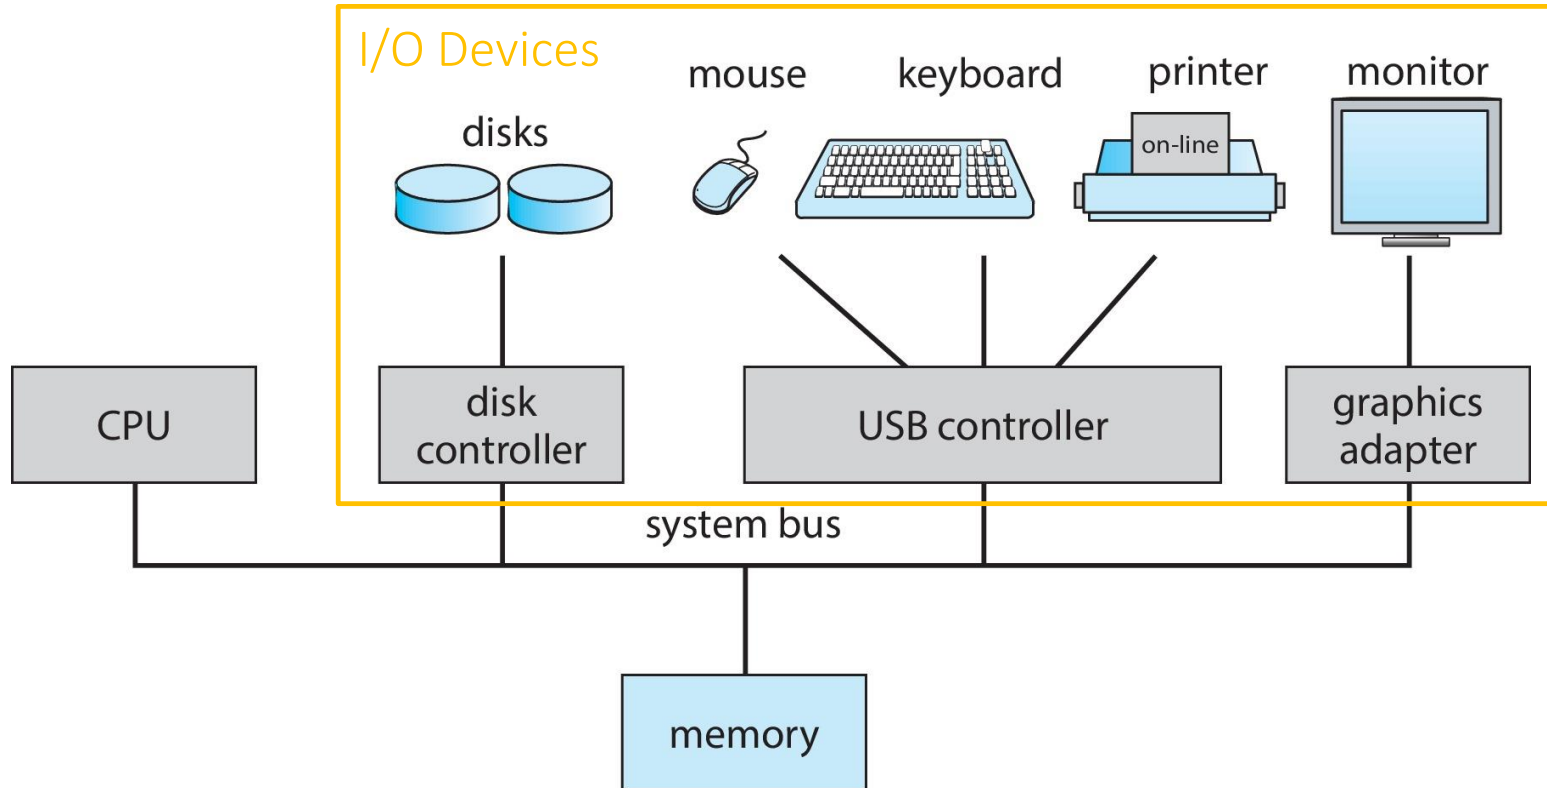
- Types

Virtualization

Free and Open Operating Systems

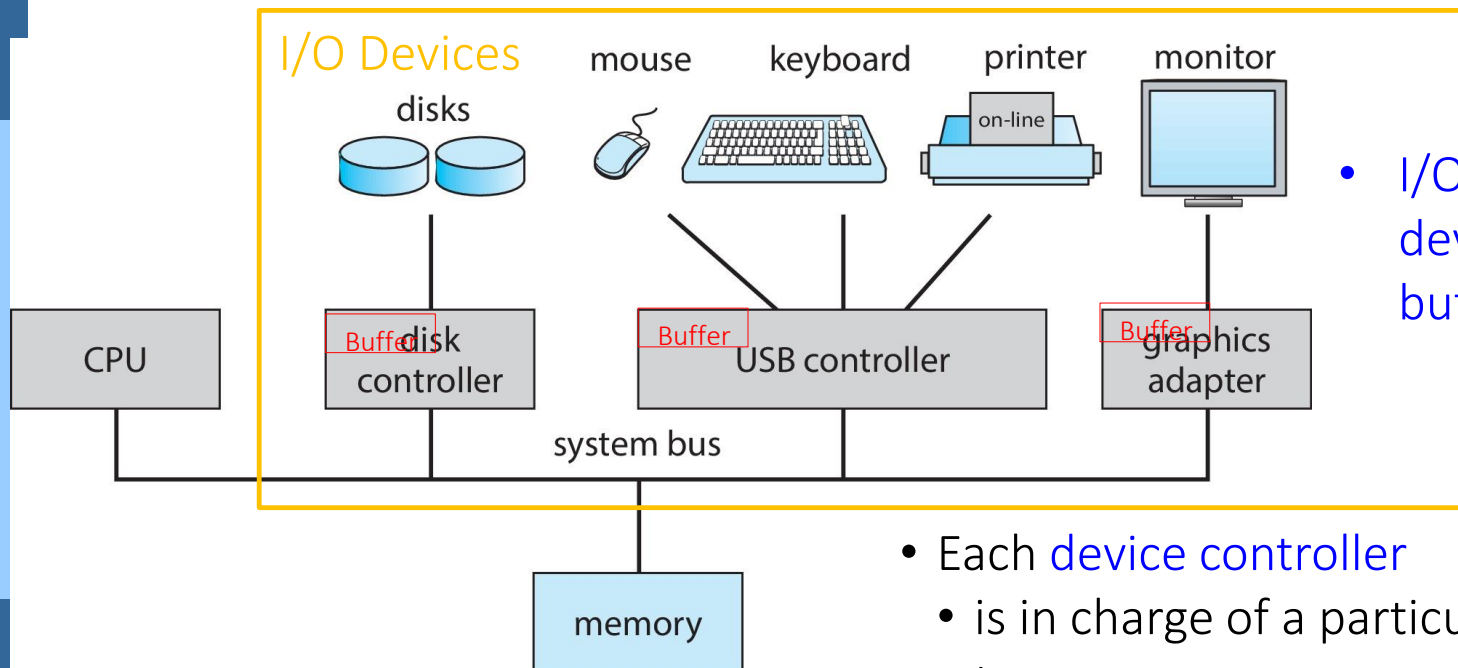
Main Data Structures Used in OS

# Computer System Architecture



# I/O Devices in Computer

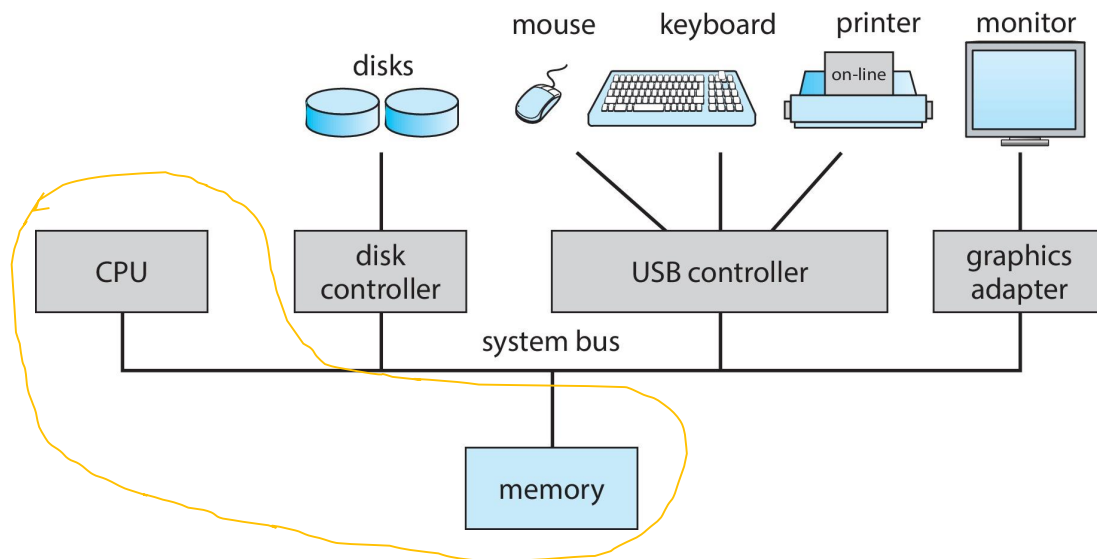
- I/O devices and the CPU can work concurrently



- I/O is between the device and local buffer of controller

- Each device controller
  - is in charge of a particular device type
  - has a local buffer (local memory)

# CPU and Memory in Computer




Single processor

or


Multiple processors

or




Dual core

or

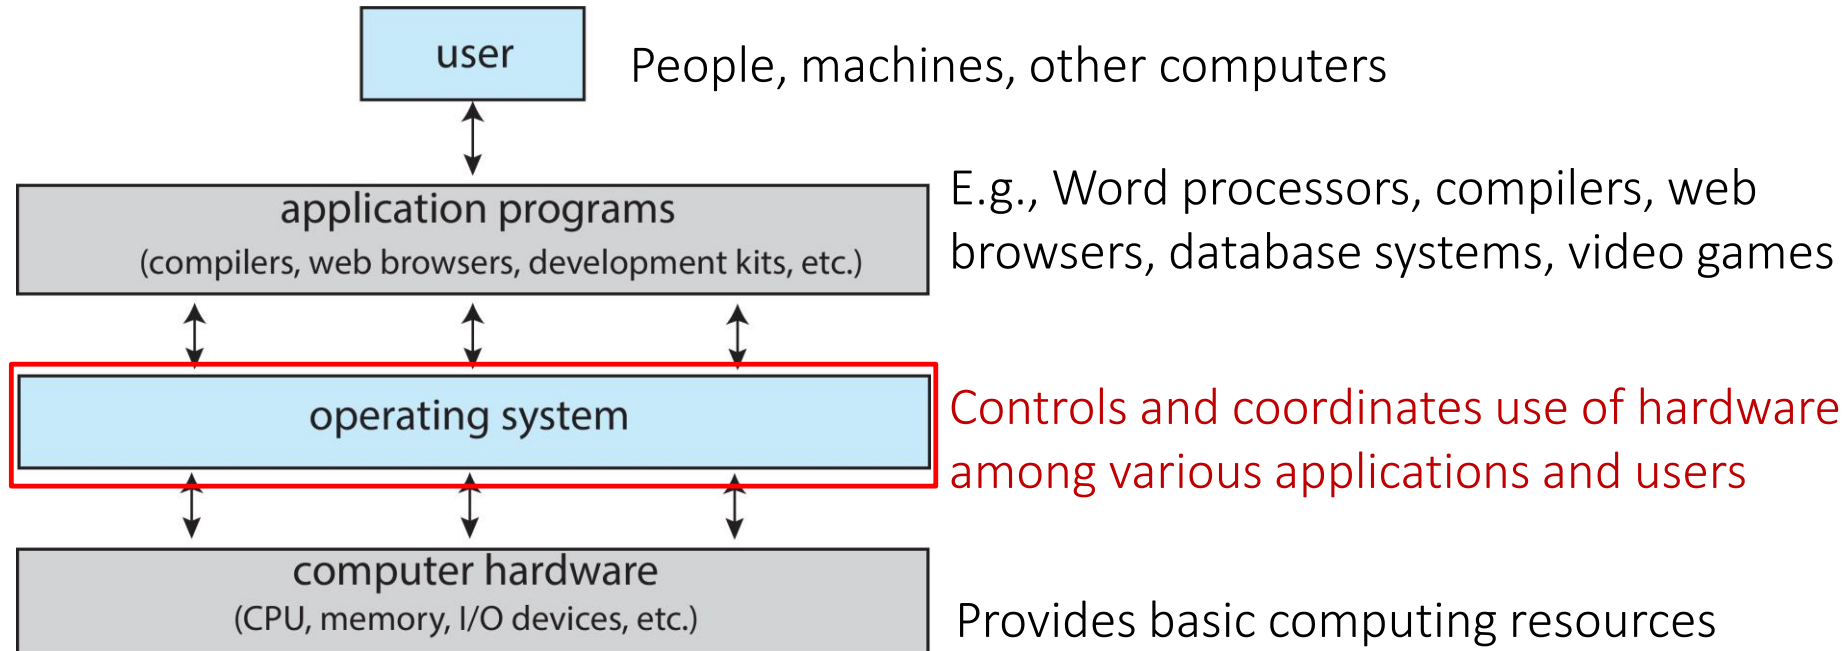


Non-Uniform Memory  
Access System (NUMA)

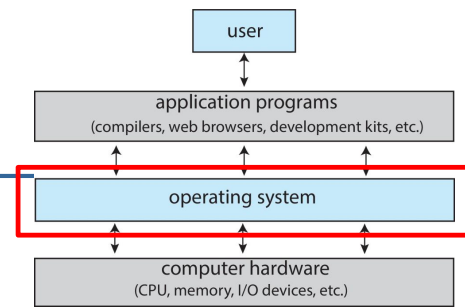


# Operating System: Intermediary between a User and the Computer Hardware

---



# Operating Systems



An operating system

is a **program** that acts as an intermediary between a **user** of a computer and the **computer hardware**

provides **interface** to a user, **runs a program** for a user

General goals of operation systems

Execute user programs and make solving user problems easier

Make the computer system convenient to use

Use the computer hardware in an efficient manner



Picture source: from bing.com

# Operating Systems Goals

---

Depends on computer/device types

**Personal** computer

- ▶ Feature: owned by individuals
- ▶ OS goal: convenience, ease of use and good performance

**Shared** computer (e.g. mainframe or minicomputer)

- ▶ Feature: shared by multiple users
- ▶ OS goal: keep all happy, efficient use of Hardware and managing user programs



# What Operating Systems Do

---

Depends on computer/device types (cont.)

**Mobile** devices (e.g. smartphones and tablets)

- ▶ Feature: resource poor
- ▶ **OS goal**: optimized for **usability and battery life**, provide user interfaces such as **touch screens, voice recognition**

**Embedded computers**

- ▶ Feature: little or no user interface
- ▶ **OS goal**: run program

**User view**: ease of use

**System view**: resource allocation

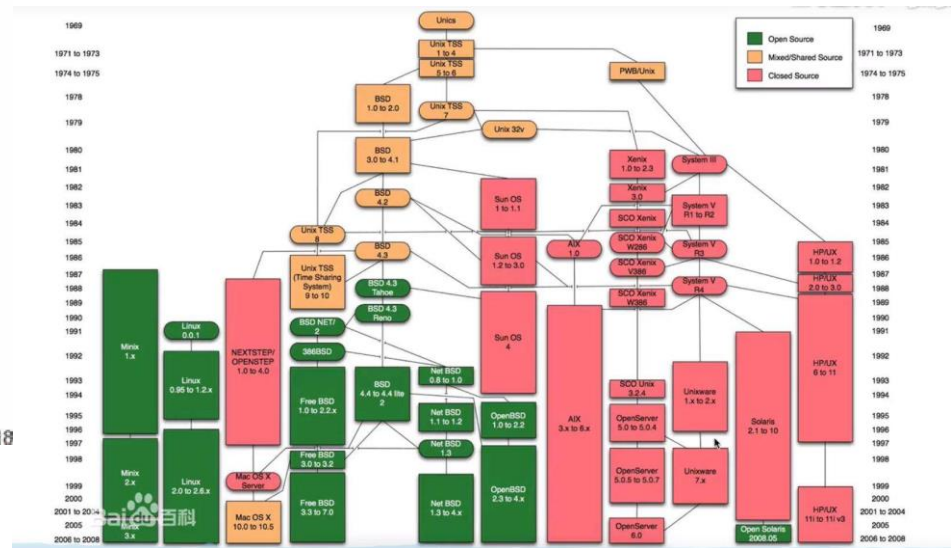
# Operating System Examples

Term OS exists in many devices.

There is a large number of OS designs



Jaceu2018



Linux

Picture resource:

[https://cn.bing.com/images/search?view=detailV2&ccid=m%2fuR5GOU&id=298444024B0761F4FD9628A1F3C02D0CF00DD3B8&thid=OIP.m\\_uR5GOUh7Mbc2NbXowUAHafJ&mediaurl=https%3a%2f%2fimg00.deviantart.net%2fece9%2f%2f2018%2f021%2fd%2fc%2fwindows\\_family\\_photo\\_by\\_jaceu-dc0r7j6.png&exph=450&expw=600&q=windows+family+pictures&simid=608037459554172544&FORM=IRPRST&ck=95ED26317FF82815C67EE21D7145788F&selectedIndex=0&idpp=overlayview&ajaxhist=0&ajaxserp=0https://baike.baidu.com/item/MacOS/8654551](https://cn.bing.com/images/search?view=detailV2&ccid=m%2fuR5GOU&id=298444024B0761F4FD9628A1F3C02D0CF00DD3B8&thid=OIP.m_uR5GOUh7Mbc2NbXowUAHafJ&mediaurl=https%3a%2f%2fimg00.deviantart.net%2fece9%2f%2f2018%2f021%2fd%2fc%2fwindows_family_photo_by_jaceu-dc0r7j6.png&exph=450&expw=600&q=windows+family+pictures&simid=608037459554172544&FORM=IRPRST&ck=95ED26317FF82815C67EE21D7145788F&selectedIndex=0&idpp=overlayview&ajaxhist=0&ajaxserp=0https://baike.baidu.com/item/MacOS/8654551)

# Operating System Definition

---

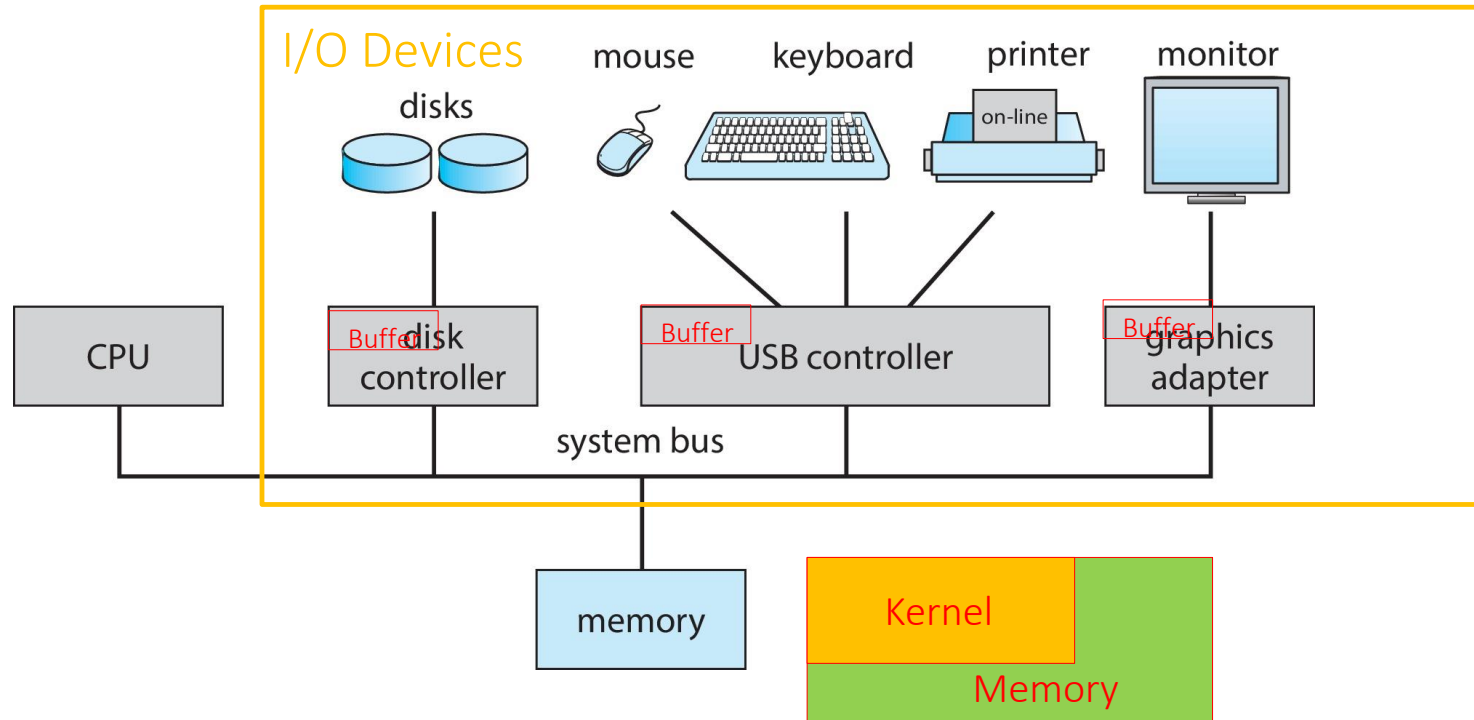
No universally accepted definition

A definition: “Everything a vendor ships when you order an operating system”

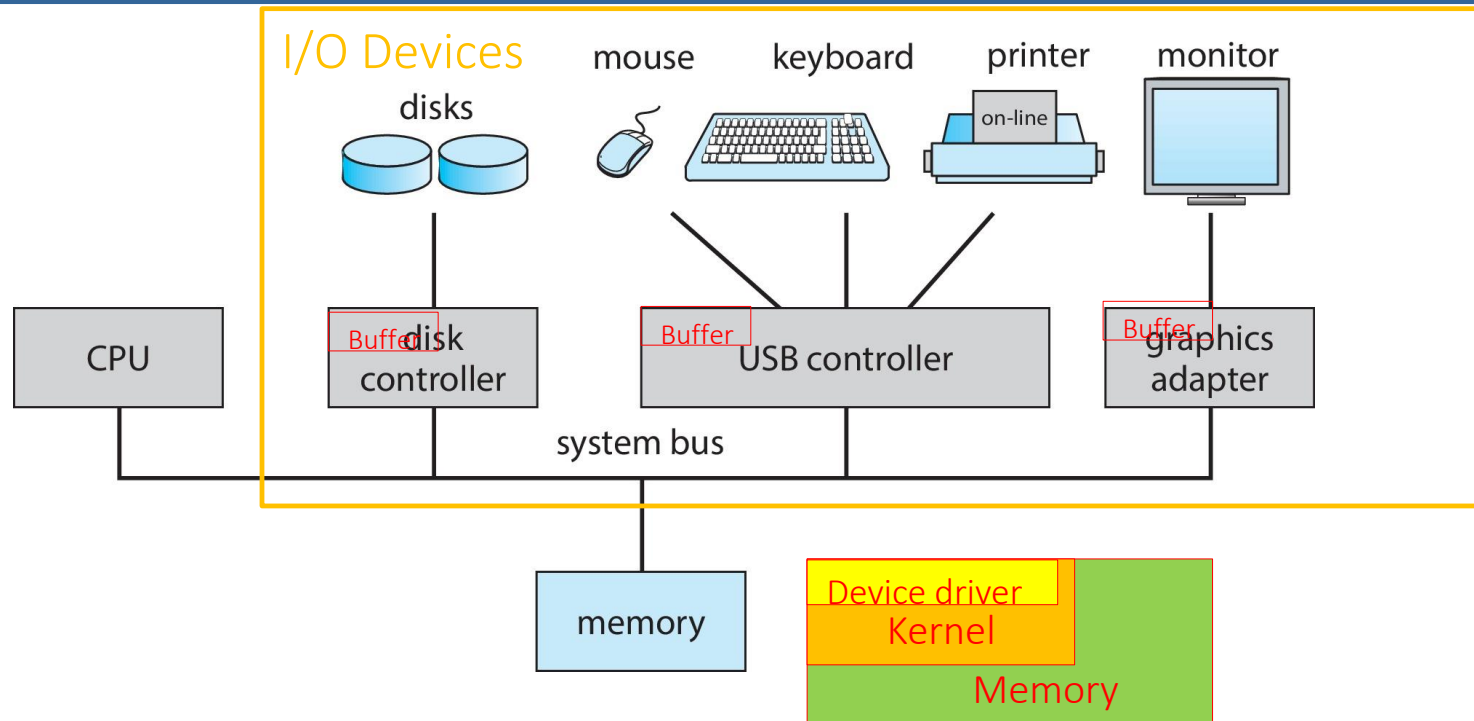
“Everything” include

- ▶ **Kernel** - part of a operating system, resides **in memory at all times** on the computer,
- ▶ others
  - **system programs** - all programs associated with the operating system, but not in **kernel**
  - **application program** - all programs not associated with the operating system
  - **middleware** – a set of **software frameworks** that provide additional services to application developers such as databases, multimedia, graphics

# Kernel in Memory



# Operating System Is Interrupt-Driven



- **Device driver** (software) inside the **kernel** knows how to talk and manage the **device**
- **Device driver** provides uniform interface between **controller** and **kernel**
- Each **device controller** informs **CPU** that it has finished its operation by causing an **interrupt**

# Interrupts

---

An **interrupt** is a signal emitted by **hardware** or **software** when a process or an event needs immediate attention

**Software**: A **trap** or **exception** is a software-generated interrupt caused either by **an error** or **a user request** (**system call**)

**Hardware**: Device

**ISR** (**Interrupt Service Routines**) **inside the kernel** determine what action should be taken **for each type of interrupt**

An operating system is **interrupt driven**

the code of the kernel is in memory all the time but the code of the kernel is only executed when there is an **interrupt**, **on demand**!

**An operating system is interrupt driven**

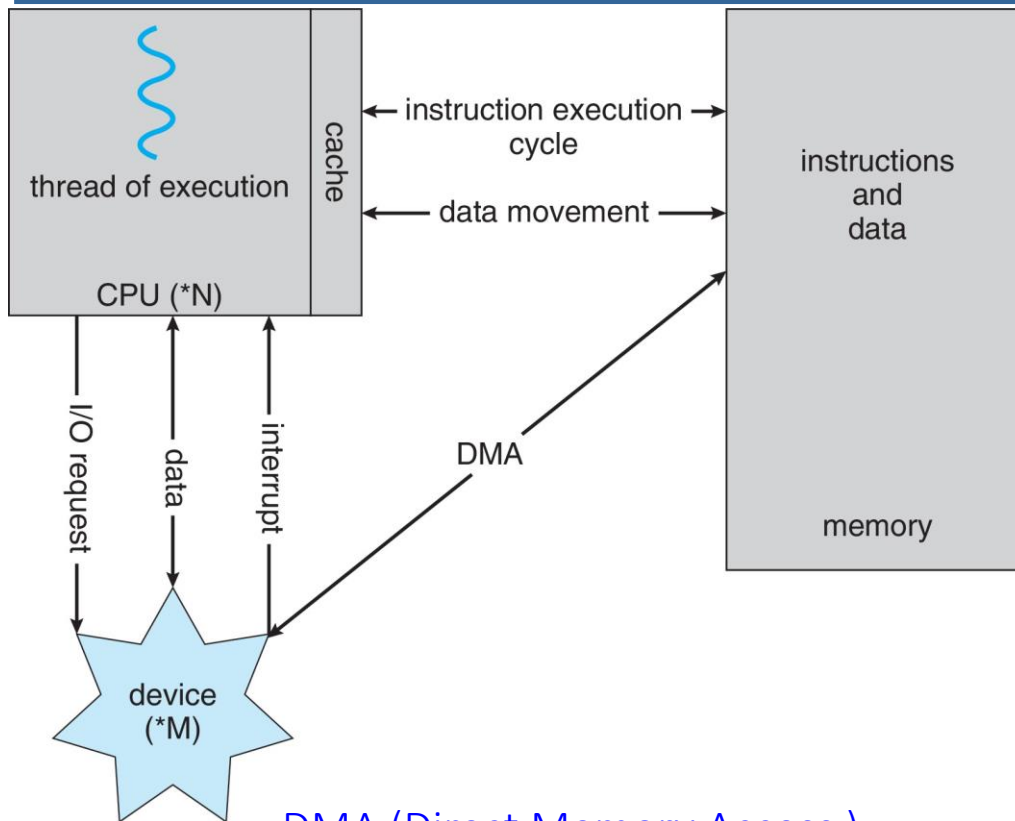
# Interrupt Handling Steps

---

When there is an **interrupt**, the operating system

1. **Preserves the state** of the CPU by storing registers and the program counter(PC) for the software that was just interrupted (so that the same software can be restarted later).
2. **Determine type of interrupt**
  - ▶ check **Interrupt vector** to get the address of corresponding **ISR** for this interrupt (**used on all modern computers**)
3. **Runs ISR in kernel** to handle **interrupt**

# How a Modern Computer Works: DMA




*A von Neumann architecture*

## DMA (Direct Memory Access )

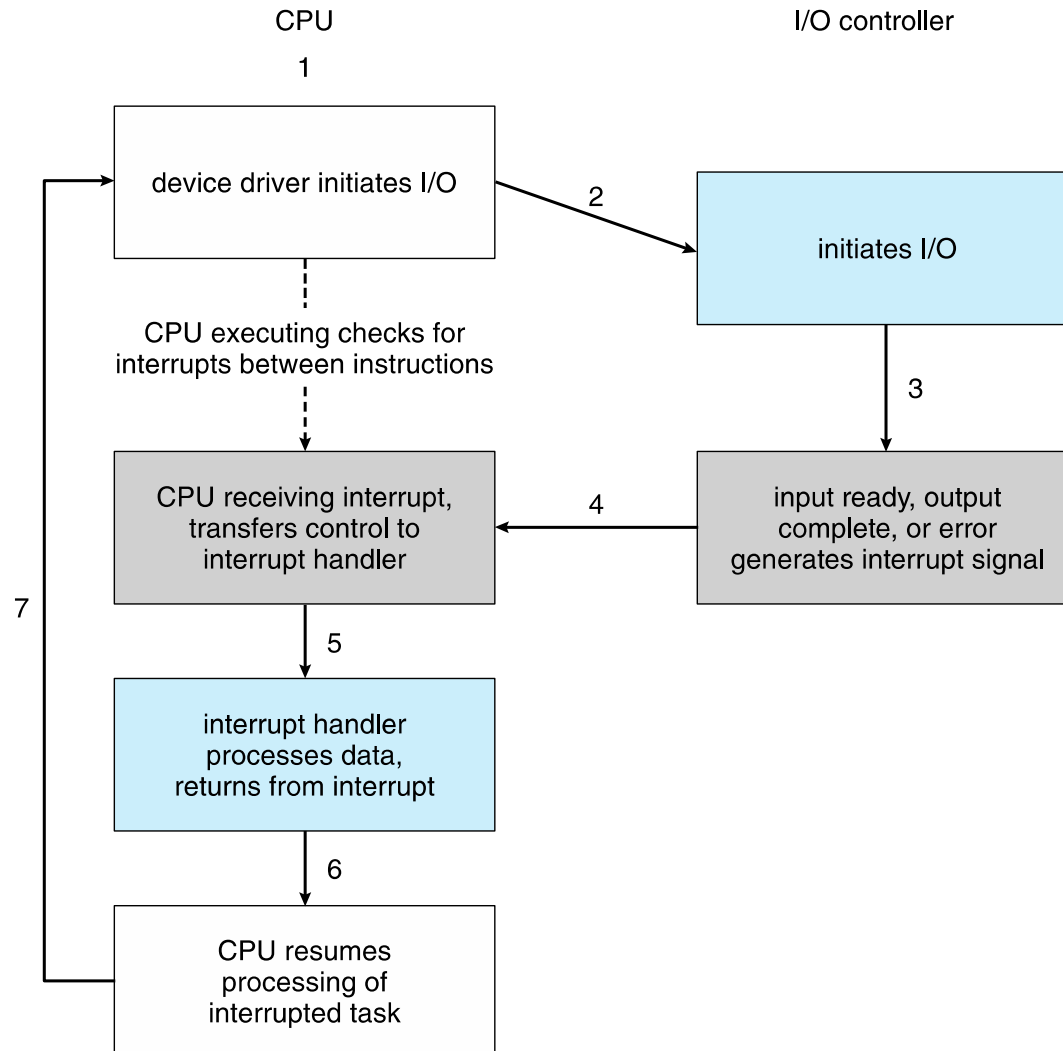
- **Device controller** transfers **blocks** of data from buffer directly to main memory without CPU intervention
  - In the meantime **the CPU can work on something else.**
- Only **one interrupt** is generated **per block**
- High transfer speed: not byte by byte, but block by block



# Interrupt-Driven I/O Cycle: Timeline View



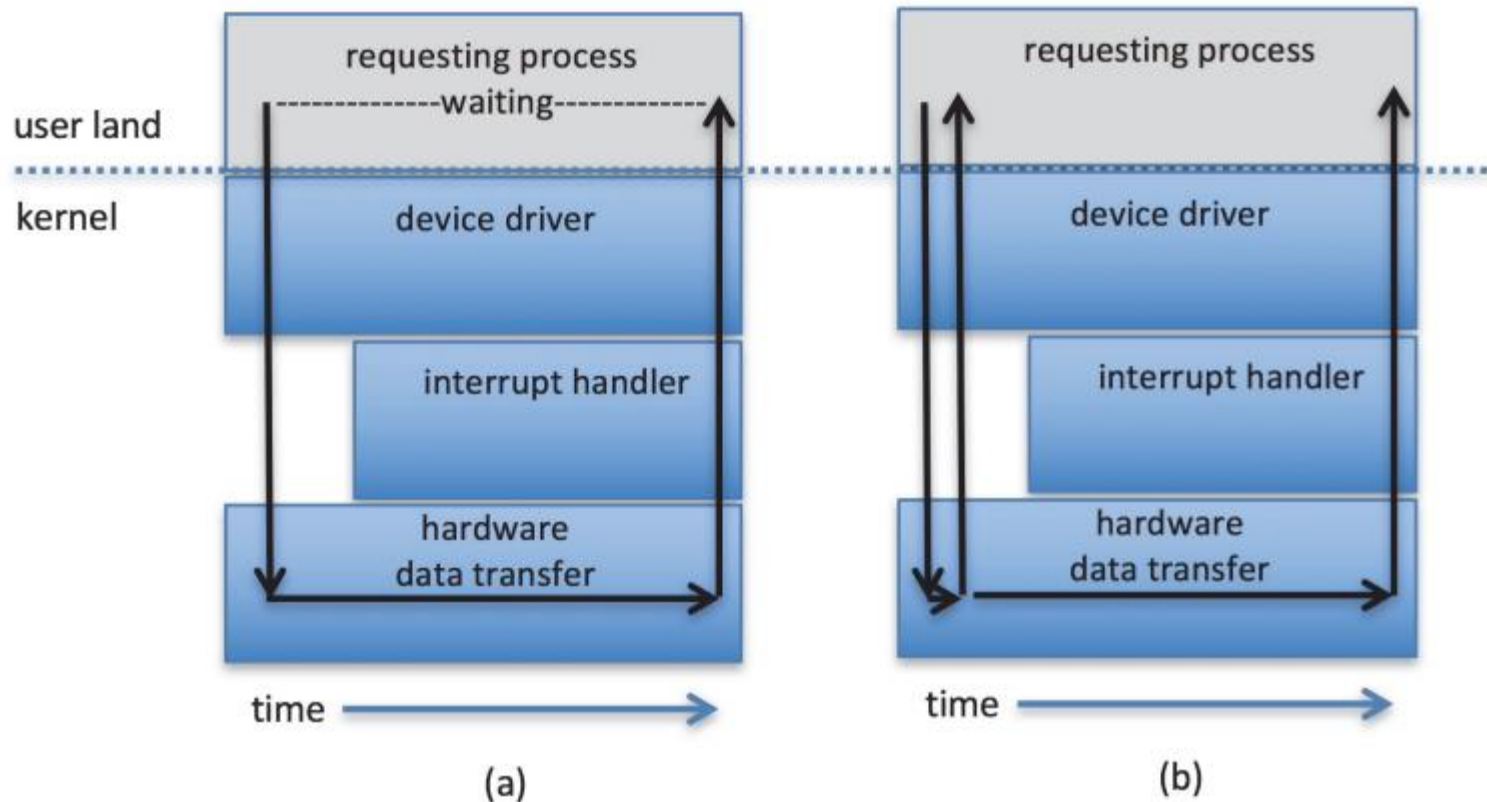
# Interrupt-Driven I/O Cycle – Workflow View



# Interrupt-Driven I/O Cycle: Two Methods

After I/O starts, control returns to user program **only upon I/O completion**

After I/O starts, control returns to user program **without waiting for I/O completion**



**Figure 12.9** Two I/O methods: (a) synchronous and (b) asynchronous.

# Operating System Activities

---

Operating system does a lot of management work

- Caching
- Process management
- Memory management
- File management
- Secondary storage management
- I/O
- Protection and security

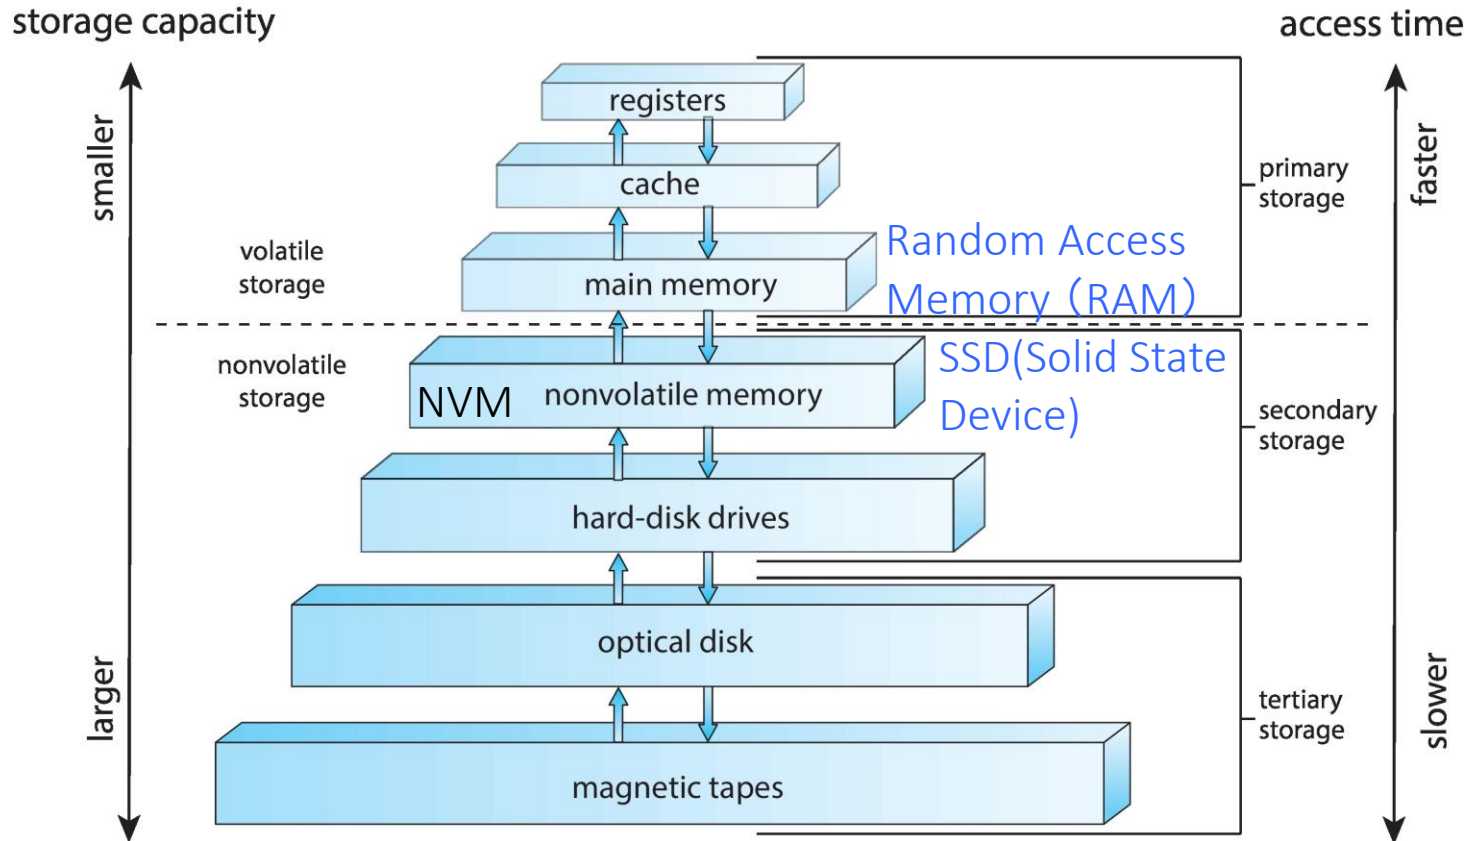
# Operating System Activities

---

Operating system does a lot of management work

- Caching
- Process management
- Memory management
- File management
- Secondary storage management
- I/O
- Protection and security

# Storage-Device Hierarchy



Comparison:

- Speed
- Cost
- Volatility

# Storage Size Units

The basic unit of computer storage is the **bit**. A bit can contain one of two values, 0 and 1. All other storage in a computer is based on collections of bits. Given enough bits, it is amazing how many things a computer can represent: numbers, letters, images, movies, sounds, documents, and programs, to name a few. A **byte** is 8 bits, and on most computers it is the smallest convenient chunk of storage. For example, most computers don't have an instruction to move a bit but do have one to move a byte. A less common term is **word**, which is a given computer architecture's native unit of data. A word is made up of one or more bytes. For example, a computer that has 64-bit registers and 64-bit memory addressing typically has 64-bit (8-byte) words. A computer executes many operations in its native word size rather than a byte at a time.

Computer storage, along with most computer throughput, is generally measured and manipulated in bytes and collections of bytes. A **kilobyte**, or KB, is 1,024 bytes; a **megabyte**, or MB, is  $1,024^2$  bytes; a **gigabyte**, or GB, is  $1,024^3$  bytes; a **terabyte**, or TB, is  $1,024^4$  bytes; and a **petabyte**, or PB, is  $1,024^5$  bytes. Computer manufacturers often round off these numbers and say that a megabyte is 1 million bytes and a gigabyte is 1 billion bytes.

Networking measurements are an exception to this general rule; they are given in bits (because networks move data a bit at a time).

# Storage Size Units (In Chinese)

中文单位	中文简称	英文单位	英文简称	进率 (Byte=1)
位元	比特	bit	b	0.125
字节	字节	Byte	B	1
千字节	千字节	KiloByte	KB	$2^{10}$
兆字节	兆	MegaByte	MB	$2^{20}$
吉字节	吉	GigaByte	GB	$2^{30}$
太字节	太	TeraByte	TB	$2^{40}$
拍字节	拍	PetaByte	PB	$2^{50}$
艾字节	艾	ExaByte	EB	$2^{60}$
泽字节	泽	ZettaByte	ZB	$2^{70}$
尧字节	尧	YottaByte	YB	$2^{80}$
珀字节	珀	BrontoByte	BB	$2^{90}$



# Characteristics of Various Types of Storage


Level	1	2	3	4	5
Name	registers	cache	main memory	solid-state disk	magnetic disk
Typical size	< 1 KB	< 16MB	< 64GB	< 1 TB	< 10 TB
Implementation technology	custom memory with multiple ports CMOS	on-chip or off-chip CMOS SRAM	CMOS SRAM	flash memory	magnetic disk
Access time (ns)	0.25-0.5	0.5-25	80-250	25,000-50,000	5,000,000
Bandwidth (MB/sec)	20,000-100,000	5,000-10,000	1,000-5,000	500	20-150
Managed by	compiler	hardware	operating system	operating system	operating system
Backed by	cache	main memory	disk	disk	disk or tape

Data moves between levels of storage hierarchy explicitly or implicitly

# Caching

**Caching:** Copy information from slower into faster storage system

**Purpose:** Speed up data access



**Speed:** slow —————→ fast

**Size:** big —————→ small

**Price:** cheap —————→ expensive

- Data access procedure: faster storage (cache) checked first
  - Data exist, information is used directly from the cache (fast)
  - Data not exist, data are copied from slower device to the faster.


Cache design consideration

size

content replacement policy

# Cache Coherency

---



Migration of data “A” from Disk to Register

**Cache coherency:** The coordination of the contents of caches such that an update to a value stored in one cache is immediately reflected in all other caches that hold that value

# Operating system activities

---

Operating system does a lot of management work

- Caching
- Process management
- Memory management
- File management
- Secondary storage management
- I/O
- Protection and security

# Process

---

## Program vs process

A process is a program in execution

- ▶ Program is a passive entity, on storage
- ▶ Process is an active entity, in memory

## Process

creation, execution needs resources

- ▶ CPU, memory, I/O, files
- ▶ Initialization data

Termination requires reclaim of any reusable resources

## Process can be

single-threaded (线程), or  
multi-threaded

Each process has a program counter to  
specify location of next instruction to execute

# Multiprogramming and Multitasking

Single process cannot always keep CPU and I/O devices busy



**Multiprogramming** : A subset of total jobs are kept in memory at the same time

**Job scheduling**: choose which jobs to load into memory.

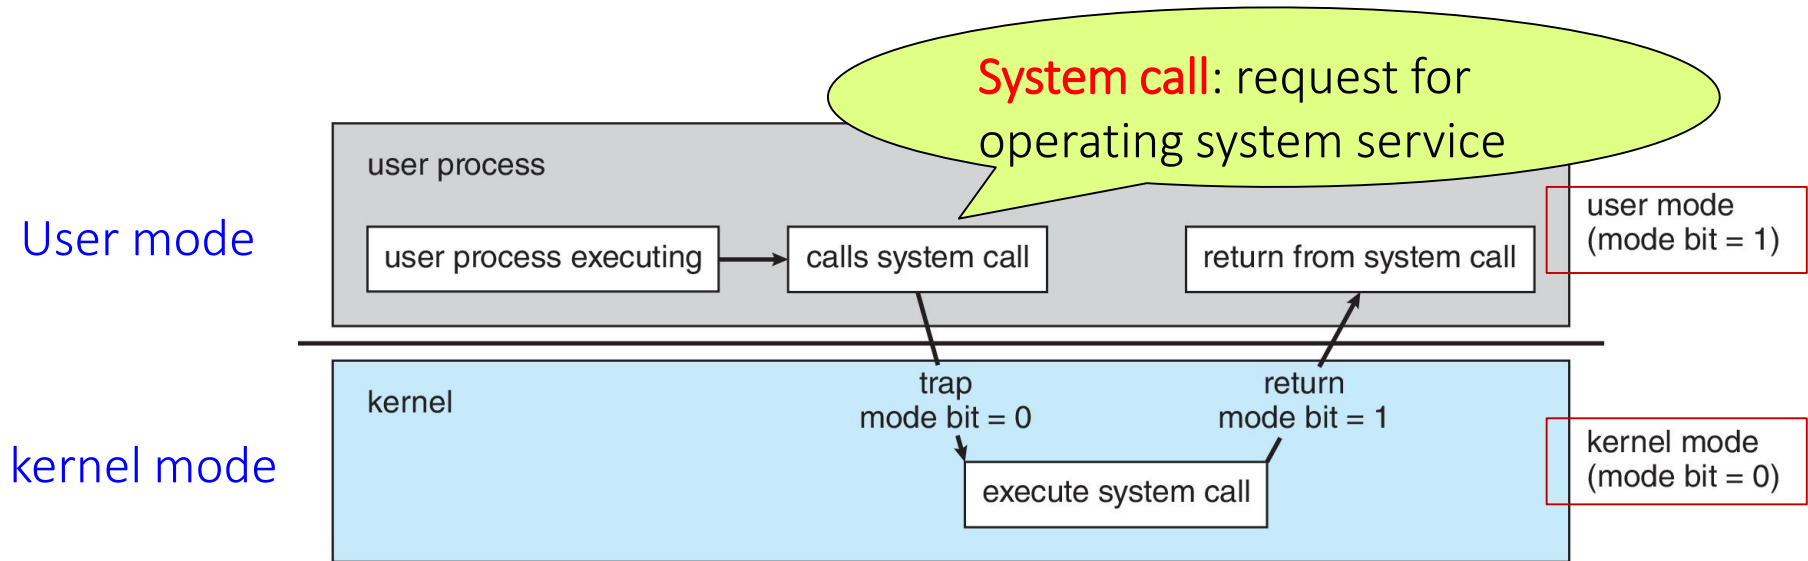
When job has to wait (for I/O for example), OS switches to another job

**Multitasking**: The CPU switches jobs so frequently to increase interactions with users

**CPU scheduling**: choose which job will run next If several jobs are ready to run at the same time



# Process Execution: Dual-mode



**Dual-mode** operation allows OS to **protect** itself and other system components

# Process Management Activities

---

Process management activities include:

- Creating and deleting both user and system processes

- Suspending and resuming processes

- Process synchronization

- Process communication

- Deadlock handling



# Operating system activities

---

Operating system does a lot of management work

- Caching
- Process management
- Memory management
- File management
- Secondary storage management
- I/O
- Protection and security

# Memory Management

---

To execute a program,

all (or part) of the **instructions** of the program must be in **memory**

All (or part) of the **data** needed by the program must be in memory

Memory management determines **what is in memory and when**

**Memory management** activities include

1. Keeping track of which parts of memory are currently being used and by whom
2. Deciding which processes and data (or part of them) to move into and out of memory
3. Allocating and de-allocating memory space as needed

# Operating system activities

---

Operating system does a lot of management work

- Caching
- Process management
- Memory management
- File management
- Secondary storage management
- I/O
- Protection and security

# File System Management

---

OS provides uniform, logical view of information storage

Abstracts physical properties to logical storage unit - file

Files are usually organized into directories

Access control determines who can access what

File system activities include

1. Creating and deleting files and directories
2. Mapping files onto secondary storage

# Secondary Storage Management

---

Disks usually are used to store data that do not fit in main memory or data that must be kept for a “long” period of time

Proper management is of central importance

Entire speed of computer operation hinges on disk subsystem and its algorithms

Secondary storage management activities

1. Mounting and unmounting
2. Free-space management
3. Storage allocation
4. Disk scheduling
5. Partitioning
6. Protection

# I/O Subsystem

---

One purpose of OS is to hide peculiarities of hardware devices from the user

I/O subsystem responsible for

Memory management of I/O including **buffering** (storing data temporarily while it is being transferred), **caching** (storing parts of data in faster storage for performance), **spooling** (the overlapping of output of one job with input of other jobs)

General device-driver interface

Drivers for specific hardware devices

Spooling: Simultaneous Peripheral Operations Online

# Protection and Security

---

## Protection

any mechanism for **controlling access** of processes or users to resources defined by the OS

## Security

**defense of the system** against internal and external attacks

- ▶ Huge range, including denial-of-service, worms, viruses, identity theft, theft of service

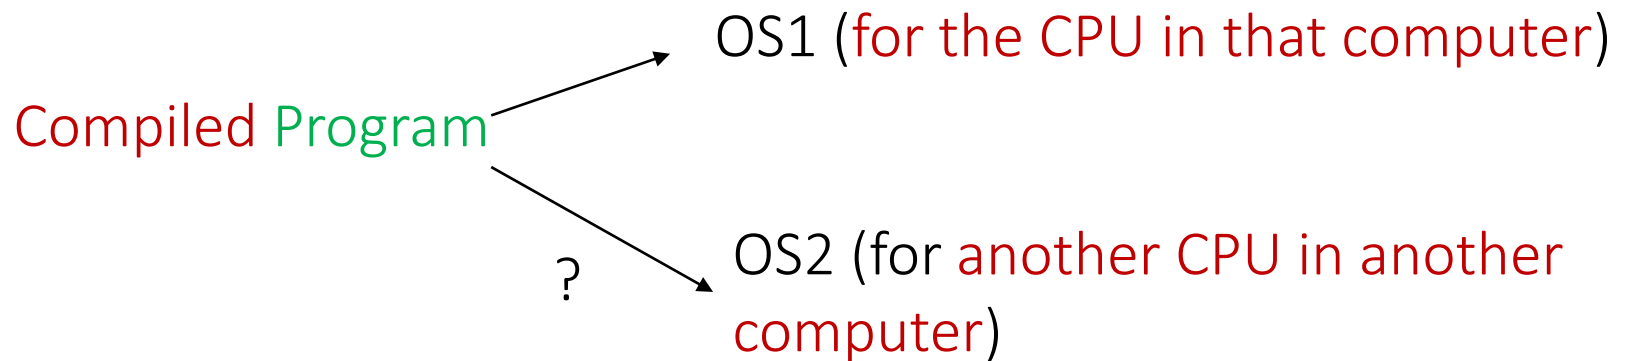
Systems generally first distinguish among users, to determine who can do what

User identities (**user IDs**, security IDs) include name and associated number, one per user

Group identifier (**group ID**) allows set of users to be defined

# Can A **Compiled** Program Be Portable

---

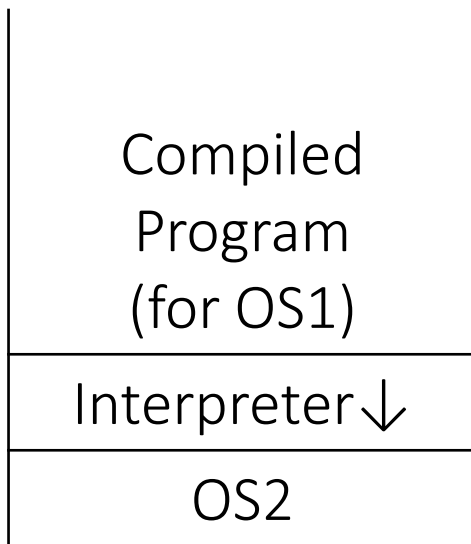


Solutions?

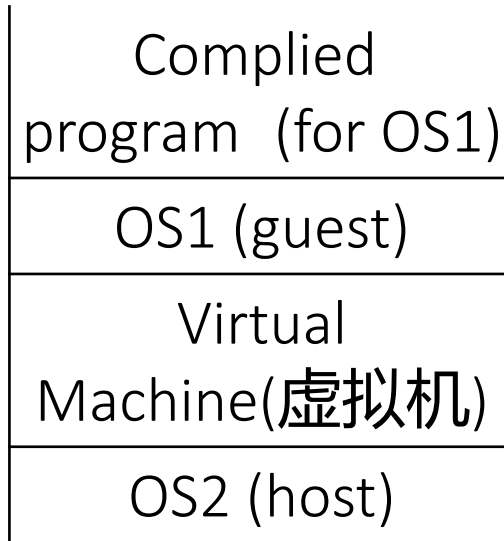


# Virtualization

Solutions: two methods.



Emulation  
(Old method)



Natively compiled OS

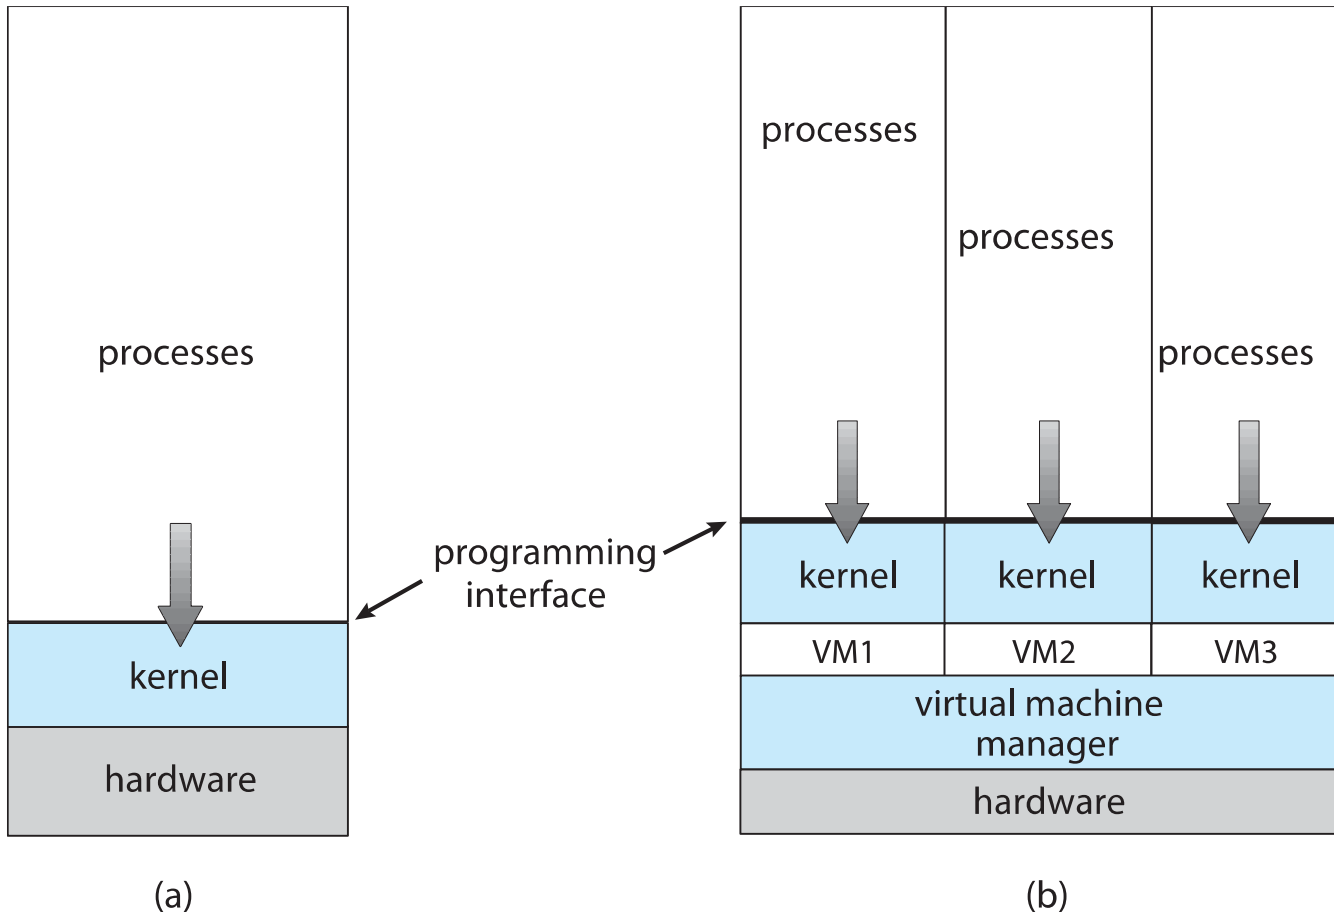
Natively compiled OS

Virtualization

Virtual Machine: VirtualBox, VMware

**Virtualization** is a technology that allows us to abstract the hardware of a single computer into several different execution environments, thereby creating the illusion that each separate environment is running on its own private computer.

# Virtualization



## VMMs (Virtual Machine Managers)

- Run natively, **no longer run on host operating systems but rather *are* the host operating systems**, providing services and resource management to virtual machine processes.
- Examples: VMware ESX and Citrix XenServer

# Virtualization Examples

---

## Use cases

- Apple laptop running Mac OS X host, Windows as a guest
- Developing apps for multiple OSes without having multiple systems
- Quality assurance testing applications without having multiple systems
- Executing and managing compute environments within data centers

After class:

Virtualization vs. Dual OS ???

# Free and Open-Source Operating Systems

---

Practical differences between **free software** (social movement) and **open source** (development technology)

all existing released free software source code would qualify as open source.

Nearly all open source software is free software, but there are exceptions.

More differences see web

<https://www.geeksforgeeks.org/difference-between-free-software-and-open-source-software/>

# Free and Open-Source Operating Systems

---

Free or open source operating systems are made available in source-code format rather than just binary closed-source and proprietary

Examples

- ▶ GNU/Linux and BSD UNIX (including core of Mac OS X), and many more
- ▶ VM like VMware Player (Free on Windows), Virtualbox (open source and free on many platforms - <http://www.virtualbox.com>)

# The Study of Operating Systems

---

There has never been a more interesting time to study operating systems, and it has never been easier. The open-source movement has overtaken operating systems, causing many of them to be made available in both source and binary (executable) format. The list of operating systems available in both formats includes Linux, BSD UNIX, Solaris, and part of macOS. The availability of source code allows us to study operating systems from the inside out. Questions that we could once answer only by looking at documentation or the behavior of an operating system we can now answer by examining the code itself.

Operating systems that are no longer commercially viable have been open-sourced as well, enabling us to study how systems operated in a time of fewer CPU, memory, and storage resources. An extensive but incomplete list of open-source operating-system projects is available from [https://curlie.org/Computers/Software/Operating\\_Systems/Open\\_Source/](https://curlie.org/Computers/Software/Operating_Systems/Open_Source/)

In addition, the rise of virtualization as a mainstream (and frequently free) computer function makes it possible to run many operating systems on top of one core system. For example, VMware (<http://www.vmware.com>) provides a free “player” for Windows on which hundreds of free “virtual appliances” can run. Virtualbox (<http://www.virtualbox.com>) provides a free, open-source virtual machine manager on many operating systems. Using such tools, students can try out hundreds of operating systems without dedicated hardware.

The advent of open-source operating systems has also made it easier to make the move from student to operating-system developer. With some knowledge, some effort, and an Internet connection, a student can even create a new operating-system distribution. Just a few years ago, it was difficult or impossible to get access to source code. Now, such access is limited only by how much interest, time, and disk space a student has.

# Free and Open-Source Operating Systems

---

Operating system used in this semester

Ubuntu

Three ways to access Ubuntu in this semester

Install it in your own computer following guidelines in iSpace (Virtual machine: Virtualbox)

▶ Virtualbox download:

<https://www.virtualbox.org/wiki/Downloads>


Access computers in CST lab rooms (Virtual machine: VMware)

Access through the network (ssh to server with user account and password).


# Kernel Data Structures

Many similar to standard programming data structures


## *Singly linked list*



## *Doubly linked list*




## *Circular linked list*






# Kernel Data Structures

Binary search tree  
left  $\leq$  right



# Kernel Data Structures

Hash function can create a hash map



**Bitmap** – string of  $n$  binary digits representing the status of  $n$  items

**Linux data structures** defined in *include* files

`<linux/list.h>`, `<linux/kfifo.h>`,  
`<linux/rbtree.h>`

# End of Chapter 1

