



Chapter 6 & 7: Process Synchronization



Process Synchronization

- Background
- The Critical-Section Problem
- Peterson's Solution
- Synchronization Hardware
- Mutex Locks
- Semaphores
- Classic Problems of Synchronization
- Monitors



Background

- Processes can execute concurrently
 - May be interrupted **at any time**, partially completing execution
 - Concurrent access to shared data may result in **data inconsistency**
- An example
 - The variable **counter** in an implementation of producer-consumer

```
item next_produced;
while (true) {
    /* produce an item in next_produced */
    while (counter == BUFFER_SIZE) //buffer is full
        /* do nothing, such as create data */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
    counter++;
```

```
item next_consumed;
while (true) {
    while (counter == 0) //buffer is empty
        /* do nothing */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    counter--; //total # of items in the buffer
    /* consume the item in next_consumed */
}
```



Race Condition

- counter++ could be implemented as

```
register1 = counter
register1 = register1 + 1
counter = register1
```
- counter-- could be implemented as

```
register2 = counter
register2 = register2 - 1
counter = register2
```
- Consider this execution interleaving with “counter = 5” initially:

S0: producer execute	register1 = counter	{register1 = 5}
S1: producer execute	register1 = register1 + 1	{register1 = 6}
S2: consumer execute	register2 = counter	{register2 = 5}
S3: consumer execute	register2 = register2 - 1	{register2 = 4}
S4: producer execute	counter = register1	{counter = 6}
S5: consumer execute	counter = register2	{counter = 4}

Race condition:

- several processes access and manipulate the **same data** concurrently
- **outcome depends on** which **order** each access takes place.



Critical Section Problem

- Consider system of n processes $\{p_0, p_1, \dots, p_{n-1}\}$
- Several processes may be **changing common variables, updating table, writing file**, etc.
 - **critical section**
 - ▶ The segment of code in a process that modifies these shared variables, tables, files
- When one process is in **critical section**, other process should not enter their critical sections for these shared data .

do {

entry section

critical section

exit section

remainder section

} while (true);



Critical Section

- General structure of process P_i

do {

entry section

Ask for permission

critical section

Protect this section

exit section

Do sth in order to allow other
processes to enter critical section

remainder section

} while (true);



Critical Section Problem

- Critical section problem is to design a protocol (协议) that the processes can use to cooperate
 - Process with critical section should follow the following steps
 1. execute entry section to ask for permission
 2. then execute critical section,
 3. execute exit section to allow other process to enter critical section
 4. then execute remainder section



Solution to Critical-Section Problem

- A solution to the critical-section problem must satisfy **three requirements**:
 1. Mutual Exclusion
 - ▶ If process P_i is executing in its **critical section**, then no other processes can be executing in their critical sections
 2. Progress
 - ▶ If **no process is executing in its critical section** and if other processes want to enter critical section, one of them must be selected. They **cannot be postponed indefinitely**
 3. Bounded Waiting
 - ▶ If a process has made a request to enter its critical section, then, before that request is granted, there is a bound to the times that others can enter critical section



Solution to Critical-Section Problem

- Assumptions in the solution
 - ▶ each process executes at a nonzero speed
 - ▶ no concerning **relative speed** of the n processes
 - ▶ the **load** and **store** machine-language instructions are **atomic** (that is, cannot be interrupted)



Example: Solution to Critical-Section Problem

Example: two students want to eat from the same plate of food using one spoon.

1. **Mutual Exclusion** - If one student is eating using the spoon, then the other student must wait for the spoon to become available: the two students cannot both use the spoon at the same time! Otherwise the two students will end up fighting and breaking the plate of food (race condition).
2. **Progress** - If one student is not hungry and the other student wants to get the spoon to eat, then the other student must be able to take the spoon and eat: the student who is not hungry is not allowed to prevent the other student from taking the spoon and eating. Otherwise the other student will wait for ever and die from hunger.
3. **Bounded Waiting** - If one student is eating and the other student is waiting for the spoon, the first student is not allowed to put down the spoon on the table and immediately take it again to eat again: the student who puts down the spoon must allow the other student to take the spoon so that the other student can eat too. In other words, the two students must be nice with each other and not be selfish. Otherwise the other student will wait for ever and die from hunger.



Critical Section: Solution 1

turn's value is initialized to be either *i* or *j*

P_i	P_j
<pre>do { while (turn == j); //entry section critical section //turn is i turn = j; //exit section remainder section } while (true);</pre>	<pre>do { while (turn == i); //entry section critical section // turn is j turn = i; //exit section remainder section } while (true);</pre>

Mutual Exclusion: Yes

Progress: No

Bounded waiting: Yes



Detailed Explanation on Example 1

1. **Mutual Exclusion: Yes** - If process P_i is inside the critical section then P_j cannot enter the critical section because of the **while** loop in the entry section that forces P_j to wait for P_i to exit the critical section.
 1. And vice-versa: if P_j is inside the critical section then P_i must wait in the entry section until P_j exits the critical section.
 2. What happens if both processes try to enter the critical section at the exact same time? The shared **turn** variable can only store **i** or **j**, but not both, so the test of the **while** loop will only be false for one of the two processes, not for both, and only one process will be able to enter the critical section, the other process will have to wait in the entry section.
 1. Which process enters the critical section first is then determined by the initial value of the **turn** variable.



Detailed Explanation on Example 1 (Continue)

2. **Progress: No** – If the **turn** variable contains the value **j** (for example) and P_i wants to enter the critical section and P_j does not want to enter the critical section (assuming there is no **do{ }while** loop around P_j 's code and P_j is busy doing something else) then P_i will wait for ever, even though P_j is not inside the critical section.
 1. P_i can only enter the critical section after P_j has been in the critical section, and P_j can only enter the critical section after P_i has been in the critical section.
$$P_i \rightarrow P_j \rightarrow P_i \rightarrow P_j \rightarrow P_i \rightarrow P_j \rightarrow P_i \rightarrow \dots$$
 2. If P_j does not want to enter the critical section again then P_i is going to get stuck and wait for ever in the entry section for P_j to exit the critical section, which P_j is never going to do.
$$P_i \rightarrow P_j \rightarrow P_i \rightarrow P_i \text{ again and gets stuck for ever.}$$



Detailed Explanation on Example 1 (Continue)

3. **Bounded Waiting: Yes** - There is a bound of 1 time on the number of times that P_j (for example) is allowed to enter the critical section after P_i has made a request to enter its critical section and before that request is granted, because when P_j exits the critical section it changes the **turn** variable to be **i**, which then immediately allows P_i to enter the critical section.
 1. Even if we assume that P_j is very fast and P_i is very slow (for example), it is not possible for P_j to beat P_i to the race and exit the critical section and immediately re-enter it again before P_i is allowed to take its turn inside the critical section.
 2. In other words, it is not possible for P_j to prevent P_i from going into the critical section by going in and out of the critical section as fast as possible.
 3. So P_i is guaranteed to be able to enter the critical section at some point in the future (right after P_j is finished) and will not have to wait for ever, even if P_j wants to re-enter the critical section all the time.



Peterson's Solution

- It is a classic software-based solution to the critical-section problem
 - Good algorithmic description of solving the problem
- Solution for two processes by using two variables:
 - ▶ `int turn;` // indicates whose turn it is to enter the critical section.
 - ▶ `Boolean flag[2]` // indicate if a process is ready to enter the critical section.
 - `flag[i] = true` implies that process P_i is ready!
 - It is initialized to `FALSE`.



Algorithm for Process P_i & P_j

do {

P_i

```
flag[i] = true; //ready  
turn = j; //allow  $P_j$  to enter
```

Ask for entry permission

```
while (flag[j] && turn == j);
```

critical section

```
flag[i] = false; //exit
```

remainder section

} while (true);

Mutual Exclusion: **Yes**

Progress: **Yes**

Bounded waiting: **Yes**

do {

P_j

```
flag[j] = true; //ready  
turn = i; //allow  $P_i$  to enter
```

Ask for entry permission

```
while (flag[i] && turn == i);
```

critical section

```
flag[j] = false; //exit
```

remainder section

} while (true);



Synchronization Hardware

- Software-based solutions (such as Peterson's) are not guaranteed to work on modern computer architecture
- Many systems provide **hardware support** for implementing the critical section code.
 - **Uniprocessors** – could disable interrupts
 - ▶ Currently running code would execute without preemption
 - **too inefficient**
 - Modern machines provide special atomic hardware instructions
 - ▶ **Atomic** = non-interruptible, the atomic hardware instruction will do the following work
 1. Test memory *word* and set value
 2. Swap contents of two memory *words*
 - ▶ E.g., test_and_set instruction in the next page



Atomic test_and_set Instruction

Definition:

```
boolean test_and_set (boolean *target)
{
    boolean rv = *target;
    *target = TRUE;
    return rv;
}
```

1. The C code here is just to explain how the hardware instruction works.
2. This instruction is executed **atomically** by CPU as a single hardware instruction.
3. In practice, *target* is a pointer to the lock itself, shared by all the processes that want to acquire the lock.
 - ▶ If *target* is **FALSE**, the return value of *rv* is **FALSE**, means lock is **FALSE** (available), *target*'s new value is **TRUE**
 - ▶ If *target* is true, the return value of *rv* is **TRUE**, means lock is **TRUE** (locked)



Solution to Critical-section Problem Using Locks

```
do {  
    acquire lock           entry section  
    critical section  
    release lock          exit section  
    remainder section  
} while (TRUE);
```

- Use the idea of locking
 - Protecting critical regions via locks
- A process that wants to enter the critical section must first get the lock.
- If the lock is already acquired by another process, the process will wait until the lock becomes available.



Solution using test_and_set()

- ❑ Shared Boolean variable **lock**, initialized to **FALSE**
- ❑ Solution using test_and_set:

```
do {  
    while (test_and_set(&lock))  
        ; /* do nothing */  
    /* critical section */  
    lock = FALSE;  
    /* remainder section */  
} while (true);
```

```
boolean test_and_set  
(boolean *target)  
{  
    boolean rv = *target;  
    *target = TRUE;  
    return rv;  
}
```

```
do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
} while (TRUE);
```

- When lock = true, keep **while** looping.
- When lock = FALSE, process can enter the critical section
- And set lock = TRUE, **block** other processes to enter.
- After finish the critical section, reset lock = FALSE, to allow other processes to enter the critical section.

Mutual Exclusion: **Yes**

Progress: **Yes**

Bounded waiting: **No**



Atomic compare_and_swap Instruction

Definition:

```
int compare_and_swap(int *value, int expected, int new_value)
{
    int temp = *value;
    if (*value == expected)
        *value = new_value;
    return temp;
}
```

- The C code here is just to explain how the hardware instruction works.
- The instruction is executed atomically by CPU as a single hardware instruction.
 1. In practice, *value* is a pointer to the lock itself, shared by all the processes that want to acquire the lock.
 2. Set **value* (the lock) to the value of the passed parameter *new_value* but only if **value == expected*. That is, the swap takes place only under this condition.
- Returns as result the original value of the lock.
- Similar to *test_and_set* but with an integer *lock* and an extra condition.



Solution using compare_and_swap

- Shared integer *lock* initialized to 0 (false);
- Solution using *compare_and_swap*:

```
do {  
    ↩ while (compare_and_swap(&lock, 0, 1) != 0)  
        ; /* do nothing */  
    /* critical section */  
    lock = 0;  
    /* remainder section */  
} while (true);
```

Mutual Exclusion: Yes

Progress: Yes


Bounded waiting: No

test_and_set(&lock)



Solution using compare_and_swap and test_and_set

- Mutual exclusion: Yes
- Progress: Yes
- Bounded Waiting: No
- Busy-waiting: Yes (in entry section, use while statement)
- Therefore the solution presented in the previous slide is not good enough either.



Bounded-waiting Mutual Exclusion with test_and_set


```
do {  
    waiting[i] = true;  
    key = true;  
    while (waiting[i] && key)  
        key = test_and_set(&lock);  
    waiting[i] = false;  
  
    /* critical section */  
  
    j = (i + 1) % n;  
    while ((j != i) && !waiting[j])  
        j = (j + 1) % n;  
    if (j == i)  
        lock = false;  
    else  
        waiting[j] = false;  
  
    /* remainder section */  
} while (true);
```

P_i

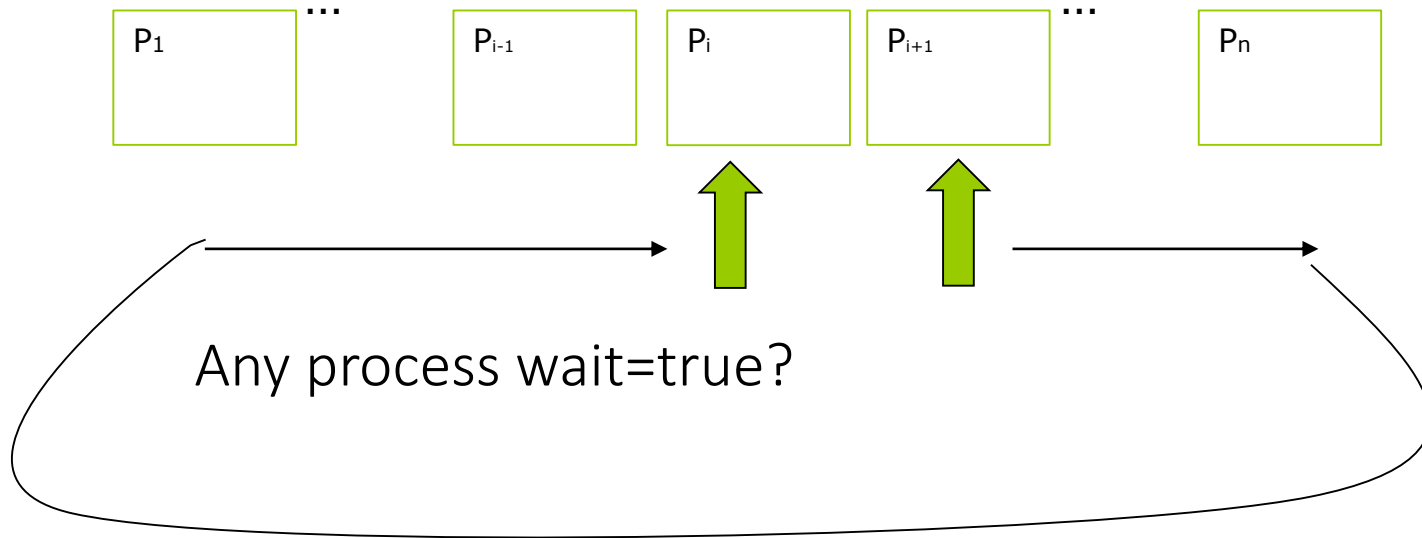



```
do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
} while (TRUE);
```

Common data structures:
boolean waiting[n];
boolean lock;
These data structures are
initialized to **false**.



Bounded-waiting Mutual Exclusion with test_and_set (Continue)





Bounded-waiting Mutual Exclusion with test_and_set (Continue)

P_i

```
do {
    waiting[i] = true;
    key = true;
    while (waiting[i] && key)
        key = test_and_set(&lock);
    waiting[i] = false;

    /* critical section
*/

    j = (i + 1) % n;
    while ((j != i) && !waiting[j])
        j = (j + 1) % n;
    if (j == i)
        lock = false;
    else
        waiting[j] = false;

    /* remainder section */
} while (true);
```

P_k

```
do {
    waiting[k] = true;
    key = true;
    while (waiting[k] && key)
        key = test_and_set(&lock);
    waiting[k] = false;

    /* critical section */

    j = (k + 1) % n;
    while ((j != k) && !waiting[j])
        j = (j + 1) % n;
    if (j == k)
        lock = false;
    else
        waiting[j] = false;

    /* remainder section */
} while (true);
```

Mutual exclusion: Yes

Progress: Yes

Bounded-waiting: Yes




Explanations on Three Reequipments for Slide 24

Mutual exclusion: **Yes**

```
waiting[i] = true;
key = true;
while (waiting[i] && key)
    key = test_and_set(&lock);
waiting[i] = false;
```

1. Process P_i can enter its critical section only if either **waiting[i] == false** or **key == false**
 - If waiting[i] is false, that means, another process finished critical section and give the turn to process P_i
 - If key is false, that means, lock was false before run test_and_set.
2. If another process (e.g., P_k) in waiting to enter, it cannot enter. Because waiting[k] is true, and lock is true.



Bounded-waiting Mutual Exclusion with test_and_set

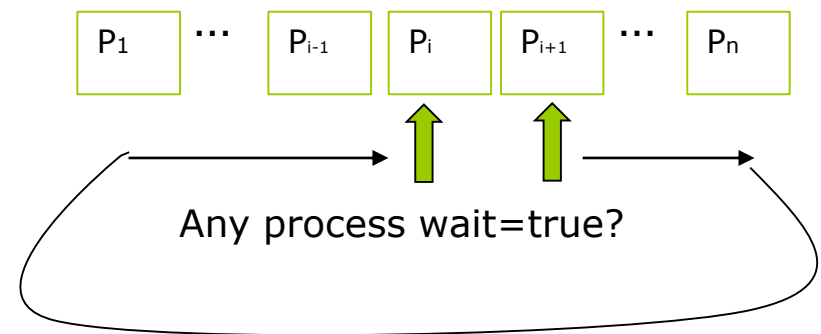
Progress: **Yes**

If P_i wants to enter the critical section, while other process do not want to enter, lock is false. So the test_and_set will return false.

Bounded waiting: **Yes**

If process P_i leaves the critical section, and another process (e.g. P_k) is waiting, the exit section in P_i will set waiting[k] to the false


```
j = (i + 1) % n;  
while ((j != i) && !waiting[j])  
    j = (j + 1) % n;  
if (j == i)  
    lock = false;  
else  
    waiting[j] = false;
```





Mutex Locks

- Previous hardware-based solutions are complicated and generally **inaccessible to application programmers**
- OS designers build high-level software tools to solve **critical section problem**
 - Simplest one of these tools is **mutex lock (互斥锁)**
 - Use **mutex lock** to protect a critical section by first `acquire()` a lock then `release()` the lock
 - Assumption: Calls to `acquire()` and `release()` must be **atomic**
 - Usually implemented via hardware atomic instructions
- This solution still requires **busy waiting**
 - This lock therefore is called a **spinlock**



acquire() and release()

```
acquire() {  
    while (!available)  
        ; /* busy wait */  
    available = false;  
}
```

```
release() {  
    available = true;  
}
```

```
do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
} while (true);
```

- This solution requires busy waiting
 - This lock therefore called a spinlock(自旋锁)



POSIX Mutex Locks

```
#include <pthread.h>
pthread_mutex_t mutex;
pthread_mutex_init(&mutex, NULL);
```

```
pthread_mutex_lock(&mutex);
```

```
pthread_mutex_unlock(&mutex);
```

```
int pthread_mutex_init(pthread_mutex_t *mutex, const pthread_mutexattr_t *mutexattr);
```

The `pthread_mutex_init()` function initializes the specified mutex. If *attr* is non-NULL, the attributes specified are used to initialize the mutex. If the attribute object is modified later, the mutex's attributes are not affected. If *attr* is NULL, the mutex is initialized with default attributes, as specified for `pthread_mutex_init()`.

A mutex can be statically initialized by assigning `PTHREAD_MUTEX_INITIALIZER` in its definition, as follows:

```
pthread_mutex_t def_mutex = PTHREAD_MUTEX_INITIALIZER;
```

A mutex must be initialized (either by calling `pthread_mutex_init()`, or statically) before it may be used in any other mutex functions.

```
do {
    acquire lock
    critical section
    release lock
    remainder
} while (true);
```



Semaphore(信号灯)

- Semaphore is a synchronization tool
 - more sophisticated than mutex locks
- Semaphore S : an integer variable
 - S can only be accessed via two atomic operations
 - ▶ `wait()` and `signal()`
 - Originally called `P()` and `V()`

```
wait(S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}
```

acquire lock

```
signal(S) {  
    S++;  
}
```

release lock



Semaphore(信号灯)

```
do {  
    wait(S);  
    /* critical section */  
    signal(S)  
    /* remainder section */  
} while (true);
```

P_i

```
do {  
    wait(S);  
    /* critical section */  
    signal(S)  
    /* remainder section */  
} while (true);
```

P_j

```
wait(S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}
```

```
signal(S) {  
    S++;  
}
```

release lock



sem_wait and sem_post

```
#include <semaphore.h>
sem_t *sem_mutex;
```

Normally used **among different processes**

```
int sem_wait(sem_t *sem);
int sem_post(sem_t *sem);
```

```
sem_wait(sem_mutex); → acquire lock
                        critical section
sem_post(sem_mutex); → release lock
                        remainder section
                        } while (true);
```



POSIX Named Semaphore

```
#include <semaphore.h>
sem_t *sem_mutex;
```

Normally used among different processes

```
sem_t *sem_open(const char *name, int oflag, mode_t mode, unsigned int value);
```

e.g., `sem_mutex = sem_open("SEM", O_CREATE, 0666, 1);`

multiple processes can easily use a common semaphore as a synchronization mechanism by the **semaphore's name**.

```
sem_wait(sem_mutex);
```

→ **acquire lock**

critical section

```
sem_post(sem_mutex);
```

→ **release lock**

remainder section

```
} while (true);
```



POSIX Unnamed Semaphore

```
#include <semaphore.h>
sem_t sem_mutex;
```

Normally used among different threads within a process

```
int sem_init(sem_t *sem, int pshared, unsigned int value);
    e.g., sem_init(&sem_mutex, 0, 1);
```

Three parameters in `sem_init`:

1. A pointer to the semaphore
2. A flag indicating the level of sharing (shared between threads in a process or between processes)
3. The semaphore's initial value

```
sem_wait(&sem_mutex);
```

```
do {
```

```
    acquire lock
```

```
        critical section
```

```
sem_post(&sem_mutex);
```

```
    release lock
```

```
        remainder section
```

```
    } while (true);
```



Semaphore Usage

- Counting semaphore – integer value can range over an unrestricted domain
- Binary semaphore – integer value can range only between 0 and 1
 - Same as a mutex lock
- Semaphore can be used to solve various synchronization problems

An example: Consider two processes P_1 and P_2 that require T_1 to happen before T_2 . Solution:

```
semaphore synch;  
synch = 0
```

```
P1 :  
    T1 ;  
    signal (synch) ;
```

```
P2 :  
    wait (synch) ;  
    T2 ;
```



Semaphore Implementation

- Advantage
 - implementation code is short
- Disadvantage:
 - Must guarantee that no two processes can execute the **wait()** and **signal()** on the same semaphore at the same time
 - ▶ S--, S++ in wait() and signal()
 - ▶ Thus, the implementation becomes the critical section problem where the **wait** and **signal** code are placed in the critical section
 - ▶ The solution now have **busy waiting** in critical section implementation
 - Both mutex lock and semaphore suffer from **busy waiting**.



Semaphore Implementation with no Busy waiting

- With each semaphore there is an associated waiting queue
 - Each semaphore has:
 - ▶ an integer value
 - ▶ A queue of processes
- Two operations:
 - **block** – place the process invoking the operation on the appropriate waiting queue
 - **wakeup** – remove one of processes in the waiting queue and place it in the ready queue

```
typedef struct{  
    int value;  
    struct process *list;  
} semaphore;
```



Implementation with no Busy waiting (Cont.)

- Again, C code is just for explanation:

Busy waiting

```
wait(S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}
```

```
signal(S) {  
    S++;  
}
```

S: initial value 1

No busy waiting

```
wait(semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        //add this process to S->list;  
        block(); //put in a waiting queue  
    }  
}  
  
signal(semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        //remove a process P from S->list;  
        wakeup(P); //put in a ready queue  
    }  
}
```

S->value: initial value 1



Classical Problems of Synchronization

- Classical problems used to test newly-proposed synchronization schemes
 - Bounded-Buffer Problem
 - Readers-Writers Problem
 - Dining-Philosophers Problem



Bounded-Buffer Problem

- Producer-Consumer problem
 - n buffers, each holds one item
 - Semaphore **mutex** initialized to the value 1
 - Semaphore **full** initialized to the value 0
 - Semaphore **empty** initialized to the value n



Bounded Buffer Problem (Cont.)

■ The structure of the producer process

```
int n;  
semaphore mutex = 1; //mutual exclusion for buffer pool access  
semaphore empty = n; //number of empty entries in the buffer pool  
semaphore full = 0; //number of items available in the buffer pool  
do {
```

```
    ...  
    /* produce an item in next_produced */
```

```
    ...  
    wait(empty);  
    wait(mutex);
```

```
    ...  
    /* code to add next_produced to the  
    buffer */
```

```
    ...  
    signal(mutex);  
    signal(full);
```

```
} while (true);
```

```
wait(S) {  
    while (S <= 0)  
        ; // busy  
wait  
    S--;  
}  
  
signal(S) {  
    S++;  
}
```

```
item next_produced;  
while (true) {  
    /* produce an item in  
    next_produced */  
    while (counter == BUFFER_SIZE)  
        ;/* do nothing*/  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;
```



Bounded Buffer Problem (Cont.)

■ The structure of the consumer process

```
do{
    wait(full);
    wait(mutex);

    ...

    /* code to remove an item from buffer to
    next_consumed */

    ...

    signal(mutex);
    signal(empty);

    ...

    /* consume the item in next_consumed */

    ...
} while (true);
```

```
item next_consumed;
while (true) {
    while (counter == 0)
        ; /* do nothing */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    counter--; //total # of item in the buffer
    /* consume the item in next_consumed */
}
```

```
wait(S) {
    while (S <= 0)
        ; // busy wait
    S--;
}

signal(S) {
    S++;
}
```



Pthreads Synchronization

■ Semaphore functions:

```
#include <semaphore.h>
```

```
sem_t sem;
```

```
/* Create the semaphore and initialize it to 1 */
```

```
sem_init(&sem, 0, 1);
```

```
/* acquire the semaphore */
```

```
sem_wait(&sem);
```

```
/* critical section */
```

```
/* release the semaphore */
```

```
sem_post(&sem);
```



Bounded Buffer Problem (Cont.)

Sample code demonstration in class
(Bounded Buffer using semaphores)

`posix_bb.c`



Bounded Buffer Problem (Cont.)

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#include <semaphore.h>
#include <deque>
#include <unistd.h>
#include <signal.h>
```

```
#define TRUE 1
```

```
/* functions for threads to call */
void *producer(void *param); //for producer
void *consumer(void *param); //for consumer
```

```
//using three semaphores
```

```
sem_t empty;
sem_t full;
sem_t mutex;
```

```
int shared_item = 0;
```

```
std::deque<int> myboundedbuffer;
pid_t pid;
```

posix_bb.c

Compile command:
g++ -o posix_bb posix_bb.c -lpthread



Bounded Buffer Problem (Cont.)

```
int main(int argc, char *argv[]) {                                posix_bb.c
    int i, scope;
    pthread_t producerID, consumerID; /* the thread identifier */
    pthread_attr_t attr; /* set of attributes for the thread */
    int n = 10;
    pid = getpid();

    //create semaphores, and initialize
    sem_init(&empty, 0, n); //n
    sem_init(&full, 0, 0); //0
    sem_init(&mutex, 0, 1); //1

    /* get the default attributes */
    pthread_attr_init(&attr);

    /* create the threads */
    pthread_create(&producerID, &attr, producer, &producerID);
    pthread_create(&consumerID, &attr, consumer, &consumerID);

    /*Now join on each thread*/
    pthread_join(producerID, NULL);
    pthread_join(consumerID, NULL);
    return 0;
}
```




Bounded Buffer Problem (Cont.)

/*The thread will begin control in this function.*/

posix_bb.c

```
void *producer(void *param) {  
    int id = *(int*)param;  
    printf("producer Thread ID = %d\n", id);  
  
    do {  
        sem_wait(&empty);  
        sem_wait(&mutex);  
  
        shared_item++;  
        if (shared_item > 20)  
            break;  
  
        printf("Producer create an item %d\n", shared_item);  
        myboundedbuffer.push_back(shared_item);  
  
        sem_post(&mutex);  
        sem_post(&full);  
    } while (TRUE);  
  
    kill(pid, SIGINT); //send a signal to kernel to terminate the process  
    pthread_exit(0);  
}
```

This program will finish
when the producer has
created 20 items



Bounded Buffer Problem (Cont.)

posix_bb.c

```
void *consumer(void *param) {  
  
    int id = *(int*)param;  
    printf("consumer Thread ID = %d\n", id);  
  
    do {  
        sem_wait(&full);  
        sem_wait(&mutex);  
  
        printf("\tConsumer process an item %d\n", myboundedbuffer.front());  
        myboundedbuffer.pop_front();  
  
        sem_post(&mutex);  
        sem_post(&empty);  
  
    } while (TRUE);  
  
    pthread_exit(0);  
}
```



Readers-Writers Problem

- A data set is shared among a number of concurrent processes
 - Readers – only read the data set; they do **not** perform any updates
 - Writers – can both read and write data
- Problem
 - Allow **multiple readers** to read at the same time
 - Only **one single writer** can access the shared data at the same time
- Several variations of how readers and writers are considered – all involve some form of **priorities**
- Shared data
 - Data set
 - Semaphore **rw_mutex** initialized to 1
 - Semaphore **mutex** initialized to 1
 - Integer **read_count** initialized to 0



Readers-Writers Problem (Cont.)

- The structure of a writer process and reader processes

/ rw_mutex: a mutual exclusion semaphore for the writers. It is also used by the first or last reader that enters or exits the critical section*/*

```
semaphore rw_mutex = 1;
```

/ mutex: ensure mutual exclusion when the variable read_count is updated. */*

```
semaphore mutex = 1;
```

/ read_counter: variable keeps track of how many processes are currently reading the object. */*

```
int read_count = 0;
```

```
do {  
    wait(rw_mutex);  
    ...  
    /* writing is performed */  
    ...  
    signal(rw_mutex);  
} while (true);
```

For writer

```
do {  
    wait(mutex);  
    read_count++;  
    if (read_count == 1)  
        wait(rw_mutex);  
    signal(mutex);  
    ...  
    /* reading is performed */  
    ...  
    wait(mutex);  
    read_count--;  
    if (read_count == 0)  
        signal(rw_mutex);  
    signal(mutex);  
} while (true);
```

1st reader process

Last reader process

For reader



Readers-Writers Problem (Cont.)

Sample code demonstration in class
(Reader-Writer using semaphores)

`posix_rw.c`



Readers-Writers Problem (Cont.)

posix_rw.c

Compile command:

gcc -o posix_rw posix_rw.c -lpthread

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#include <semaphore.h>
#include <unistd.h>
#include <signal.h>
```

```
#define TRUE 1
```

```
/* the thread runs in this function */
```

```
void *writer(void *param);
void *reader(void *param);
```

```
sem_t rw_sem;
sem_t mutex;
int read_count = 0;
```

```
int shared_data = 0;
pid_t pid;
```



Readers-Writers Problem (Cont.)

```
int main(int argc, char *argv[]) {                                posix_rw.c
    int i, scope;
    pthread_t writerID, readerID1, readerID2;                      /* the thread identifier */
    pthread_attr_t attr;                                           /* set of attributes for the thread */

    pid = getpid();

    //create semaphores, and initialize to 1
    sem_init(&rw_sem, 0, 1);
    sem_init(&mutex, 0, 1);

    /* get the default attributes */
    pthread_attr_init(&attr);

    /* create 3 threads */
    pthread_create(&writerID, &attr, writer, &writerID);
    pthread_create(&readerID1, &attr, reader, &readerID1);
    pthread_create(&readerID2, &attr, reader, &readerID2);

    /* Now join on each thread */
    pthread_join(writerID, NULL);
    pthread_join(readerID1, NULL);
    pthread_join(readerID2, NULL);
    return 0;
}
```



Readers-Writers Problem (Cont.)

posix_rw.c

```
void *writer(void *param) {
    int id = *(int*)param;
    printf("Writer Thread ID = %d\n", id);
    shared_data = 0;
    do {
        sem_wait(&rw_sem);
        shared_data++;
        printf("Writer: new data = %d\n", shared_data);
        sem_post(&rw_sem);
        sleep(1);
    } while (TRUE);

    pthread_exit(0);
}
```




Readers-Writers Problem (Cont.)

posix_rw.c

```
void *reader(void *param) {
    int id = *(int*)param;
    do {
        sem_wait(&mutex);
        read_count++;
        if (read_count == 1)
            sem_wait(&rw_sem);
        sem_post(&mutex);

        printf("\tReader_TID(%d): data(%d)\n", id, shared_data);

        if (shared_data > 10)
            break;

        sem_wait(&mutex);
        read_count--;
        if (read_count == 0)
            sem_post(&rw_sem);
        sem_post(&mutex);
        sleep(1);
    } while (TRUE);

    kill(pid, SIGINT); //send a signal to kernel to terminate the process
    pthread_exit(0);
}
```




Readers-Writers Problem Variations

- **First** variation – no reader kept waiting unless writer has permission to use shared object
- **Second** variation – once writer is ready, it performs the write ASAP
- Both may have **starvation** leading to even more variations
- Problem is solved on some systems by kernel providing **reader-writer locks**



Dining-Philosophers Problem



- Philosophers spend their lives alternating **thinking** and **eating**
- Don't interact with their neighbors, occasionally try to pick up 2 chopsticks (one at a time) to eat from bowl
 - Need both chopsticks to eat, then release both when done
- In the case of 5 philosophers
 - ▶ **Semaphore chopstick [5]** initialized to 1



Dining-Philosophers Problem Algorithm

- The structure of Philosopher *i*:

```
do {  
    wait (chopstick[i] );           //left chopstick  
    wait (chopstick[ (i + 1) % 5] ); //right chopstick  
    // eat state  
    signal (chopstick[i] );  
    signal (chopstick[ (i + 1) % 5] );  
    // think state  
} while (TRUE);
```



- What is the problem with this algorithm?
 - Although this solution guarantees that no two neighbors are eating simultaneously, but it could create a deadlock.
 - Suppose that all five philosophers become hungry at the same time and each grabs the left chopstick. No chopstick left. When each philosopher tries to grab the right chopstick, he/she will be delayed forever.
 - **Deadlock** – two or more processes are waiting **indefinitely** for an event that can be caused by only one of the waiting processes



Dining-Philosophers Problem Algorithm (Cont.)

■ Deadlock solutions

1. Allow **at most 4** philosophers to be sitting simultaneously at the table.
2. Allow a philosopher to pick up the chopsticks only if both are available (picking must be done in a critical section.)
3. Use an asymmetric solution -- an odd-numbered philosopher picks up first the left chopstick and then the right chopstick. Even-numbered philosopher picks up first the right chopstick and then the left chopstick.

■ A **deadlock-free** solution using **monitor** is provided later (appendix slide 74-80).

■ Any satisfactory solution must guard against the possibility that one of the philosophers may starve to death.

■ **A deadlock-free** does not necessarily eliminate the possibility of **starvation**.




Problems with Semaphores

- Deadlock and starvation are possible.
- Incorrect use of semaphore operations:
 - signal (mutex) wait (mutex)
 - wait (mutex) ... wait (mutex)
 - Omitting of wait (mutex) or signal (mutex) (or both)
- It is the responsibility of the application programmer to use semaphores correctly.



Deadlock and Starvation

- Let S and Q be two semaphores initialized to 1




- Can cause **starvation** and **priority inversion**
 - **Starvation** – **indefinite blocking**
 - ▶ A process may never be removed from the semaphore queue in which it is suspended)
 - **Priority Inversion** – Scheduling problem when lower-priority process holds a lock needed by higher-priority process



Deadlock and Starvation

■ priority inversion example:

- Assume three processes
 - ▶ L , M , and H with the priority order: $L < M < H$.
- Assume that
 - ▶ process H requires resource R , which is currently hold by process L .
 - ▶ process H would wait for L to finish using resource R .
- However, process M becomes runnable, and preempts process L .
- Consequence:
 - ▶ process M (with middle priority) has affected how long process H must wait for L to relinquish resource R .

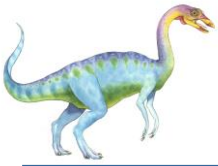




Deadlock and Starvation

□ Solution


- Use priority-inheritance protocol
 - All processes that are accessing resources needed by a higher-priority process inherit the higher priority until they are finished with the resources in question.
 - When they are finished, their priorities revert to their original values.
- This protocol solve the previous priority inversion problem



Monitors

■ Monitor (管程)

- A high-level abstraction that provides a convenient and effective mechanism for process synchronization
 - ▶ *Abstract data type*, internal variables only accessible via procedures
 - ▶ Only **one process may be active** within the monitor at a time
 - ▶ Can utilize **condition** variables to suspend or resume processes



monitor monitor-name

{

// shared variable declarations

procedure P1 (...) { }

procedure Pn (...) {.....}

Initialization code (...) { ... }

}



Monitors

<https://baike.baidu.com/item/%E7%AE%A1%E7%A8%8B/10503922?fr=aladdin>



管程



百度一下

网页

图片

视频

文库

知道

贴吧

资讯

地图

采购

更多

百度为您找到相关结果约8,930,000个

搜索工具

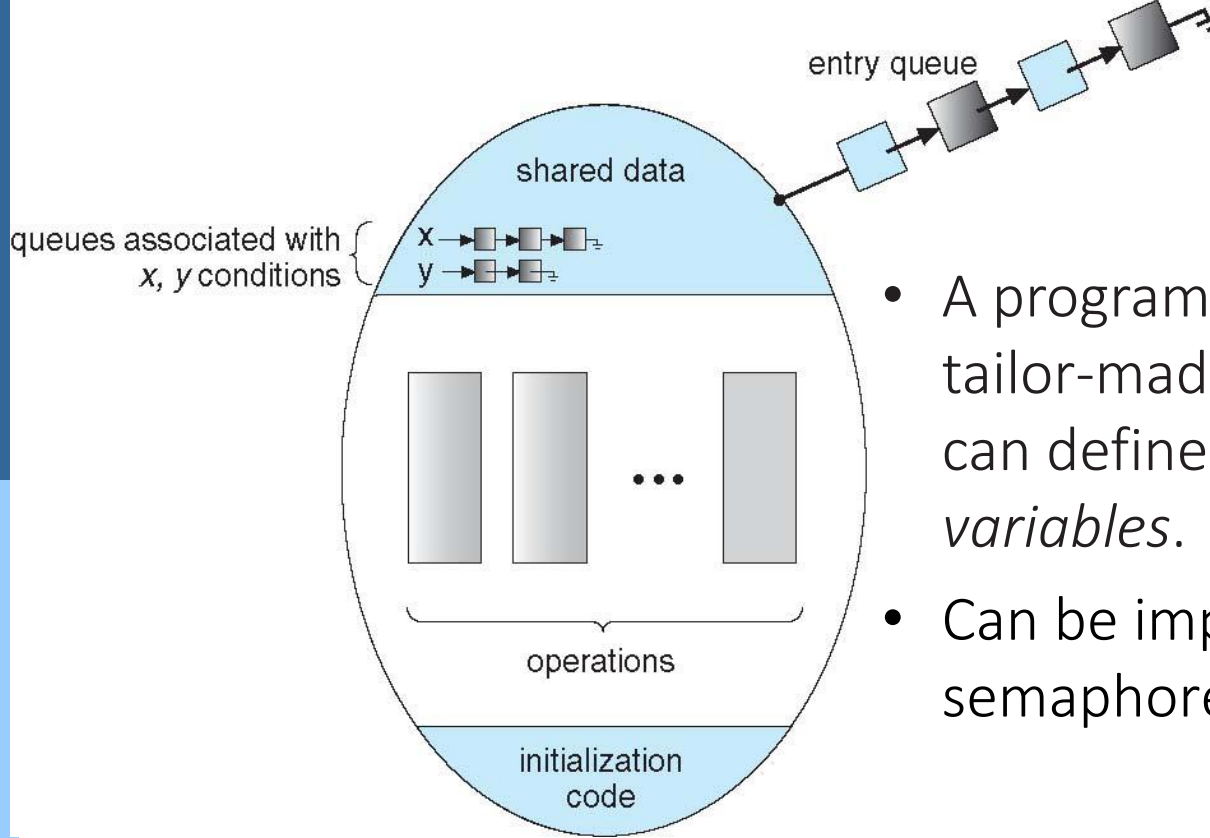
[管程 - 百度百科](#)

管程

管程在功能上和信号量及PV操作类似，属于一种进程同步互斥工具，但是具有与信号量及PV操作不同的属性。

简介 [管程结构分析](#) [管程的实现](#)

百度百科



- A programmer who needs to write a tailor-made synchronization scheme can define one or more **condition variables**.
- Can be implemented with semaphores

■ condition x, y;

■ Two operations are allowed on a condition variable:

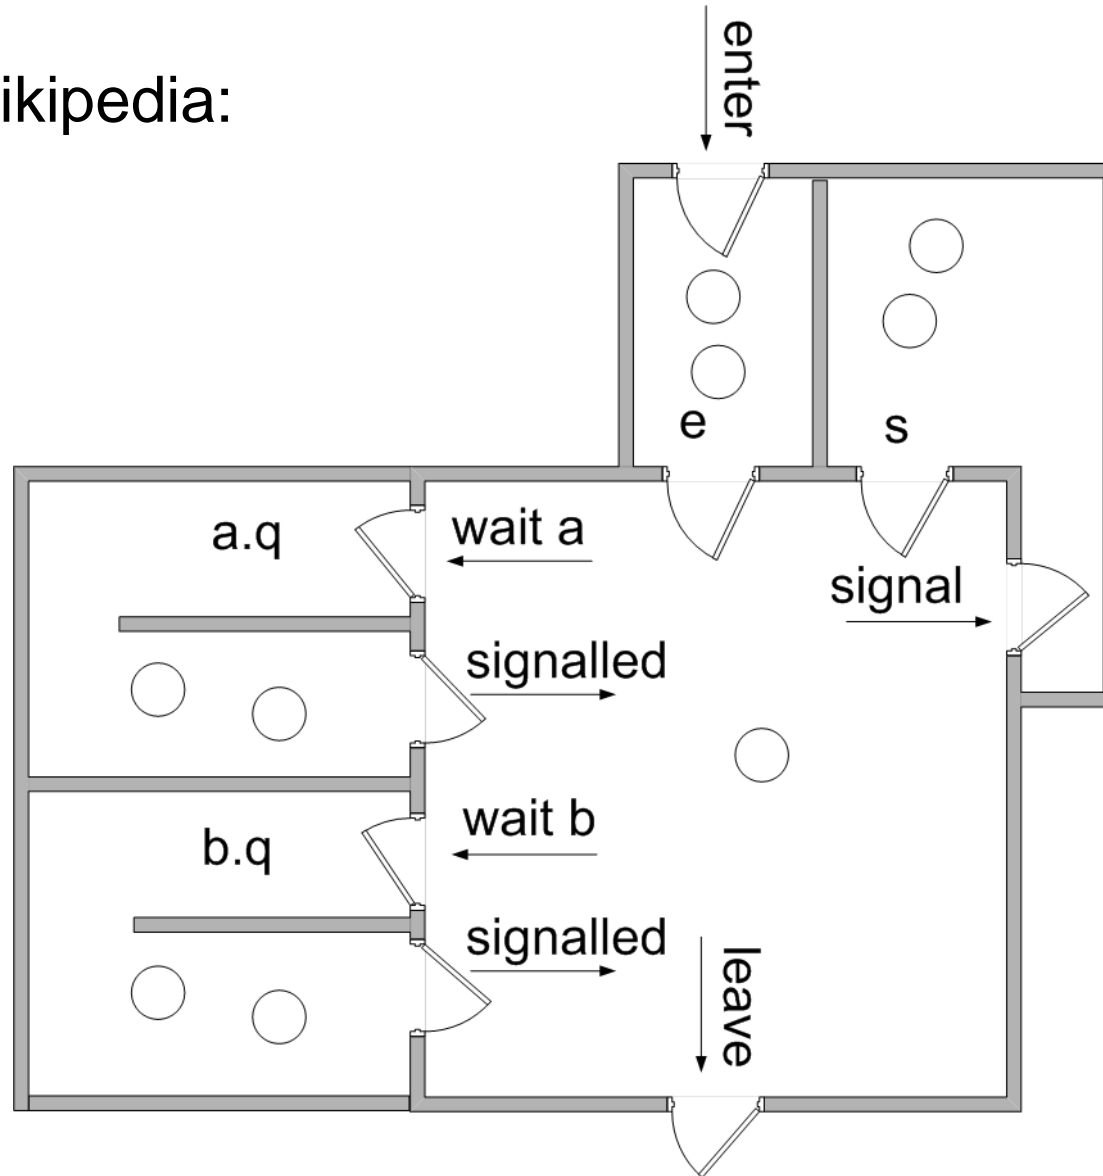
- x.wait() – a process that invokes the operation is suspended until x.signal()
- x.signal() – resumes one suspended process (if any)

If no x.wait() on the variable, then it has no effect on the variable



Monitor with Condition Variables

From Wikipedia:





End of Chapter 6 & 7



Appendices

- The appendix parts are for students who are interested in knowing more about the programming related to Dining Philosophers using Dining Philosophers using Monitor and Monitor implemented using semaphores introduced in this lecture.



Monitors

[https://en.wikipedia.org/wiki/Monitor_\(synchronization\)](https://en.wikipedia.org/wiki/Monitor_(synchronization))

In concurrent programming, a **monitor** is a synchronization construct that allows threads to have both **mutual exclusion** and the ability to **wait** (block) for a certain condition to become true.

Monitors also have a mechanism for signaling other threads that their condition has been met.

A monitor consists of a **mutex (lock)** object and **condition variables**.

A **condition variable** is basically a **container of threads** that are waiting for a certain condition.

Monitors provide a mechanism for threads to temporarily give up exclusive access in order to wait for some condition to be met, before regaining exclusive access and resuming their task.

Another definition of **monitor** is a **thread-safe** class, object, or module that uses wrapped mutual exclusion in order to safely allow access to a method or variable by more than one thread.

The defining characteristic of a monitor is that its methods are executed with **mutual exclusion**: At each point in time, at most one thread may be executing any of its methods.

By using one or more condition variables it can also provide the ability for threads to wait on a certain condition.



Condition Variables Choices

■ If process P invokes **`x.signal()`** , and process Q is suspended in **`x.wait()`** , what should happen next?

- Both Q and P cannot execute in parallel. (only one process can access) If Q is resumed, then P must wait.


■ Options include

- **Signal and wait** – P waits until Q either leaves the monitor or it waits for another condition
- **Signal and continue** – Q waits until P either leaves the monitor or it waits for another condition
- Both have pros and cons (优缺点) – language implementer can decide
- Monitors implemented in **Concurrent Pascal** compromise
 - ▶ P executing signal immediately leaves the monitor, and Q is immediately resumed
 - i.e., `signal()` is the last sentence
- Monitor is Implemented in other languages including C#, Java



Monitor Solution to Dining Philosophers

```
monitor DiningPhilosophers {
    enum { THINKING; HUNGRY, EATING} state[5] ;
    condition self[5];
    void pickup (int i) {
        state[i] = HUNGRY;
        test(i);
        while (state[i] != EATING) self[i].wait();
    }
    void putdown (int i) {
        state[i] = THINKING;
        test((i+4)%5); //left
        test((i+1)%5); //right
    }
    void test (int i) {
        if ((state[(i+4)%5] != EATING) && (state[i] == HUNGRY)
        &&
            (state[(i+1)%5] != EATING) ) {
            state[i] = EATING ;
            self[i].signal();
        }
    }
    initialization_code() {
        for (int i = 0; i < 5; i++)
            state[i] = THINKING;
    }
}
```





Solution to Dining Philosophers (Cont.)

- Each philosopher i invokes the operations `pickup()` and `putdown()` in the following sequence:

```
while (true) {  
    DiningPhilosophers.pickup(i);  
  
    EAT();  
  
    DiningPhilosophers.putdown(i);  
    Think();  
}
```

- No deadlock, but starvation is possible



Code demonstration in class

Dining Philosophers using Monitor

`posix_dpm1.c`



Solution to Dining Philosophers (Cont.)

Filename: posix_dpm1.c

compile command: g++ -o dpm1 posix_dpm1.c -lpthread

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#include <semaphore.h>
#include <unistd.h>

const int N = 5;
int phil_num[N]={0,1,2,3,4};
int eat_count[N] = { 0, 0, 0, 0, 0 };
enum State { THINKING, HUNGRY, EATING};

void *philosopher(void *num);
```



Solution to Dining Philosophers (Cont.)

//monitor struct for dining philosophers implemented with pthread condition variable

typedef struct Monitor {

public:

int state[N];

pthread_cond_t self[N]; //pthread condition variables

pthread_mutex_t mutex_lock;

void pickup(int i) { //external function

pthread_mutex_lock(&mutex_lock);

state[i] = HUNGRY;

test(i);

while (state[i] != EATING) {

int pthread_cond_wait(pthread_cond_t *cond, pthread_mutex_t *mutex)

pthread_cond_wait(&self[i], &mutex_lock);

}

pthread_mutex_unlock(&mutex_lock);

}

void putdown(int i) { //external function

pthread_mutex_lock(&mutex_lock);

state[i] = THINKING;

test((i + N-1) % N); //left

test((i + 1) % N); //right

pthread_mutex_unlock(&mutex_lock);

}

void test(int i) { //internal function

if ((state[(i + N-1) % N] != EATING) && (state[i] == HUNGRY) &&

(state[(i + 1) % N] != EATING)) {

state[i] = EATING;

pthread_cond_signal(&self[i]);

}

}

void initialization() {

pthread_mutex_init(&mutex_lock, NULL);

printf("use pthread conditional variables\n");

for (int i = 0; i < N; i++){

state[i] = THINKING;

pthread_cond_init(&self[i], NULL);

}

}

} DiningPhilosophers;



Solution to Dining Philosophers (Cont.)

DiningPhilosophers dp;

```
int main() {  
    int i;  
    pthread_t thread_id[N];  
  
    dp.initialization();  
  
    for(i = 0; i < N; i++)  
        pthread_create(&thread_id[i], NULL, philosopher, &phil_num[i]);  
  
    for(i = 0; i < N; i++)  
        pthread_join(thread_id[i], NULL);  
  
    for (i = 0; i < N; i++)  
        printf("Philosopher %d eat %d times\n", i + 1, eat_count[i]);  
  
    return 0;  
}
```



Solution to Dining Philosophers (Cont.)

```
void *philosopher(void *num) {  
    int loops = 0;  
    int i = *(int*)num;  
    while (loops < 5) {  
  
        //hungry and want to eat  
        printf("Philosopher %d is hungry\n",i+1);  
        dp.pickup(i);  
  
        //eating time  
        printf("\tPhilosopher %d is eating\n",i+1);  
        eat_count[i]++;  
        sleep(2);  
  
        //putting down forks  
        dp.putdown(i);  
        printf("\t\tPhilosopher %d is putting down forks\n",i+1);  
  
        //thinking time  
        printf("Philosopher %d is thinking\n",i+1);  
        sleep(1);  
        ++loops;  
    }  
    pthread_exit(0);  
}
```




Monitor Implementation Using Semaphores

■ Variables

- semaphore mutex; // (initially = 1)
 - ▶ A process must
 - execute wait(mutex) before entering the monitor and
 - execute signal(mutex) after leaving the monitor.
- semaphore next; // (initially = 0)
 - ▶ A signaling process must wait until the resumed process either leaves or waits, an additional semaphore, next is used.
 - ▶ The signaling processes can use next to suspend themselves.
- int next_count = 0;
 - ▶ count the number of processes suspended on semaphore

next



Monitor Implementation Using Semaphores

- Each external procedure F will be replaced by

```
wait(mutex);  
.....  
  body of F;  
  
if (next_count > 0) //some processes suspended  
  signal(next)  
else  
  signal(mutex);
```

Mutual exclusion within a monitor is ensured



Monitor Implementation – Condition Variables

- For each condition variable x , we have:

```
semaphore x_sem; // (initially = 0)
int x_count = 0;
```

- The process P calls x .signal() to wake up processes waiting for x .
The following code shows the implementation of the x .signal() method:

```
x_count++; // Because P is adding itself to the wait queue of x.
if (next_count > 0) // If other process in global wait queue
    signal(next); // allow it to execute inside monitor
else
    signal(mutex); // allow other process to enter monitor;
wait(x_sem); // then P falls asleep in wait queue for x.
// Some time later another process calls x.signal
x_count--; // then P leaves the wait queue for x after waking up.
```



Monitor Implementation (Cont.)

- The operation **`x.signal`** can be implemented as:

```
if (x_count > 0) { // If there is a Q in the wait queue for x:  
    next_count++; // P is going to add itself to the global wait queue  
    signal(x_sem); // P wakes up Q  
    wait(next);    // P falls asleep in the global wait queue  
                  // Some time later another process calls signal(next)  
    next_count--; // then P leaves the global wait queue after waking up.  
}
```

This implementation is applicable to the definitions of monitors given by both Hoare and Brinch-Hansen.



Resuming Processes within a Monitor

- If several processes queued on condition `x`, and `x.signal()` executed, which should be resumed?
 - FCFS (First Come First Served; FIFO) frequently not adequate
 - Priority based: conditional-wait construct of the form `x.wait(c)`
 - ▶ Where `c` is **priority number**
 - ▶ When `x.signal()` is called, the process in the wait queue of `x` which has the lowest number (highest priority) is going to be woken up first.

To illustrate this new mechanism, consider the ResourceAllocator monitor shown in next slide, which controls the allocation of a single resource among competing processes.



Code demonstration in class

Dining Philosophers using Monitor
implemented using semaphores

`posix_dpm2.c`




Code demonstration in class

```
/**
 * A pthread program illustrating POSIX dining-philosophers problem
 * using monitor struct and implementation using semaphores
 *
 * To compile:
 *      g++ -o dpm2 posix_dpm2.c -lpthread
 */

#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#include <semaphore.h>
#include <unistd.h>

const int N = 5;
int phil_num[N]={0,1,2,3,4};
int eat_count[N] = {0,0,0,0,0};
enum Philosophers_State { THINKING, HUNGRY, EATING};

void *philosopher(void *num);
```



Code demonstration in class (Cont.)

```
//monitor struct for dining philosophers
typedef struct Monitor {
    int state[N];
    //use semaphore to implement condition variables
    sem_t x_sem[N];
    int x_count[N];
    sem_t mutex;
    sem_t next;
    int next_count; //count the number of threads suspended on next

    //wait function for condition variable x_sem
    void wait(int i) {
        x_count[i]++;
        if (next_count > 0)
            sem_post(&next);
        else
            sem_post(&mutex);
        sem_wait(&x_sem[i]);
        x_count[i]--;
    }

    //signal function for condition variable x_sem
    void signal(int i){
        if (x_count[i] > 0){
            next_count++;
            sem_post(&x_sem[i]);
            sem_wait(&next);
            next_count--;
        }
    }

    //external function
    void pickup(int i) {
        sem_wait(&mutex);
        state[i] = HUNGRY;

        test(i);

        while (state[i] != EATING)
            wait(i);

        if (next_count > 0)
            sem_post(&next);
        else
            sem_post(&mutex);
    }
}
```

```
//external function
void putdown(int i) {
    sem_wait(&mutex);

    state[i] = THINKING;

    test((i + N-1) % N); //left
    test((i + 1) % N); //right

    if (next_count > 0)
        sem_post(&next);
    else
        sem_post(&mutex);
}

//internal function
void test(int i) {
    if ((state[(i + N-1) % N] != EATING) && (state[i] == HUNGRY)
        && (state[(i + 1) % N] != EATING)) {
        state[i] = EATING;
        signal(i);
    }
}

void initialization() {
    printf("use semaphore to implement condition variables\n");
    for (int i = 0; i < N; i++) {
        state[i] = THINKING;
        x_count[i] = 0;
        sem_init(&x_sem[i], 0, 1); //1
    }
    next_count = 0;
    sem_init(&next, 0, 0); //0
    sem_init(&mutex, 0, 1);
}
} DiningPhilosophers;
```




Code demonstration in class (Cont.)

DiningPhilosophers dp;

```
int main() {  
    int i;  
    pthread_t thread_id[N];  
  
    dp.initialization();  
  
    for(i = 0; i < N; i++)  
        pthread_create(&thread_id[i], NULL, philosopher, &phil_num[i]);  
  
    for(i = 0; i < N; i++)  
        pthread_join(thread_id[i], NULL);  
  
    for (i = 0; i < N; i++)  
        printf("Philosopher %d eat %d times\n", i + 1, eat_count[i]);  
    return 0;  
}
```



Code demonstration in class (Cont.)

```
void *philosopher(void *num) {
    int i = *(int*)num;
    int loops = 0;
    while (loops < 5) {

        //hungry and want to eat
        printf("Philosopher %d is hungry\n",i+1);
        dp.pickup(i);

        //eating time
        printf("\tPhilosopher %d is eating\n",i+1);
        eat_count[i]++;
        sleep(2);

        //putting down forks
        dp.putdown(i);
        printf("\t\tPhilosopher %d is putting down forks\n",i+1);

        //thinking time
        printf("Philosopher %d is thinking\n",i+1);
        sleep(1);
        ++loops;
    }
    pthread_exit(0);
}
```



A Monitor to Allocate Single Resource

```
monitor ResourceAllocator {  
    boolean busy;  
    condition x;  
    void acquire(int time) {  
        if (busy)  
            x.wait(time);  
        busy = TRUE;  
    }  
    void release() {  
        busy = FALSE;  
        x.signal();  
    }  
    initialization code() {  
        busy = FALSE;  
    }  
}
```

```
monitor ResourceAllocator {  
    boolean busy;  
    pthread_cond_t x;  
    void acquire(int time) {  
        if (busy)  
            pthread_cond_wait(&x, NULL);  
        busy = TRUE;  
    }  
    void release() {  
        busy = FALSE;  
        pthread_cond_signal(&x);  
    }  
    initialization code() {  
        pthread_cond_init(&x, NULL);  
        busy = FALSE;  
    }  
}
```



Synchronization Examples

- Windows
- Linux
- Pthreads



Synchronization Examples

- Windows
- Inside kernel
 - Uses **interrupt masks ()** to protect access to global resources on uniprocessor systems
 - Uses **spinlocks** on multiprocessor systems
 - Spinlocking-thread will never be preempted
- Outside kernel
 - Provides **dispatcher objects** (event, mutex, semaphore, timer, and thread) for thread synchronization
 - Synchronization using mutexes, semaphores, events, and timers
 - ▶ **Events**
 - An event acts much like a condition variable
 - ▶ **Timers** notify one or more thread when time expired
 - A document online: [Introduction to Kernel Dispatcher Objects - Windows drivers | Microsoft Docs](#)



Windows Synchronization

- (Mutex)Dispatcher objects either **signaled-state** (object available, thread will not block) or **non-signaled state** (thread will block)




Figure 7.8 Mutex dispatcher object.

If a thread tries to acquire a mutex dispatcher object that is in a non-signaled state, that thread will be suspended and placed in awaiting queue for the mutex object. When the mutex moves to the signaled state (because another thread has released the lock on the mutex), the thread waiting at the front of the queue will be moved from the waiting state to the ready state and will acquire the mutex lock.



Windows Synchronization Cont.

- A **critical-section object** is a user-mode mutex that can be used without kernel intervention.
 - Document: [win32/critical-section-objects.md at docs · MicrosoftDocs/win32 · GitHub](#)

A *critical section object* provides synchronization similar to that provided by a mutex object, except that a critical section can be used only by the threads of a single process. Critical section objects cannot be shared across processes.

Event, mutex, and semaphore objects can also be used in a single-process application, but critical section objects provide a slightly faster, more efficient mechanism for mutual-exclusion synchronization (a processor-specific test and set instruction). Like a mutex object, a critical section object can be owned by only one thread at a time, which makes it useful for protecting a shared resource from simultaneous access. Unlike a mutex object, there is no way to tell whether a critical section has been abandoned.



Linux Synchronization

- Linux:
 - Prior to kernel Version 2.6, disables interrupts to implement short critical sections
 - Version 2.6 and later, fully preemptive(抢占式)
- Linux provides:
 - atomic integers
 - ▶ all math operations using atomic integers are performed without interruption
 - Mutex locks
 - Spinlocks and semaphores
 - reader-writer versions of Spinlocks and semaphores



Linux Synchronization Cont.

- On single-cpu system, spinlocks are replaced by **enabling and disabling kernel preemption**

Single Processor	Multiple Processors
Disable kernel preemption	Acquire spin lock
Enable kernel preemption	Release spin lock

Linux provides two simple system calls:

`preempt_disable()`

`preempt_enable()`

for disabling and enabling kernel
preemption

when **a lock is held for a short duration**, **spinlocks & enabling and disabling kernel preemption** are used;

When **a lock must be held for a longer period**, **semaphores or mutex locks** are used.



Pthreads Synchronization

- Pthreads API is **OS-independent**
- It provides:
 - mutex locks
 - condition variable

```
#include <pthread.h>
pthread_mutex_t mutex;
pthread_cond_t cond_var;

//create a mutex lock
pthread_mutex_init(&mutex, NULL);
//create a condition
pthread_cond_init(&cond_var, NULL);

pthread_cond_wait(&cond, NULL);
pthread_cond_signal(&cond);

/* acquire the mutex lock */
pthread_mutex_lock(&mutex);

/* critical section */

/* release the mutex lock */
pthread_mutex_unlock(&mutex);
```



Pthreads Synchronization

- Non-portable extensions include:
 - read-write locks
 - Spinlocks
 - semaphore



Pthreads Synchronization

■ Read-write locks functions

```
#include <pthread.h>
```

```
int pthread_rwlock_init(pthread_rwlock_t *rwptr, const pthread_rwlockattr_t *attr);  
int pthread_rwlock_destroy(pthread_rwlock_t *rwptr);
```

```
int pthread_rwlockattr_init(pthread_rwlockattr_t *attr);  
int pthread_rwlockattr_destroy(pthread_rwlockattr_t *attr);
```

```
int pthread_rwlock_rdlock(pthread_rwlock_t *rwptr);  
int pthread_rwlock_wrlock(pthread_rwlock_t *rwptr);  
int pthread_rwlock_unlock(pthread_rwlock_t *rwptr);
```

■ Spinlocks functions

```
#include <pthread.h>
```

```
int pthread_spin_init(pthread_spinlock_t *lock, int pshared);  
int pthread_spin_lock(pthread_spinlock_t *lock);  
int pthread_spin_trylock(pthread_spinlock_t *lock);  
int pthread_spin_unlock(pthread_spinlock_t *lock);  
int pthread_spin_destroy(pthread_spinlock_t *lock);
```



Pthreads Synchronization

■ Semaphore functions:

```
#include <semaphore.h>
sem_t sem;
```

```
/* Create the semaphore and initialize it to 1 */
sem_init(&sem, 0, 1);
```

```
/* acquire the semaphore */
sem_wait(&sem);
```

```
/* critical section */
```

```
/* release the semaphore */
sem_post(&sem);
```



Alternative Approaches

- Explore various features provided in both programming languages and hardware that support designing thread-safe concurrent applications.
 - Transactional Memory
 - OpenMP
 - Functional Programming Languages



Transactional Memory

- A **memory transaction** is a sequence of read-write operations to memory that are performed **atomically**.

```
void update() {  
    acquire();  
    /* read/write memory */  
    release();  
}
```

Deadlock is
possible



```
void update() {  
    atomic{  
        /* read/write memory */  
    }  
}
```

- (1) The transactional memory system is responsible for guaranteeing atomicity.
- (2) Because **no locks** are involved, **deadlock is not possible**.



OpenMP

- **OpenMP** is a set of compiler directives and API that support parallel programming.

```
void update(int value)
{
    #pragma omp critical
    {
        count += value
    }
}
```

- The code contained within the `#pragma omp critical` directive is treated as a critical section and performed atomically.



Functional Programming Languages

- Functional programming languages offer a different paradigm than procedural languages in that they do not maintain state.
- Variables are treated as immutable (不变的) and cannot change state once they have been assigned a value.
- There is increasing interest in functional languages such as Erlang and Scala for their approach in handling data races.



Summary

- Given a collection of cooperating sequential processes that share data, mutual exclusion must be provided to ensure that a critical section of code is used by only one process or thread at a time.
- Hardware provides several operations that ensure mutual exclusion, but too complicated for most developers to use.
- Mutex locks and semaphores: to solve various synchronization problems and can be implemented efficiently, especially atomic operations are available.
- Various synchronization problems are discussed:
 1. bounded-buffer problem
 2. readers–writers problem
 3. dining-philosophers problem



Summary cont.

- **Monitors** provide a synchronization mechanism for **sharing abstract data types**.
- A **condition variable** provides a method by which a **monitor function** can **block its execution** until it is **signaled to continue**.
- OS provides support for synchronization. For example, Windows, Linux, provide mechanisms such as **semaphores**, **mutex locks**, **spinlocks**, and **condition variables** to control access to shared data.
- The Pthreads API provides support for **mutex locks** and **semaphores**, as well as **condition variables**.
- Several alternative approaches focus on synchronization for multicore systems.
 - using transactional memory;
 - using the compiler extensions offered by OpenMP;