

Chapter 9: Main Memory





Outline

Background

Memory Allocation

Contiguous Memory Allocation

Frame

- ▶ Page Table
- ▶ Address Translation

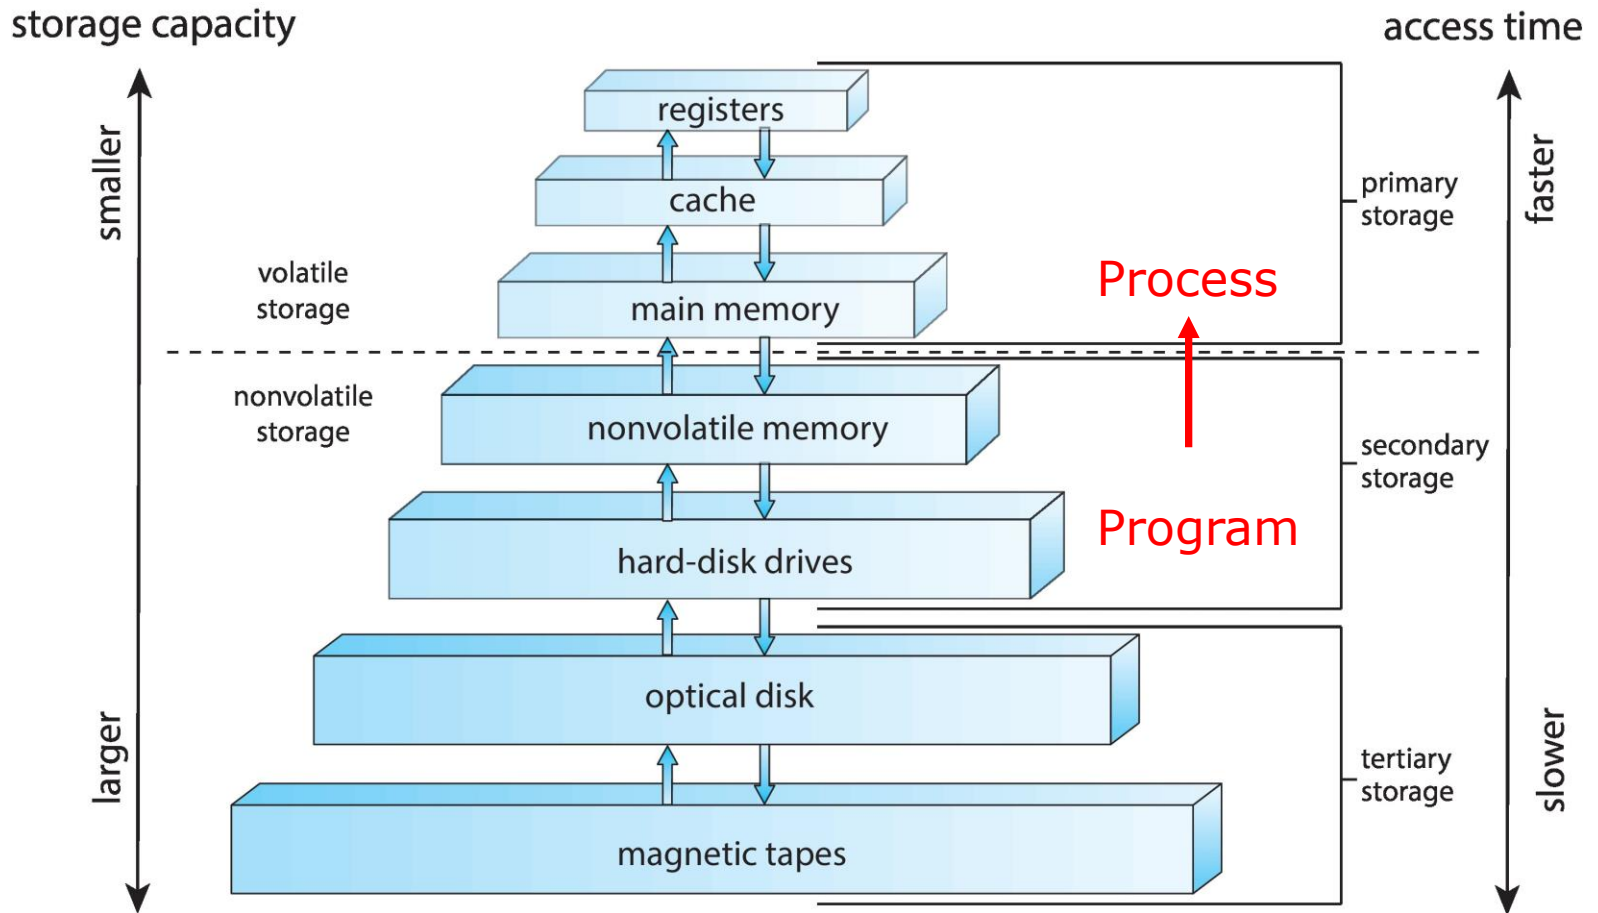
Page Table Schema

- ▶ Hierarchical Page Table
- ▶ Hash Table
- ▶ Inverted Table

Examples



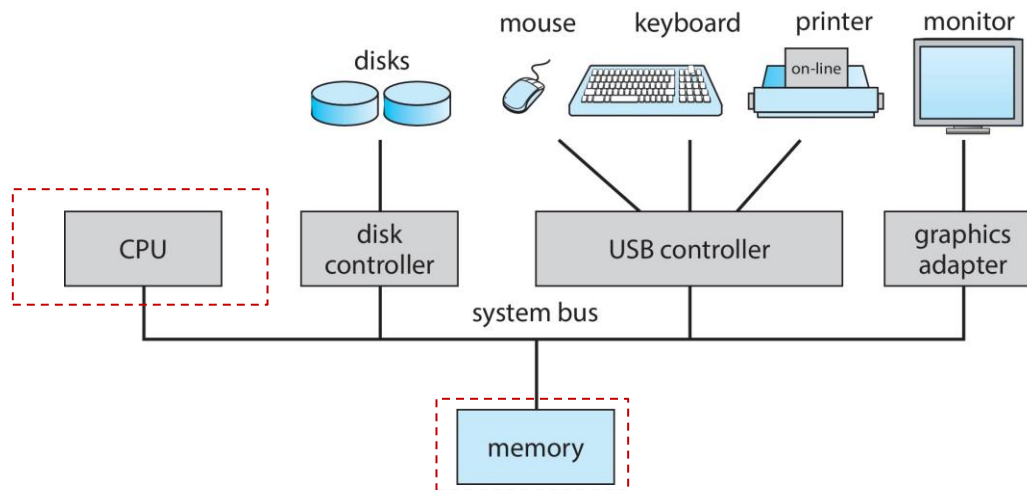
Review: Storage Hierarchy



Then, how CPU accesses the process in the memory?

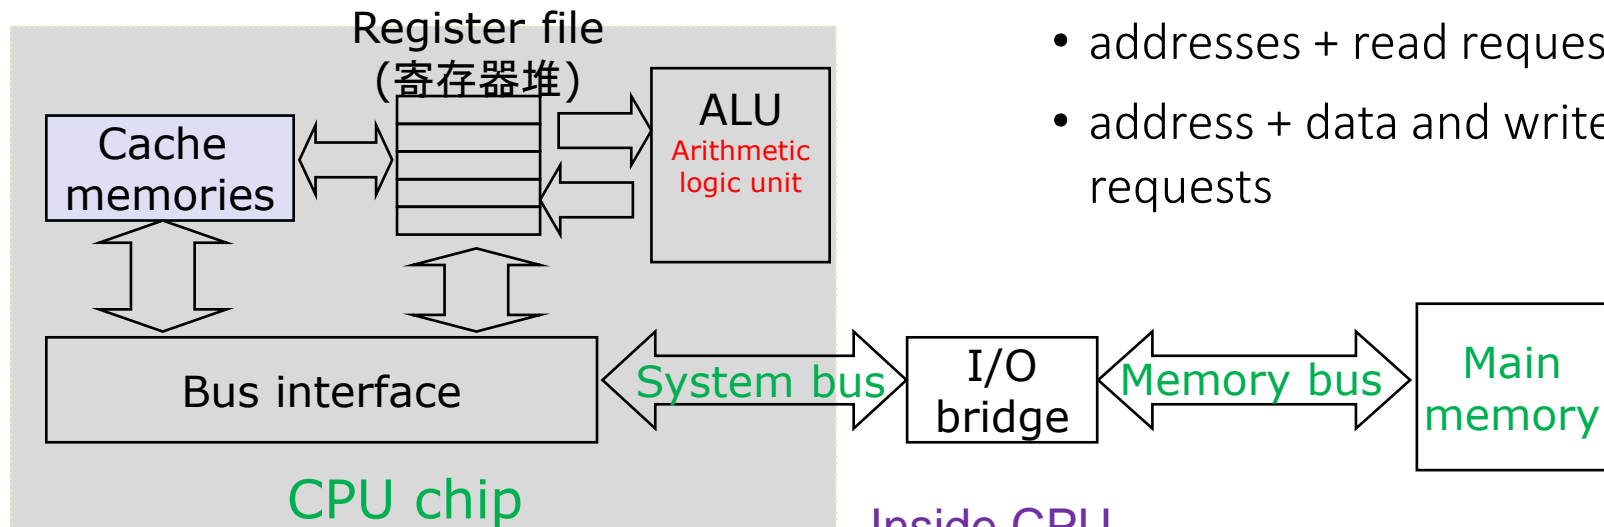


Relationship of CPU and Memory



A PC System Structure

- Main memory and registers are only storage that CPU can access directly
- Register access is done in one CPU clock (or less)
- Main memory can take many cycles, causing a stall (delay)
- Memory unit only sees from CPU a stream of:
 - addresses + read requests, or
 - address + data and write requests





Protection

CPU notifies memory **where** to read or write

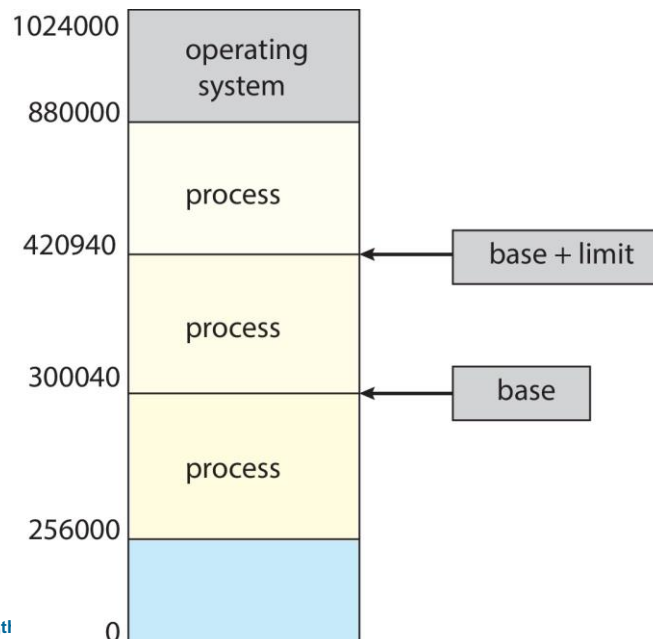
Is this “where” reliable?

Will the instructions go to other process’s land?

Protection is needed

An old but simple method

- ▶ a pair of **base register** and **limit register** for each process to address of the process





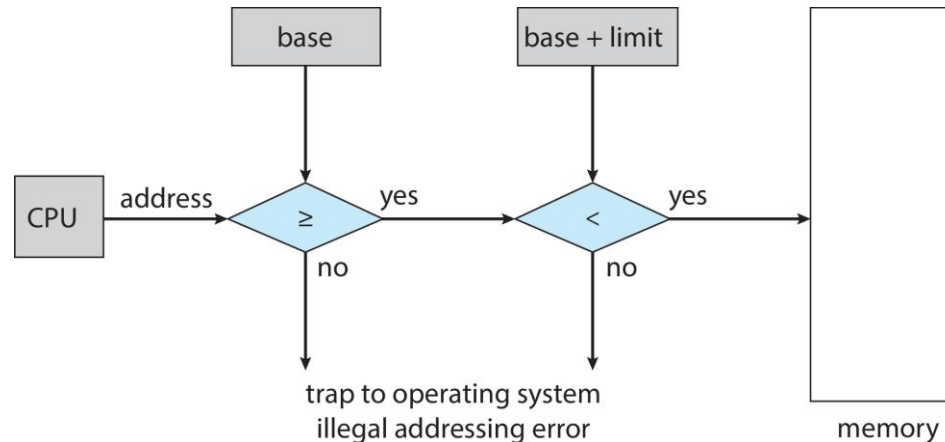
Hardware Address Protection

Two registers are needed

Base register

Limit register

CPU must check every memory access generated in user mode to be sure it is between base and limit for that user

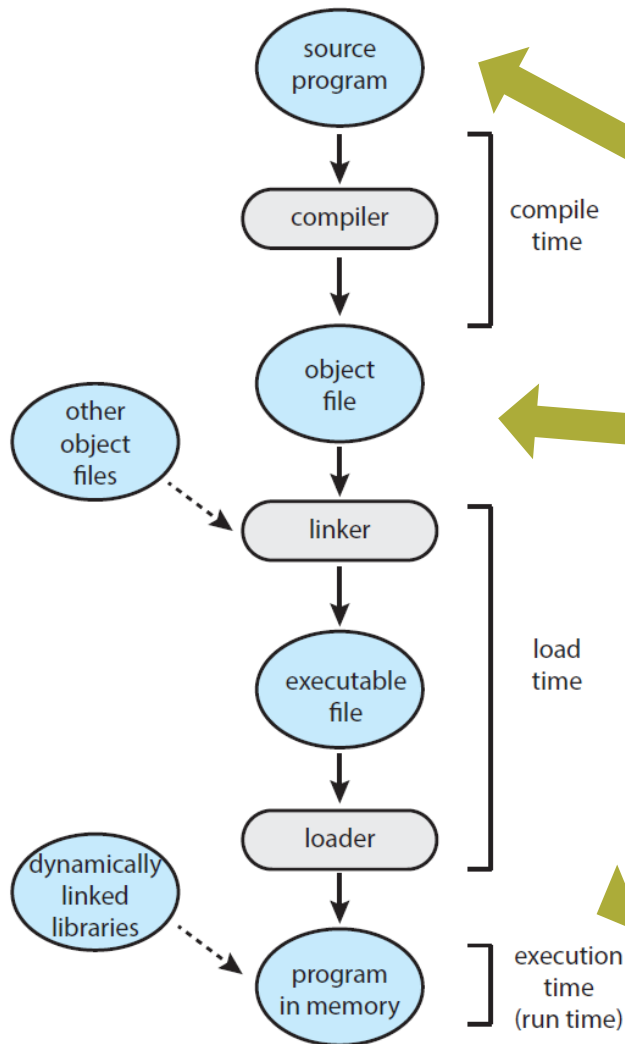


The instructions to load base and limit registers are privileged

No process can spy on another one by changing its own base register to be the same as other's.



Address Binding in A Typical Multistep Processing of a User Program



How are process addresses determined in memory?

Addresses are represented in different ways at different stages of a program's life

Source code addresses usually **symbolic**

- ▶ i.e. `int count; &count`

Compiled code addresses **bind** to relocatable addresses (**addresses relative to the start of the code**).

- ▶ i.e. "14 bytes from beginning of this code module"

Linker or loader will **bind** relocatable addresses to absolute addresses (**addresses relative to the lowest address 0000 in memory**).

- ▶ i.e. 74014

Each **binding** maps one address space to another address space



Linking: Dynamic vs Static

Static linking –system libraries and program code are combined by the loader into the binary image

Every program includes library: wastes memory

Dynamic linking –linking is postponed until execution/run time

Operating system checks if routine (子程序) is in process memory address, If not in address space, add to address space

Small piece of code, **stub (存根)**, is used to locate the appropriate memory-resident library routine

- ▶ Stub replaces itself with the address of the routine and executes the routine

Particularly useful for **libraries which are used by many processes**

- ▶ Keep only one copy of the library in memory which is shared by all the processes.
- ▶ known as **shared libraries** (.dll on Windows and .so on Linux, and .a is Linux static library)



Binding of Instructions and Data to Memory

Types of systems **determine when** the binding to memory address happens

Compile time: Absolute address is used in code

- ▶ Example: **embedded systems** with no kernel.
- ▶ Disadvantage: code must be re-compiled if starting location in memory changes

Load time: Relocatable code is generated if memory location is not known at compile time

- ▶ Example: **very old multiprogramming systems**
- ▶ Program's code has a relocation table of relative addresses of all the things in the code
- ▶ The real memory address is calculated according to the given location.
- ▶ Disadvantage: the address in memory is fixed after the program is loaded into memory.



Binding of Instructions and Data to Memory

Types of systems determine when the binding to memory address happens

Execution time: Binding is delayed until run time.

- ▶ Example: **most computers today**
- ▶ A process can be moved **during its execution** from one memory segment to another
- ▶ Need **hardware support** for address maps
 - e.g., **MMU**

This lecture talks about the binding **in execution time**.



Logical vs. Physical Address Space

Memory management is to

bind a logical address space to a separate physical address space

► Logical address

- generated by the Program/CPU
- also referred to as virtual address
- all the addresses used by the CPU
- Addressed in bytes
- Size of space depends on the CPU bits

► Physical address

- all the addresses used by the memory hardware
- Addressed in bytes
- Size of space depends on real size of physical memory

Thinking:

If a computer has 32-bit CPU and memory address is 16-bit, how big is the logical address space and the physical memory address?



Logical vs. Physical Address Space

Logical and physical addresses are the same in
compile-time address binding, and
load-time address-binding schemes

logical (virtual) and physical addresses differ in
execution-time address-binding scheme

This lecture talks about the binding of logical and
physical addresses at execution time.

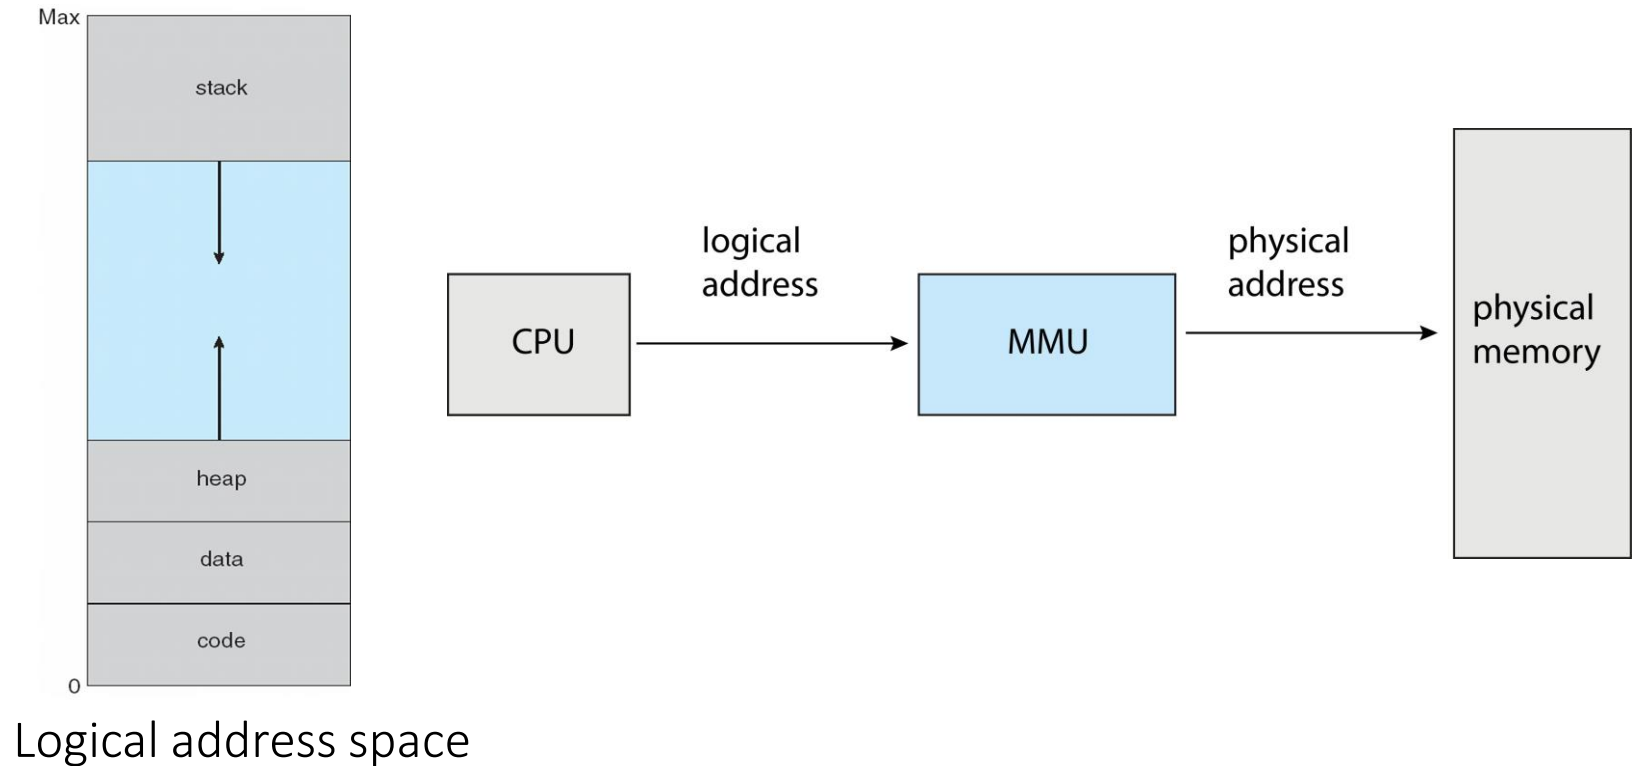


Memory-Management Unit (MMU)

MMU

Memory Management Unit

Hardware device that at run time maps (translates) logical (virtual) address to physical address



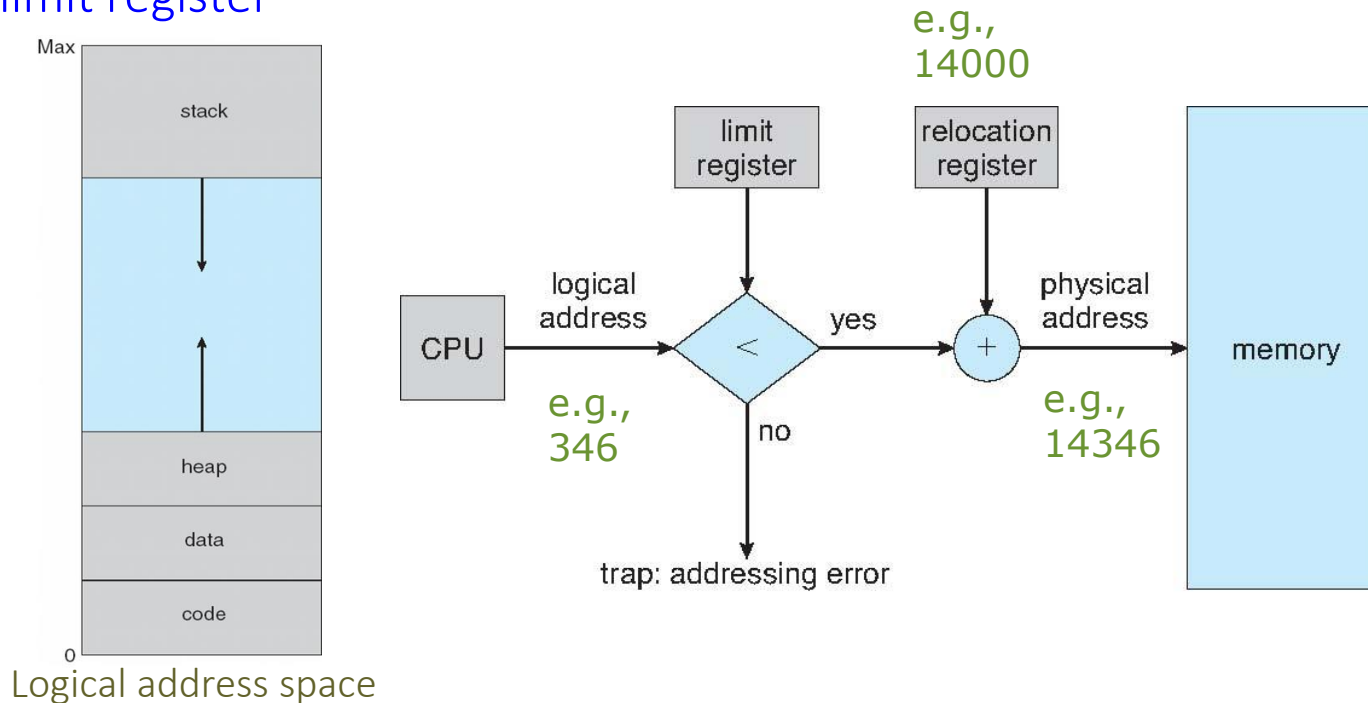


Memory-Management Unit (Cont.)

A simple mapping scheme

relocation register

limit register



Logical addresses of all processes start at 0000

Two processes can use the same logical addresses, these addresses will be mapped to different places in the physical memory

The kernel changes the relocation and limit register during a **context switch**



Dynamic Loading

Should the entire program be loaded in to the memory to start the execution?

No!

Code is loaded from the hard disk into memory only when it is needed, **on demand**

Dynamic loading

The loading of a process routine when it is called rather than when the process is started.

Why

- Better memory-space utilization

- Unused routine is never loaded

- Some large amounts of code handle infrequently occurring cases

How

- All routines are kept on disk in relocatable load format

- Implemented through program design by programmers, not OS

- OS can help by providing libraries to implement dynamic loading



Memory Allocation

Memory can be shared by many processes, including **kernel** and **user processes**


It should be allocated efficiently

Resident **kernel** usually is held in low memory


User processes is held in high memory

Memory for each process can be

Contiguous (early method)



Not contiguous (modern method)





Continuous Partition

Memory is divided into multiple partitions

Size of each partition is fixed

- ▶ **Multiple-partition allocation**

- ▶ Disadvantages:

- Degree of multiprogramming is limited by number of partitions
- Memory is wasted if a process does not fully use its partition

Size of each partition is variable

- ▶ **Variable-partition**

- ▶ Propositional to a given process' size



Continuous Partition: Variable Partition

Kernel maintains information about
allocated partitions
free partitions (hole)

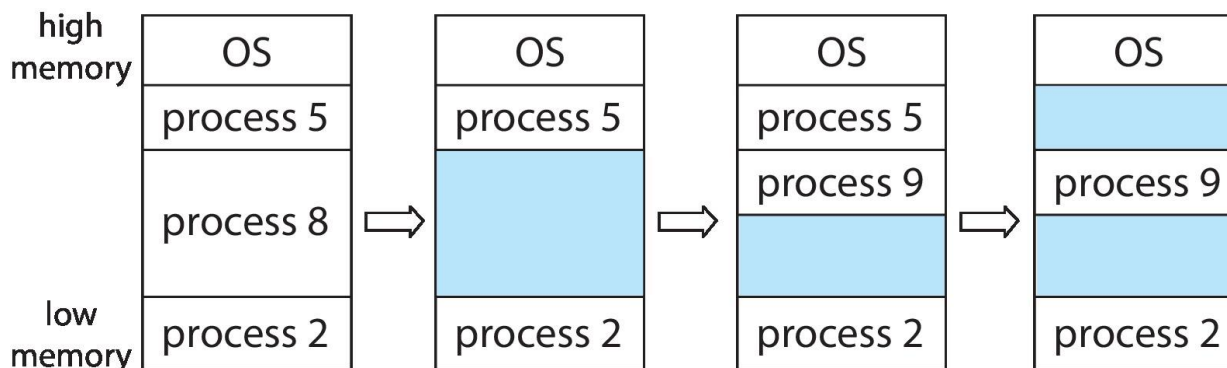
Hole

block of **available memory**

holes of various size are **scattered** throughout memory

When a process arrives, it is allocated memory from a hole
large enough to accommodate it

When a process exits, its partition is freed, adjacent free
partitions are combined (coalesced/merged)



Thinking:
Memory after **process 9**
is released?



Continuous Partition: Variable Partition

How to satisfy a request of size n from a list of **free** holes?

Three allocation algorithms

First-fit: Allocate the *first* hole that is big enough

Best-fit: Allocate the *smallest* hole that is big enough;

Worst-fit: Allocate the *largest* hole

Comparison of three algorithms

Speed

- ▶ **First-fit** > **best-fit** and **worst-fit**
 - Both **best-fit** and **worst-fit** need to search the whole list

Memory usage

- ▶ **Best-fit** > **first-fit** > **worst-fit** (theoretically)
- ▶ **First-fit** > **best-fit** (case studies)
 - Reason: **best-fit** produces many small holes that can never be used



Fragmentation

External Fragmentation

Total memory space exists to satisfy a request, but it is not contiguous

Small holes might be wasted because the holes are too small to fit any process.

Internal Fragmentation


Allocated memory may be slightly larger than requested memory; this size difference is memory internal to a partition, but not being used.

Any memory which is allocated to a process but that the process did not ask for is wasted.

P3: 31K
 $20+20>31$,
but not contiguous



External fragmentation



Internal fragmentation

P3: 31K
 $2^5\text{K}=32\text{K}$ will be allocated,
but 1k is not used



Fragmentation

First fit analysis reveals that given N blocks allocated, $0.5 N$ blocks lost to fragmentation

50-percent rule

- ▶ the memory lost to fragmentation is about 50% the size of the allocated memory

1/3 of RAM may be unusable



Fragmentation

Solutions to fragmentation problems

External fragmentation

- ▶ Compaction
 - Shuffle memory contents to place all free memory together in one large block
 - possible only if relocation is dynamic, and is done at execution time
 - Pay attention to I/O problem while there is I/O operations on the affected memory
- ▶ Use **non-contiguous** memory allocation

Internal fragmentation

- ▶ No way
- ▶ Memory is allocated in the block of 2^n bytes, rather than byte by byte.



Non-continuous Partition

Physical address space of a process can be **noncontiguous**
Process is allocated physical memory whenever the **total free memory blocks are enough**

Avoids external fragmentation

Frame

Divide **physical** memory into fixed-sized blocks

Each block is called **a frame**

The size of each frame is power of 2, between 512 bytes and 16 Mbytes

Kernel keeps track of all free frames

Thinking:

If RAM's address uses 16 bits, each frame is 4K, how many frames are in the memory?




Non-continuous Partition

0000	Frame 0	0000 0000 0000 1111 1111 1111
0001	Frame 1	0000 0000 0000 1111 1111 1111
0010	Frame 2	0000 0000 0000 1111 1111 1111
1111	Frame 15	0000 0000 0000 1111 1111 1111

16 bits: 2^{16} , 4K: $4 * 2^{10} = 2^{12}$

Total frames: $2^{16} / 2^{12} = 2^4$



Thinking:

If RAM's address uses 16 bits, each frame is 4K, how many frames are in the memory?




Non-continuous Partition

0000	Frame 0	0000 0000 0000 0000
0001	Frame 1	0000 1111 1111 1111 0001 0000 0000 0000
0010	Frame 2	0001 1111 1111 1111 0010 0000 0000 0000
		0010 1111 1111 1111
1111	Frame 15	1111 0000 0000 0000 1111 1111 1111 1111

16 bits: 2^{16} , 4K: $4 * 2^{10} = 2^{12}$

Total frames: $2^{16} / 2^{12} = 2^4$



Thinking 1:

If a datum is in the address

0100 0000 0000 0001?

Which frame is it in? where is it in the frame?

Thinking 2:

If a process's size is 8K + 1, how many frames are allocated? Is there internal fragmentation?



Non-continuous Partition

How a logical address maps to a physical address?

Pages

Divide logical memory into blocks of same size called pages

Page size = frame size

To run a program of size N pages, need to find N free frames and load program

Frames can be anywhere(non-contiguous) in RAM!


A page table is needed to translate logical to physical addresses

Thinking:

If logical space uses 32 bits, each frame is 4K, how many pages are there in logical space? $2^{32} / 2^{12} = 2^{20}$



Address Translation Scheme: Page Table




logical
memory

0	1
1	4
2	3
3	7

page table

frame
number



physical
memory

The data in physical memory is **not in the same order** as in the logical address space of the process because **a page** can be stored **in any frame** in RAM!


Each process has a page table

Process access the memory indicated by its page table.

Page table is

part of the process's **PCB**.

invisible to the process itself.




Address Translation Scheme: Page Table

Given a logical address space 2^m and **page (frame) size 2^n** , there are $2^m/2^n = 2^{m-n}$ pages in a process's logical address space

Logical address generated by CPU (m bits) is divided into

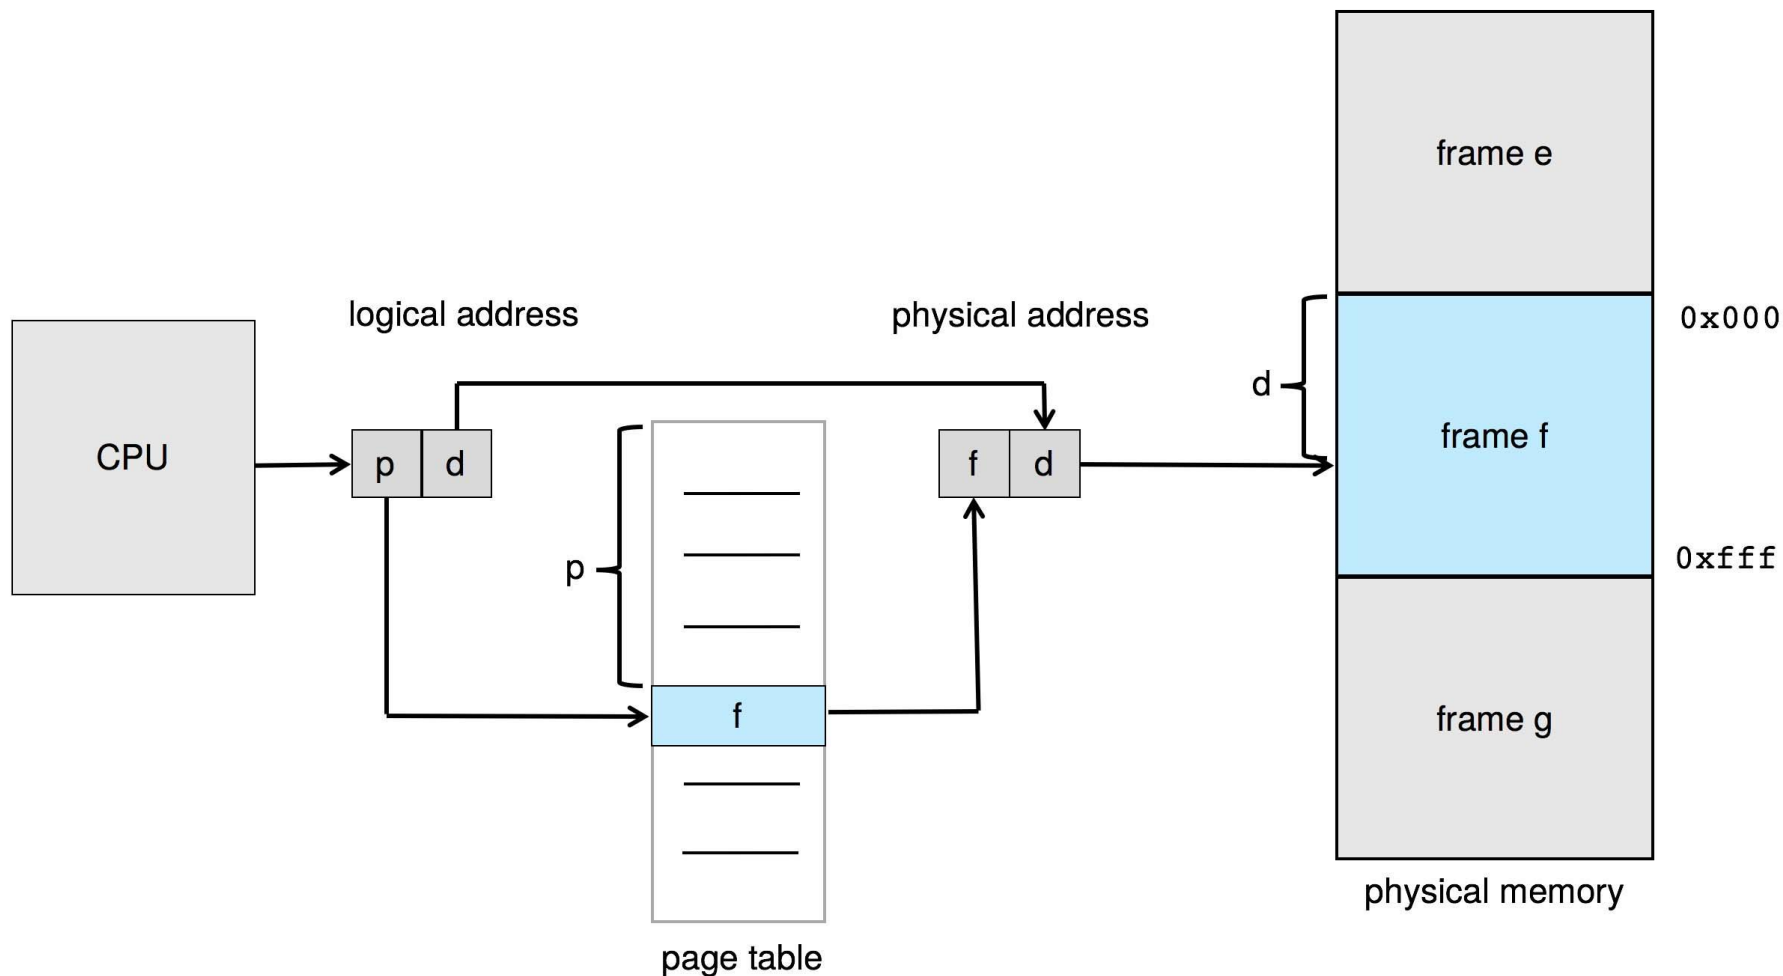
Page number (p) – used as an index into a **page table** which contains base address (frame number) of each page in physical memory

Page offset (d) – combined with base address (frame number) to define the physical memory address that is sent to the memory unit





Paging Hardware



Page number p to frame number f translation



An Paging Example

Logical address: 4 bits (16 bytes)

Each page: 4 bytes (2 bits)

$m = 4$ and $n = 2$

Page No.	Page Offset
p	d
$m - n$	n

Logical address:

Page 00	0000	0	a
	0001	1	b
	0010	2	c
	0011	3	d
Page 01	0100	4	e
	0101	5	f
	0110	6	g
	0111	7	h
Page 10	1000	8	i
	1001	9	j
	1010	10	k
	1011	11	l
Page 11	1100	12	m
	1101	13	n
	1110	14	o
	1111	15	p

logical memory

Physical address: 5 bits (32 bytes)

Each frame: 4 bytes (2 bits)

Physical address:

Frame 000	0	00000
		00001
		00010
		00011
Frame 001	4	00100
		00101
		...
		...
Frame 010	8	01000
		01001
		01010
		01011
Frame 011	12	01100
		01101
		01110
		01111
Frame 100	16	10000
		10001
		10010
		10011
Frame 101	20	10100
		10101
		10110
		10111
Frame 110	24	11000
		11001
		11010
		11011
Frame 111	28	11100
		11101
		11110
		11111

physical memory

Thinking:

Given a logical address is 1011,
what is its physical address?



Paging -- Calculating internal fragmentation

Suppose a process:

page size = 2,048 bytes

process size = 72,766 bytes

Thinking 1:

How many frames are needed to store this process?

Thinking 2:

What is the loss of memory due to the internal fragmentation?

Thinking 3:

What is the worst case in the internal fragmentation?

Thinking 4:

Is it a good idea to reduce the size of page to minimize the internal fragmentation?



Page Size


Fact: page sizes growing over time

Some systems support two page sizes – 8 KB and 4 MB

X86-64: **4 KB** (common), **2 MB** (“huge” for servers), **1GB** (“large”)



Free Frame Allocation to A New Process



Before allocation

After allocation



Implementation of Page Table

Page table is

part of the process's **PCB**

kept **inside the kernel's own memory**

- ▶ Two registers are in the **MMU**
 - **Page-table base register (PTBR)** points to the **start position** of page table in memory
 - **Page-table length register (PTLR)** indicates **length** of the page table
- ▶ These two registers are used by the MMU to find the page table in main memory.

Thinking:

How many memory accesses are needed for one data access?



Implementation of Page Table

Problem

Two memory accesses

- ▶ one for the **page table** (logical address -> physical address)
- one for the **data / instruction** (physical address -> data)

Too slow!

Solution

TLB(Translation Look-aside Buffer)

- ▶ also called **associative memory**
- ▶ hardware cache **inside the MMU**.



TLB

TLB table

Page #	Frame #

Given a logical page number p

If p is in TLB (called **TLB hit**)

- ▶ get corresponding frame number from it

If p is not in the TLB table (called **TLB miss**)

- ▶ get frame number from page table in memory
- ▶ then update the TLB to contain this page # and frame #

— Reason

- » if that page was just used by the process, it is probably going to be used again soon.



TLB

TLB is

typically small

- ▶ 64 to 1,024 entries

Shared by multiple processes

- ▶ Store **address-space identifiers (ASIDs)** in each TLB entry to uniquely identify each process to provide address-space protection for each process
- ▶ Reduce the fresh time at every context switch between processes

How is a TLB entry replaced by a new page-frame pair?

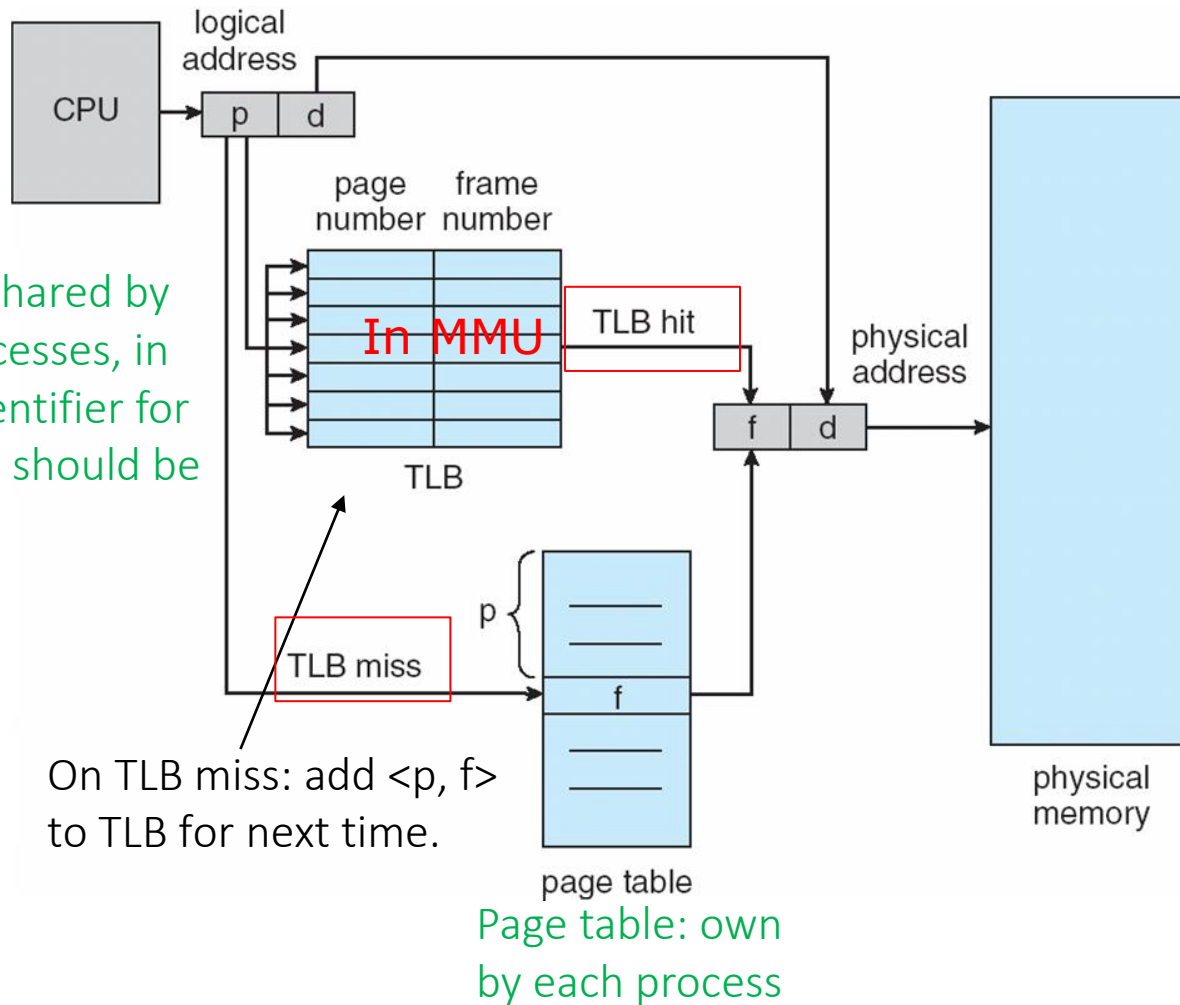
Replacement **policies** must be considered

Some entries can be **wired down** (never be replaced) for permanent fast access (for key kernel code)



Paging Hardware With TLB

TLB: can be shared by multiple processes, in that case, identifier for each process should be added.





Effective Access Time: An Example

Hit ratio α is

percentage of times that a page number is found in the TLB

Effective memory access time

$$EAT = \alpha \times (c+m) + (1-\alpha) \times (c+2m)$$

c : TLB access time, m : memory access time

An example

α : 80%, c : 10 nanoseconds, m : 90 nanoseconds

- ▶ $EAT = 0.80 \times (10+90) + 0.20 \times (10+90 \times 2) = 118 \text{ ns}$
- ▶ 18% slowdown in access time

If α : 99%

$$EAT = 0.99 \times (10+90) + 0.01 \times (10+90 \times 2) = 100.9\text{ns}$$

only 0.9% slowdown in access time



Shared Pages Example

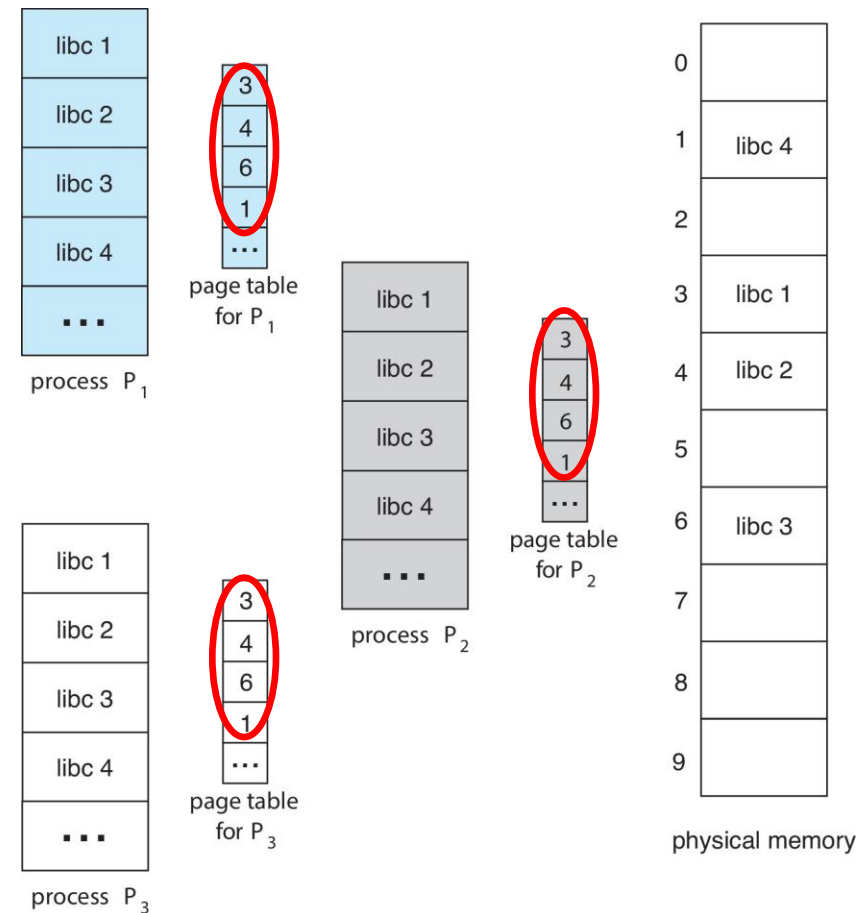
Some code (included libraries) can be shared by multiple processes

E.g., the code of the printf function is shared by all the processes executing C programs, so there is only one copy of the code of printf in physical memory.

Shared code

one copy of read-only code is shared among processes.

can stay at logical locations in the logical address spaces of the different processes,



Process p_1 , p_2 and p_3 share libc 1, 2, 3, 4



Structure of the Page Table

Problem

Memory structures for paging can get huge

For example

- ▶ logical address space: 32-bit ($m=32$), page size: 4 KB ($n=12$)
- ▶ page table length: $2^{m-n} = 2^{20}$ (1M)
- ▶ if each entry is 4 bytes \rightarrow 4 MB of physical address space
- ▶ need **a single hole** of free memory which is big enough!
- ▶ **More: 4M is only for one process.**

Solutions

Hierarchical Paging

Hashed Page Tables

Inverted Page Tables

Page No.	Page Offset
p	d
$m - n$	n



Hierarchical Page Tables

Split one (big) page table into multiple (small) page tables


two-level page table

► Outer page table

- Each entry indicates the position of small page table

► Inner (small) page tables

- Each small page table (also called page of page table) has the page size
- Different small page tables can be put in different places
- Only the part of the page table that the process requires right now needs to be in memory!





Two-Level Paging Example

logical address space: 32-bit ($m=32$), page size: 4 KB ($n=12$)

Originally

- ▶ page table length: $2^{m-n} = 2^{20}$

Two levels

p_1 : a 10-bit page number (1KB)

- ▶ an index into the **outer page table**

p_2 : a 10-bit page offset (1KB)

- ▶ the displacement within the page of the **inner page table**

Page No.	Page Offset
p	d


$m - n$
20 **12**

page number		page offset
p_1	p_2	d
10	10	12

logical address


p_1	p_2	d
-------	-------	---

p_1



outer page table

p_2




page of page table

physical address

f	d
---	---

d



- **Advantage**: save memory
 - only part of the page table needs to be in memory
- **Disadvantage**: three memory access

Address translation process



Paging Example

8 Frames

$m=4, n=2$

--	--

2	2
00	00
01	01
10	10
11	11

0000
0001
0010
0011
0100
0101
0110
0111
1000
1001
1010
1011
1100
1101
1110
1111

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

logical memory

0	5	101
1	6	110
2	1	001
3	2	010

page table

0		00000 00001 00010 00011 00100 00101 ...
1	i j k l	
2	m n o p	
3		
4		
5	a b c d	
6	e f g h	
7		11111

physical memory

p_1	p_2	d
-------	-------	-----

1 1 2

0 0
1 1

0	Table 0 address
1	Table 1 address

Outer table

Page table 0

0	5	101
1	6	110

Page table 1

0	1	001
1	2	010

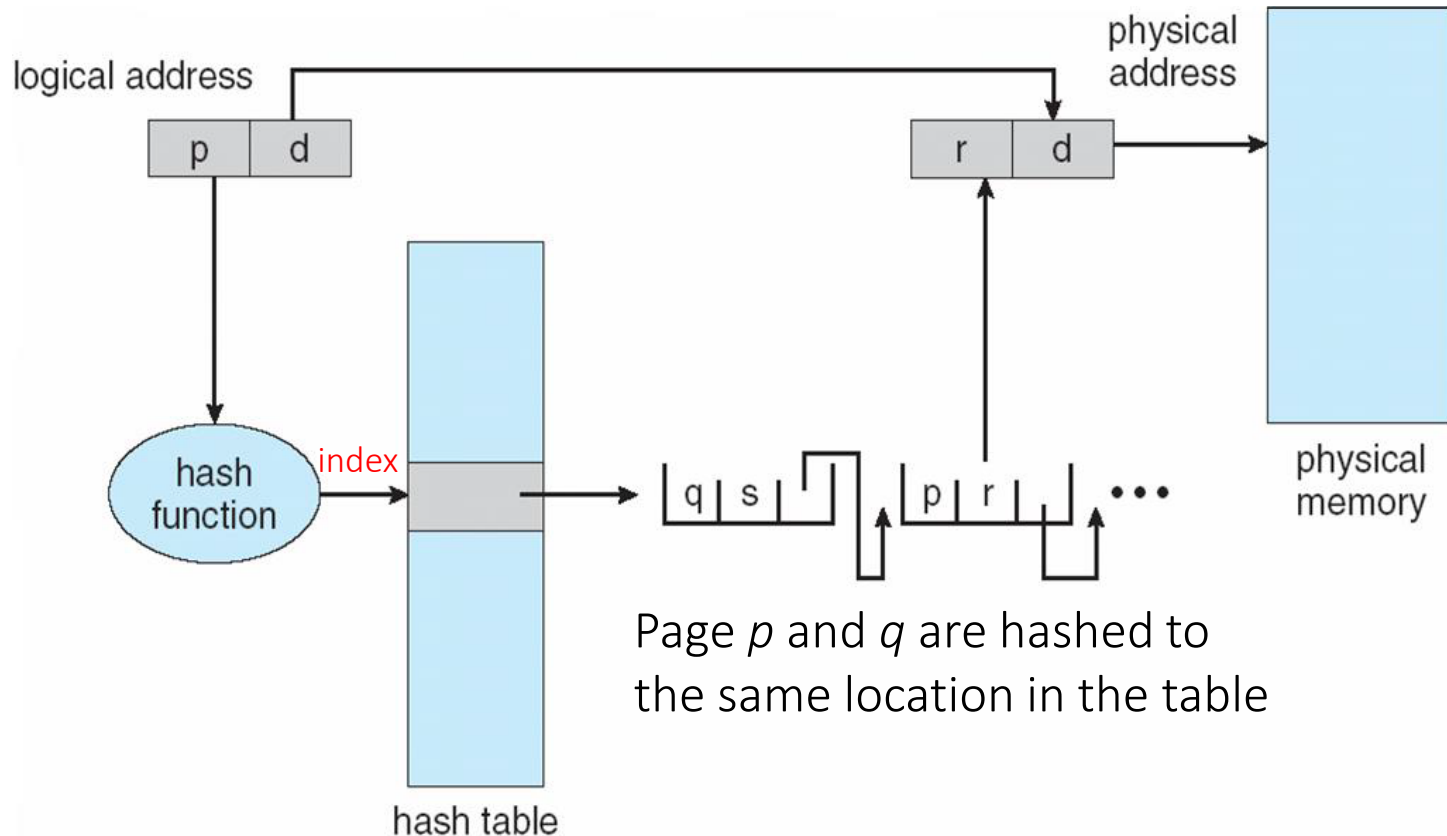
Separated page tables

Thinking:

Given a logical address 1011,
how to find its physical address?



Hashed Page Tables




- An entry in the hash table contains a chain of elements hashing to the same location
- Each element contains
 1. the virtual page number (e.g., *q*)
 2. the value of the mapped page frame (e.g., *s*)
 3. a pointer to the next element.
- This method is commonly used for the address with more than 32 bits



Paging Example

8 Frames

$m=4, n=2$



2	2
00	00
01	01
10	10
11	11

0000
0001
0010
0011
0100
0101
0110
0111
1000
1001
1010
1011
1100
1101
1110
1111

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

logical memory

0	5
1	6
2	1
3	2

page table


0	
4	i j k l
8	m n o p
12	
16	
20	a b c d
24	e f g h
28	

00000
00001
00010
00011
00100
00101
...


11111

physical memory

0 00
1 01
2 10
3 11



0, 2 -> 0
1, 3 -> 1

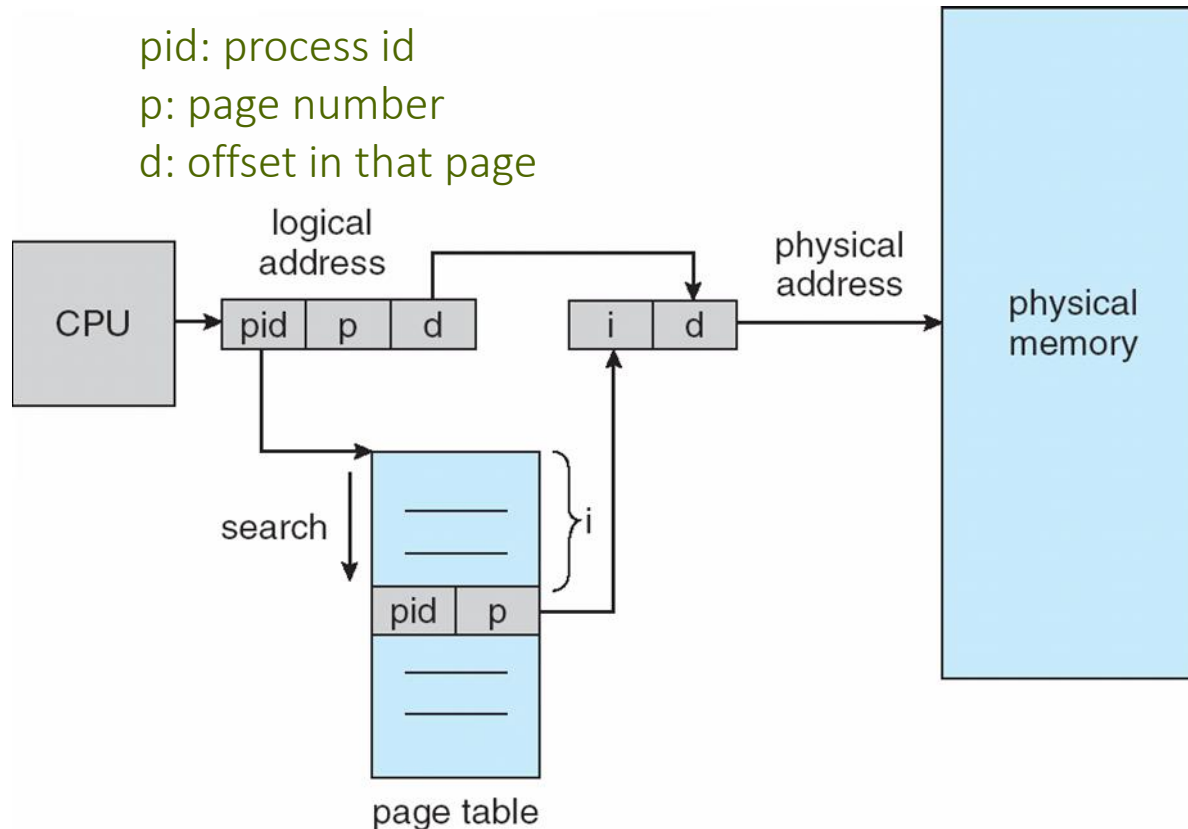


Thinking:

Given a logical address 1011,
how to find its physical address?



Inverted Page Table



The length of page table is the same as the size of physical memory
All processes share one page table
The page table is actually a reflection of layout of the physical memory frames



Inverted Page Table Example

vpn: virtual page number

Thinking:

Advantage??? Disadvantage???

- Memory
- Search Time
- Shared memory implementation

pid	vpn	offset
0	0x1	0x123

Index	PID	VPN
0x0	1	0xA63
...
0x18F1B	0	0x1
0x18F1C	3	0x31AB
...

Search one by one until
vpn = 0x1 and pid = 0

0x18F1B	0x123
ppn	offset

ppn: physical frame(page) number



Example: The Intel 32 and 64-bit Architectures

Dominant industry chips

IA-32 architecture (x86 or i386) (CPUs are 32-bit)

IA-64 architecture (x86-64 or amd64) (CPUs are 64-bit)

Many variations in the chips



Example: The Intel IA-32 Architecture

Supports both segmentation (pieces of partition) and segmentation with paging

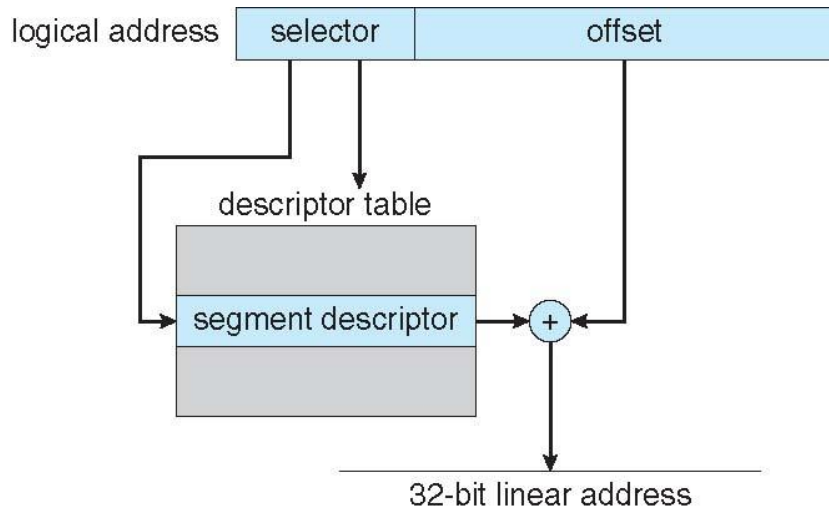


Each process has up to 16 K segments

- ▶ First partition has up to 8 K segments (private to process, kept in **local descriptor table (LDT)**)
- ▶ Second partition has also up to 8K segments (shared among all processes, kept in **global descriptor table (GDT)**)



Each segment can be 4 GB



Logical to Physical Address Translation in IA-32



Segmentation unit



Two sizes of pages are used

Paging unit



Example: Intel x86-64


64 bits is enormous (> 16 exabytes)

In practice only implement 48 bit addressing

Page sizes of 4 KB, 2 MB, 1 GB

Four levels of paging hierarchy

Can also use PAE (page address extension) so virtual addresses are 48 bits but support 52 bits physical addresses (4096 terabytes)





Example: ARM Architecture

Dominant mobile platform chip (Apple iOS and Google Android devices for example)

ARM

Modern, energy efficient, 32-bit CPU

Two page sizes: 4 KB or 16 KB pages (two levels)

1 MB or 16 MB pages (one level)

One-level paging for sections, two-level for smaller pages

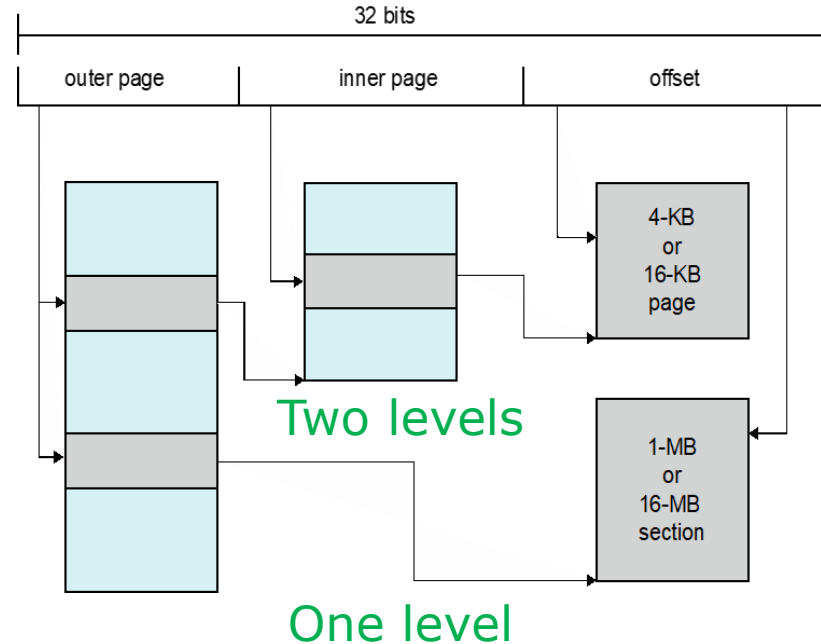
Two levels of TLBs

inner level has two micro TLBs (one data, one instruction)

outer is single main TLB

inner is checked first, on miss, outers are checked, and on miss again, page table is accessed

64 bit CPUs now exist too



End of Chapter 9

