

Chapter 10: Virtual Memory





Outline

Background

Demand Paging

Copy-on-Write

Page Replacement

Allocation of Frames

Thrashing

Allocating Kernel Memory

Other Considerations

Example



Objectives

- To describe the benefits of a virtual memory system
- To explain the concepts of demand paging, page-replacement algorithms, and allocation of page frames
- To discuss the principle of the working-set model
- To explore how kernel memory is managed




Multiprogramming

Multiprogramming

Multiple processes can reside in the memory at the same time

Problem: if a program needs to be run, but there is no enough free memory

One solution: swap out (换出) a process and swap in (换入) the target process



Backing store (后备存储器)

- is fast hard disk large enough to accommodate copies of all process memory images for all users
- has direct access to memory images



Swapping

A **complete** process can be **swapped** temporarily out of memory to a backing store, and then brought back into memory for continued execution

A **swapper** manipulates **entire** process swapping

Advantage

- Increase multiprogramming

Disadvantage

- Context switch time can then be **very high**

- Total **context switch time** includes **swapping time** for the whole process

Can the swapping time be reduced by reducing the size of memory swapped?




Swapping with **Paging**

When memory is low, unused pages are swapped to disk
(**instead of whole processes**)

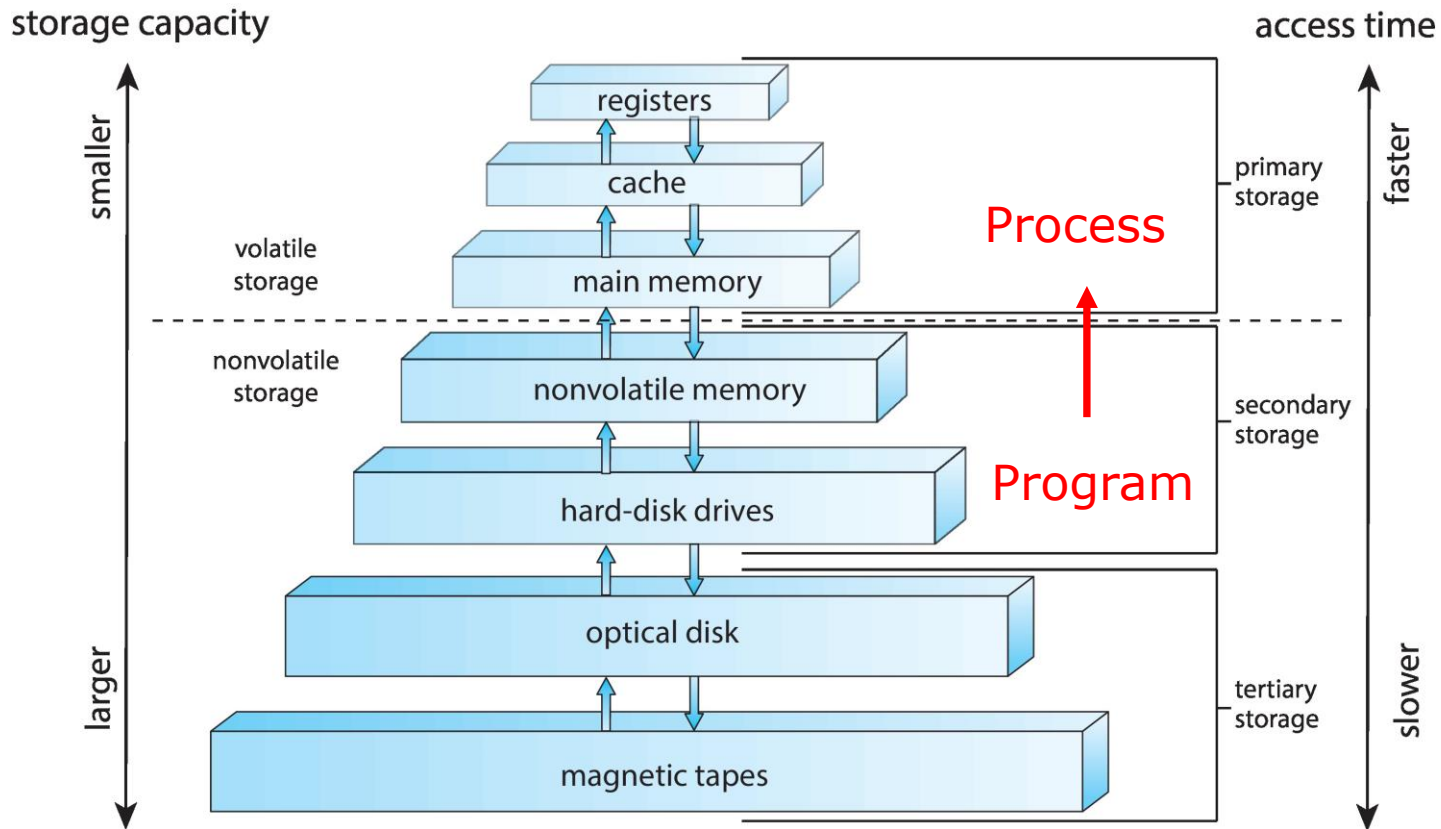
How to determine which pages are unused?

What happens when a process suddenly tries to access a page that was swapped out?





Review: Storage Hierarchy



Chapter 9: how CPU accesses the process in the memory?

Chapter 10: Are all the code and data in the program needed at the same time in the memory?



Background

Program can be partially loaded into memory

Error handling code, unusual routines, large data structures may not be needed for most of time

Program length is no longer constrained by limits of physical memory


Each program takes less memory → more programs run at the same time

Less I/O needed to load programs into memory → each user program runs faster

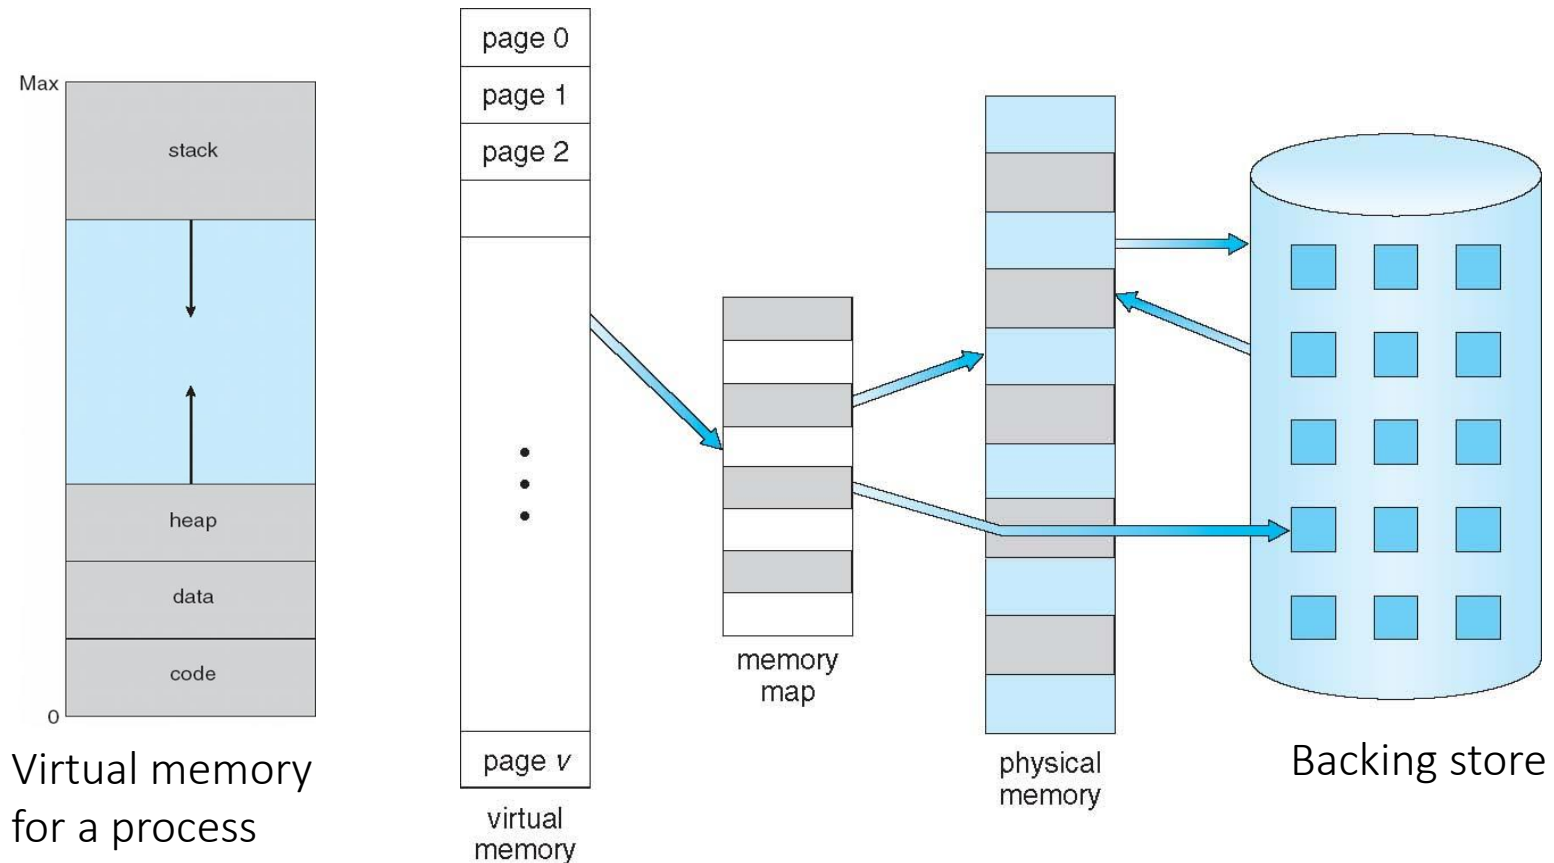
Virtual memory – separation of user logical memory from physical memory

Logical address space can therefore be much larger than physical address space

Only needed pages are loaded into memory



Virtual Memory That is Larger Than Physical Memory



On a modern computer,
the logical address space (virtual memory) is 2^{64} bytes = 16 billion GB, but
the physical address space (RAM) is only a few GB (a few hundred GB for
the biggest computers on earth).



Demand Paging

Demand paging

Could bring a page into memory only when it is needed

Advantages:

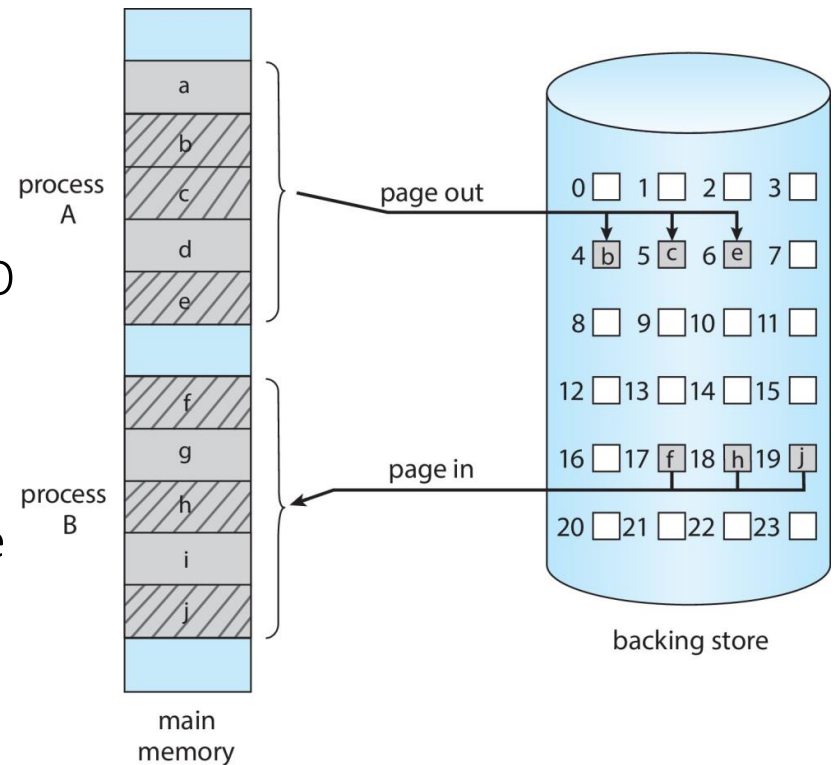
- ▶ Less I/O needed, no unnecessary I/O
- ▶ Less memory needed
- ▶ Faster response
- ▶ More users/processes

Page is needed (called a **reference** to the page

invalid reference might be caused by

- ▶ wrong address \Rightarrow process abort
- ▶ not-in-memory \Rightarrow bring to memory

A **pager** can swap in and out one page at a time (not a whole process like **swapper** does).





Demand Paging: Basic Concepts

The pager brings in only those pages into memory that the process actually wants to use

Question: How to determine that set of pages?

- ▶ Hardware support: MMU
 - If pages needed are already memory resident
 - » No difference from non-demand paging
 - If pages needed are not in memory
 - » Need to detect and load the pages into memory from storage



Valid-Invalid Bit

To check if a reference is valid or not, associate each page table entry with a **valid–invalid bit**

v: the page is in-memory

i: the page is

- ▶ either not-in-memory, or
- ▶ a wrong logical address


Initially valid–invalid bit is set to **i** on all entries

Frame #	valid-invalid bit
	v
	v
	v
	i
...	
	i
	i

page table



An Example: Some Pages Are Not in Main Memory



Pages 0, 2, and 5 are in the memory

Pages 1, 3, 4, 6, and 7 are not in the memory



Page Fault

Page fault

Caused when a process makes a **reference to a page which is not in RAM**,

MMU detects that the page is marked as **invalid** in the process's page table,

Two possibilities

the process is trying to access a page that **does not exist** at all


- ▶ E.g., the process has pages for 1 to 10, but CPU wants to access page 12
- ▶ Solution: Process **aborts**

The page is **not in memory** but is **somewhere on the hard disk**

- ▶ E.g., the process has pages from 1 to 10, CPU wants to access page 2, but this page has not been loaded into RAM
- ▶ Solution: Load the page from hard disk to RAM



Steps in Handling a Page Fault



1. CPU makes a reference to a page which is marked *i* in the page table
2. MMU traps to kernel
 - a) if the reference page does not exist at all, stop the process;
 - b) otherwise, go to 3
3. Find the page in the hard backing store
4. Find a free frame in the memory and load the page into that frame
5. Reset the page table such that the valid bit is *v*
6. Restart the process by re-executing the instruction which caused the page fault.



Aspects of Demand Paging

Pure demand paging


Start process with **no pages** in memory (extreme case)

Immediate page fault for the first instruction in the process

A given instruction could access multiple pages -> multiple page faults

E.g., suppose the following code

```
int d; //global variable
int main(){
    int i;    //local variable
    int *p;   // pointer
    p = (int*) malloc(sizeof(int))
    ... ..
    //Thinking: how many pages are usually
    //accessed for the following instruction
    *p= i + d;
}
```





Performance of Demand Paging (Cont.)

When there is a page fault, the **time** to load a page from hard disk is enormous compared with memory access

Page Fault Rate $0 \leq p \leq 1$

if $p = 0$ no page faults

if $p = 1$, every memory reference is a page fault

Effective Access Time (EAT)

$$\text{EAT} = (1 - p) * \text{memory_access_time} + p * \text{page_fault_time}$$

► Page_fault_time

- time for page fault overhead in MMU and kernel , plus
- time to swap a page out, plus
- time to swap a page in



An Example: EAT

Suppose

Memory access time = 200 nanoseconds

Average page-fault service time = 8 milliseconds

Then

$$\begin{aligned} \text{EAT} &= (1 - p) \times 200 + p \times 8 \text{ milliseconds} = (1 - p) \times 200 + p \times 8,000,000 \\ &= 200 + p \times 7,999,800 \end{aligned}$$

$$\begin{aligned} 1\text{s} &= 10^3 \text{ milliseconds} \\ &= 10^6 \text{ microseconds} \\ &= 10^9 \text{ nanoseconds} \end{aligned}$$

If one access out of 1,000 causes a page fault, i.e., $p = 0.001$,

$$\text{EAT} \approx 200 + 0.001 \times 7,999,800 \approx 8.2 \text{ microseconds}$$

A slowdown by a factor of 40 compared with one memory access

To make the degradation < 10%, i.e., $\text{EAT} < 200 + 200 \times 10\%$

$$200 + 7,999,800 \times p < 220$$

- ▶ $p < 0.0000025$, i.e., one page fault in every 400,000 memory accesses at most

To achieve small p

locality of reference




Locality Reference

Locality of reference

The tendency of processes to reference memory in patterns rather than randomly

Results in reasonable performance from demand paging



Locality in a memory-reference pattern

- It will fault for the pages in its locality until all these pages are in memory
- It will not fault again until it changes localities
- Process migrates from one locality to another over time
- Localities may overlap (over time)



Locality Reference

Working set


The set of pages in the most recent page references

Locality model

A model for page replacement based on the working-set strategy
At any point in time, a process usually uses only a small part of all its pages

page reference table

... 2 6 1 5 7 7 7 7 5 1 6 2 3 4 1 2 3 4 4 4 3 4 3 4 4 4 1 3 2 3 4 4 4 3 4 4 4 ...



Working-set model

Working-set window (Δ): a fixed number of page references



Keeping Track of the Working Set

Approximate the working-set model

Use a fixed-interval timer interrupt + a reference bit

Example: $\Delta = 10,000$ references

- ▶ Timer interrupts after every 5000 references
- ▶ Keep in memory **2 bits** for each page, plus **1** reference bit
- ▶ Whenever a timer interrupts
 - shifting the first bit to become the second bit
 - copy reference bit into the first of the 2 bits
 - reset the value of all reference bits to 0
- ▶ If one of the bits in memory is 1 \Rightarrow page in working set

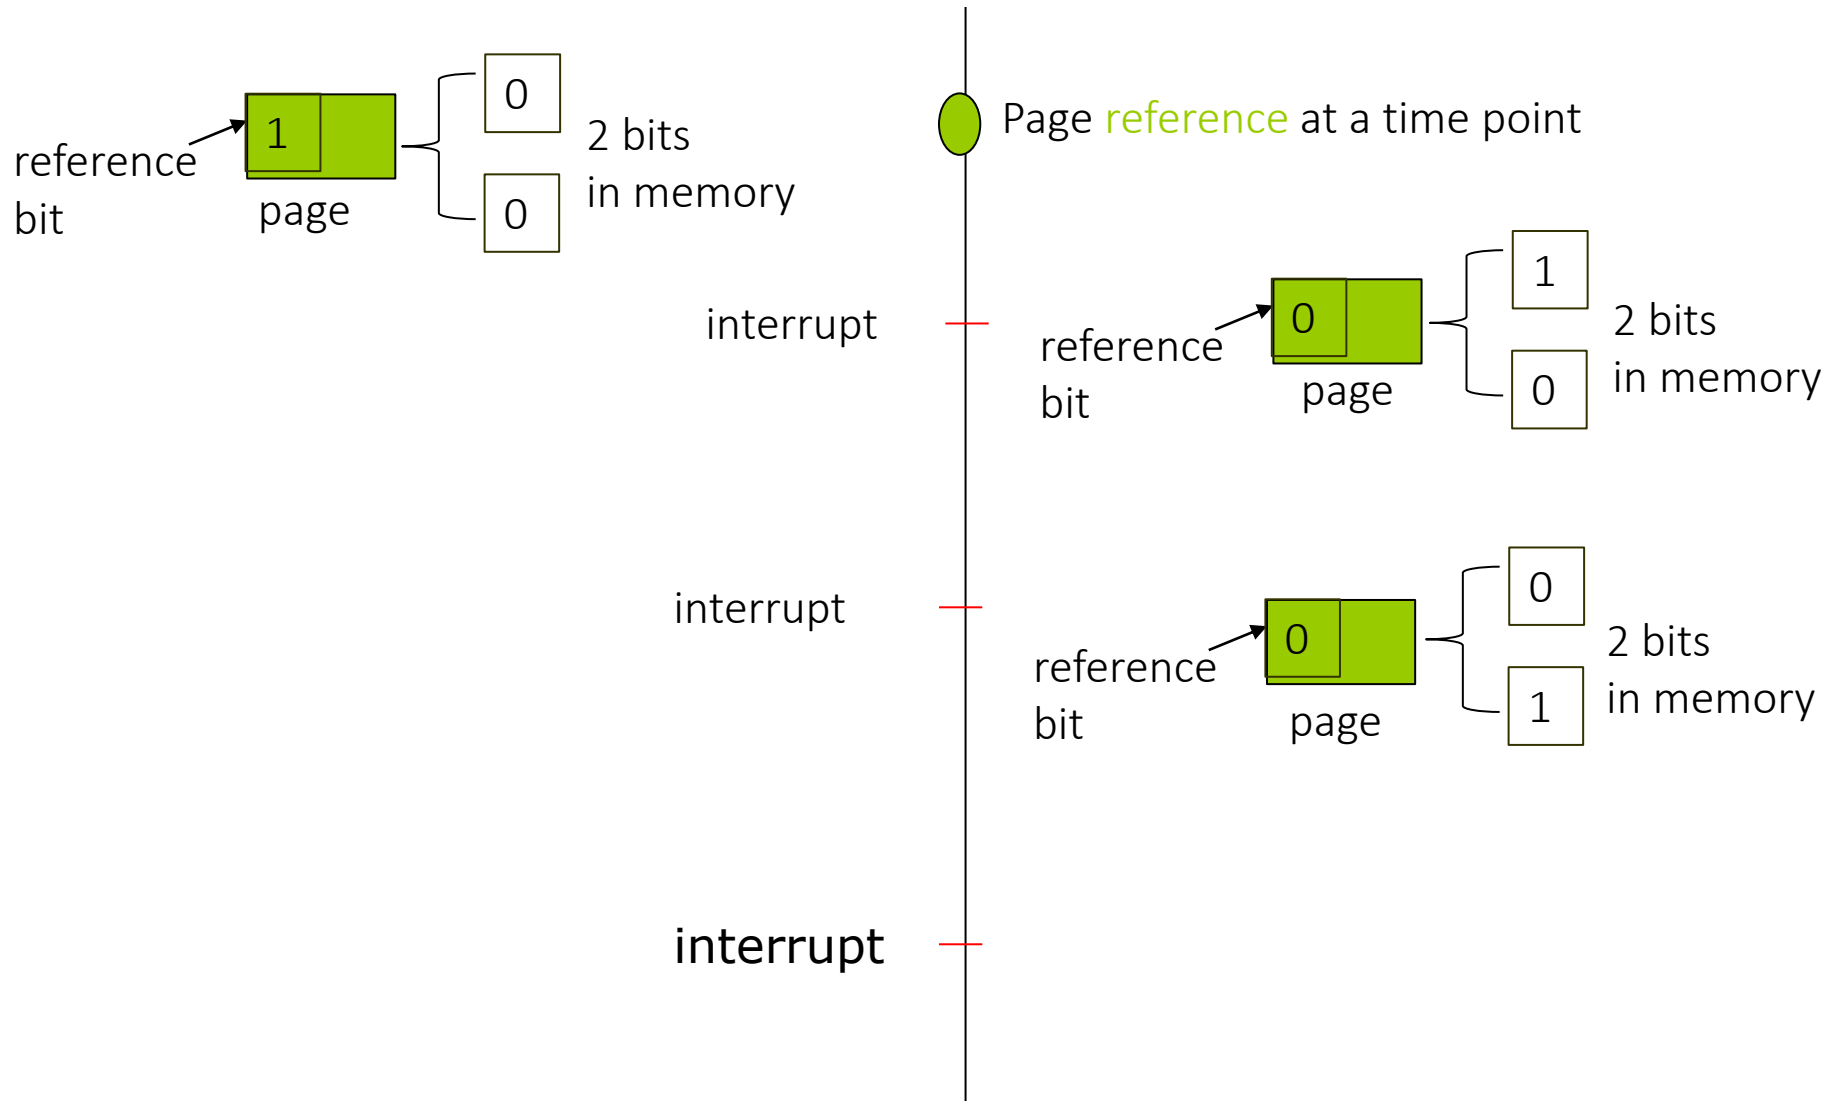
This method is not completely accurate

Unable to tell where, within an interval of 5,000, a reference occurred.



An Example

$\Delta = 10,000$ references, interrupts after every 5000 references





Demand Paging Optimizations

Optimizations

Set a swap space on hard disk (not supported in mobile systems)

- ▶ I/O faster than file system I/O
- ▶ Less management is needed than file system
- ▶ Copy entire process image to swap space at process load time, then page in and out of swap space (e.g., in older BSD Unix)

Discard the pages if they are not modified in the memory

- ▶ When a page is needed to page out to free a frame, if its content is not changed, no need to transfer out

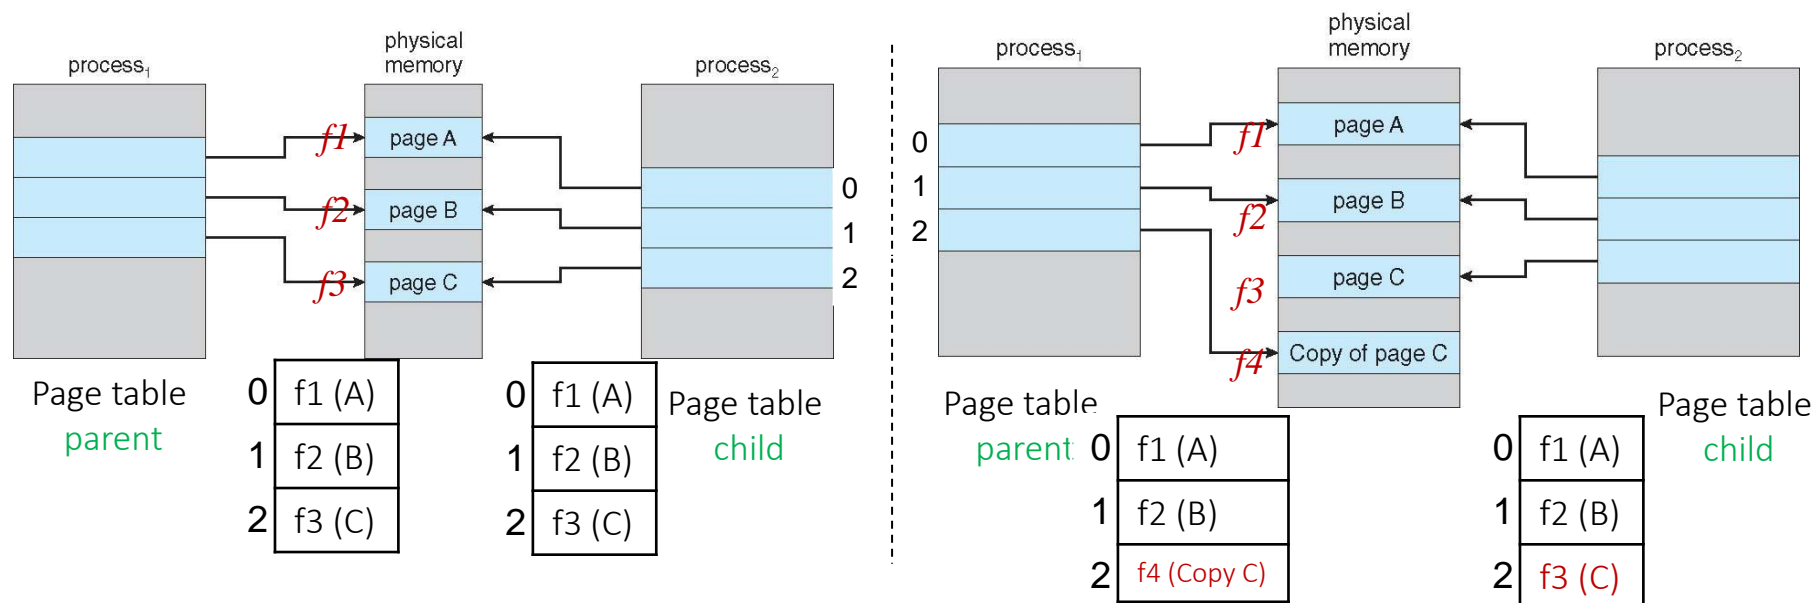


Copy-on-Write

Copy-on-Write (COW) allows both parent and child processes to initially *share the same pages* in memory (*vfork* instead of *fork*)

When a child process is initially created,

- ▶ only copy the page table of its parent
- ▶ share parents' pages
- ▶ the shared pages are write-protected (using protection bits of page tables).
- ▶ If either process modifies a shared page, then copy that page





What Happens if There is no Free Frame?

When there is a page fault, a new page should be loaded in

Find a free frame for this new page to load in

Question: Memory has used up, i.e., no frame is available for a new page

Solution: swap pages out from memory

Page replacement

Find some page in memory, but not really in use, page it out

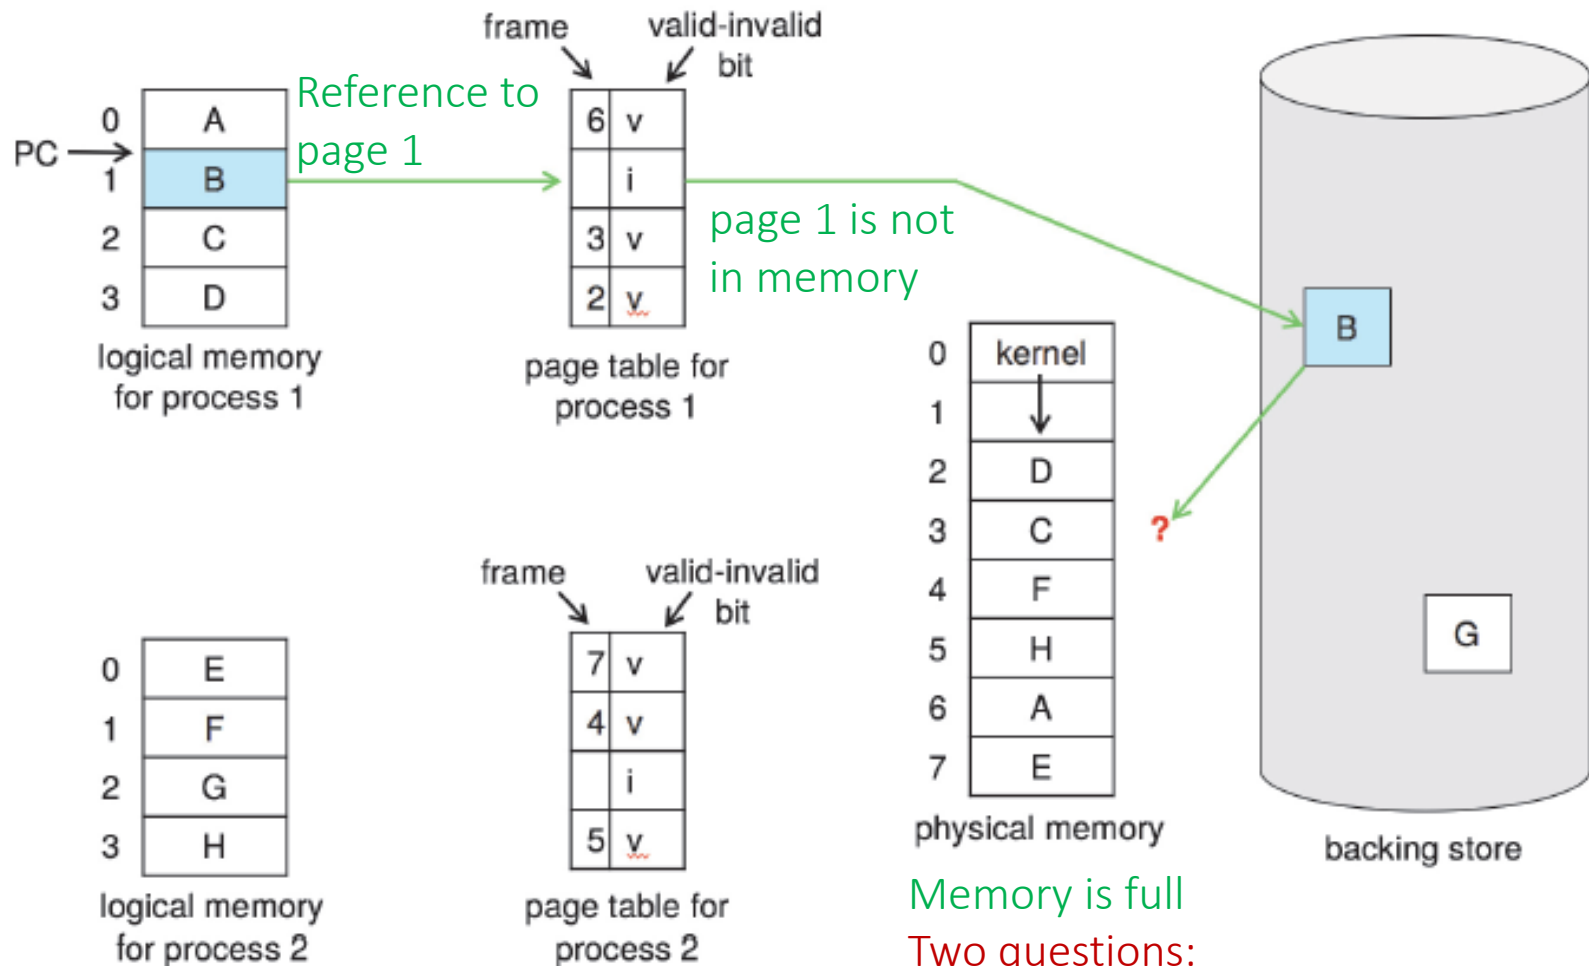
Algorithm consideration

- ▶ How to know which frames are free?
- ▶ If no frame is free, which frames should be freed (i.e., which pages should be replaced, a **victim**)
 - expectation: minimum number of page faults

Same pages may be brought into memory several times



Need For Page Replacement



Two questions:

1. Replace page from p1 or p2?
2. Which page to replace?



Global vs. Local Allocation

Question 1: from which process a page is replaced?

Global allocation

- ▶ Process selects a replacement frame from the set of **all** frames
- ▶ One process can **take a frame from another process**.
- ▶ The # of frames allocated to each process will **change**.

Local allocation

- ▶ Each process selects from only **its own** set of allocated frames.
- ▶ The # of frames allocated to each process **does not change**.



Global vs. Local Allocation

Which is better: Global or Local

Global replacement

- ▶ Process execution time can vary greatly
- ▶ But greater throughput, so more commonly used

Local replacement

- ▶ More consistent per-process performance
- ▶ But possibly under-utilized memory



Page Replacement

Question 2: Which pages to replace (be removed from RAM)?

Two kinds of pages are good candidates

- ▶ pages which are not used by a process right now
 - If a page has not been used for a while, maybe it will not be used in near future
- ▶ Pages that are not modified (called **dirty**) recently
 - These kinds of pages do not need to be paged out

An **algorithm** is needed to find the right frame to free


- ▶ which pages have not been used recently?
- ▶ which pages have not been modified recently?
- ▶ “recent”: how to decide the period for “recent”?

Victim page

- ▶ The page which will be replaced



Page Replacement





Page and Frame Replacement Algorithms

Frame-allocation algorithm determines

How many frames to give each process

Page-replacement algorithm

Which frames to replace

Want lowest page-fault rate on both first access and re-access

Algorithm evaluation

Run it on a particular string of memory references (reference string) and compute the number of page faults on that string

String is just page numbers, not full addresses

Results depend on number of frames available



Page Replacement Algorithms

Algorithms

FIFO

Optimal

Least Recently Used (LRU)

Second chance

Enhanced second chance

Reference string used in the examples

7,0,1,2,0,3,0,4,2,3,0,3,2,1,2,0,1,7,0,1

Each number is the page number to reference



FIFO

Rule

Replace the oldest one

Implementation

OS maintains a circular queue of all pages

- ▶ Page at head of the list: Oldest one
- ▶ Page at the tail: Recent arrival

When there is a page fault

- ▶ Page at the head is removed
- ▶ New page added to the tail




First-In-First-Out (FIFO) Algorithm

Reference string: 7,0,1,2,0,3,0,4,2,3,0,3,2,1,2,0,1,7,0,1

3 frames (i.e., 3 pages can be in memory at a time)

reference string



page frames


Total page faults: 15

Thinking:

Can adding more frames reduce the number of page faults?



Graph of Page Faults Versus The Number of Frames



Intuition but not Reality



FIFO Illustrating Belady's Anomaly


Consider 1,2,3,4,1,2,5,1,2,3,4,5

Adding **more frames** might cause **more page faults**!

1	2	3	4	1	2	5	1	2	3	4	5	page fault
1	1	1	4	4	4	5	5	5	5	5	5	
	2	2	2	1	1	1	1	1	3	3	3	9
		3	3	3	2	2	2	2	2	4	4	
1	1	1	1	1	1	5	5	5	5	4	4	
	2	2	2	2	2	2	1	1	1	1	5	10
		3	3	3	3	3	3	2	2	2	2	
			4	4	4	4	4	4	3	3	3	

3 frames

4 frames



Belady's Anomaly

For some page-replacement algorithms, the **page-fault rate** may **increase** as the number of allocated **frames increases**.




Optimal Algorithm

Replace the page that **will not be used** for longest period of time

reference string

7 0 1 2 0 ③ 0 ④ 2 3 ① 3 2 ① 2 0 1 ⑦ 0 1



page frames

Total page faults: 9

Problem

Can't read the future

This method can only be used to measure how other algorithms are close to the optimal



reference string



LRU is a good algorithm and frequently used

LRU and OPT don't have Belady's Anomaly



LRU Algorithm Implementation 1: Counter

Counter implementation

There is a global counter that increases by 1 whenever a page is referenced. Every page in the memory has its own counter

When a page is referenced, its counter is synchronized with the global counter

When a page needs to be replaced, find the page in the memory with the **smallest** value

reference string


7 0 1 ② 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	1	7	1	7	1	2	4	2	4	2	4	2	4	2	4	4	8	4	8	4	8	0	11	0	11	0	11	1	14	1	14	1	14	1	17	1	17	1	17	1	20
				0	2	0	2	0	2	0	2	0	5	0	7	0	7	0	7	3	10	3	10	3	12	3	12	3	12	3	12	0	16	0	16	0	16	0	19	0	19
				1	3	1	3	1	3	1	3	3	6	3	6	3	6	2	9	2	9	2	9	2	9	2	13	2	13	2	15	2	15	2	15	7	18	7	18	7	18

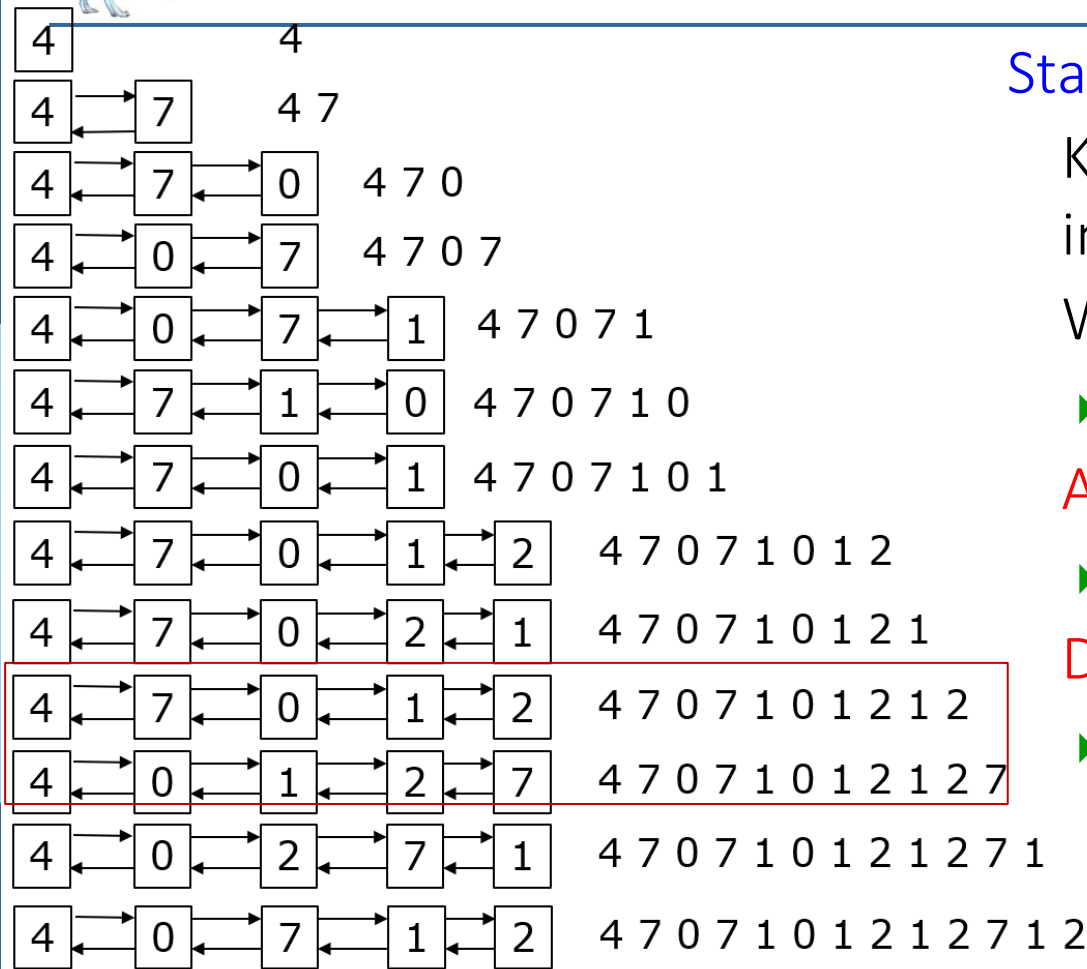
page frames

Thinking:

What is the disadvantage of this algorithm (time issue)?



LRU Algorithm Implementation 2: Stack



Stack implementation

Keep a stack of page numbers in a double link form:

When a page referenced:

- move it to the top

Advantage

- No search for replacement

Disadvantage

- Each update of the stack might requires 6 pointers to be changed

Reference string:

4 7 0 7 1 0 1 2 1 2 7 1 2

Thinking:

1. In the stack, which one is least recently used?
which one is the most recently used
2. Why needs 6 pointer changes for each update?



LRU Approximation Algorithms

LRU **approximation** algorithm 1

Usage of a reference bit



- ▶ Associate each page with a reference bit, initial value 0
- ▶ When a page is referenced, set the bit 1
- ▶ When a page is needed to be replaced, find the ones with reference bit 0 if exist one

Thinking:

What is the problem with this algorithm?



LRU Approximation Algorithms



LRU **approximation** algorithm 2

Second-chance algorithm

- ▶ Use a circular **FIFO** and **reference bit** for **each page** in memory
- ▶ If page to be replaced has
 - reference bit = 0
 - » replace it
 - reference bit = 1
 - » set reference bit to 0
 - » search for the next page



LRU Approximation Algorithms

Enhanced second-chance algorithm

Improve algorithm by using **reference bit** and **modify bit** for each page in memory, i.e., an ordered pair (reference, modify)

All pages in memory fall into **four classes**

1. (0, 0) neither recently used nor modified – **best page to replace**
2. (0, 1) not recently used but modified – not quite as good, must write out before replacement
3. (1, 0) recently used but clean – probably will be used again soon
4. (1, 1) recently used and modified – probably will be used again soon and need to write out before replacement – **worst page to replace**

When page replacement called for, replace page in **lowest non-empty class** in clock scheme

Disadvantage

- ▶ Might need to search circular queue several times



Other Algorithms

Counting algorithms

Keep a counter of the number of references that have been made to each page

- ▶ **Least Frequently Used (LFU) Algorithm:**

- replaces page with **smallest count**

- ▶ **Most Frequently Used (MFU) Algorithm:**

- based on the argument that the page with the smallest count was probably just brought in and has yet to be used

None of these counting algorithms are **common**



Page-Buffering

More strategies (algorithms) in improving the efficiency of page replacement

Page-buffering consideration

Keep a pool of free frames always

- ▶ Whenever it is possible, select a victim to evict (逐出) and add it to free pool
- ▶ When convenient, evict the victim
- ▶ When frame needed, read page into free frame
- ▶ Advantage: Reduce time to find a free frame at a page fault

Expansion

- ▶ Keep a list of modified pages
 - Whenever the paging device is idle, write the modified page to the disk. Its modify bit is then reset to 0
- ▶ Keep free frame contents intact (完整的, 内容未被清除的)
 - Reduce penalty if wrong victim frame was selected
 - If the page is referenced again, no need to load from the disk



Allocation of Frames

For performance reason, each process needs **minimum number** of frames

This number is defined by the computer architecture (**CPU**)

- ▶ Example: IBM 370 – 6 pages to handle SS MOVE instruction
 - instruction might span 2 pages
 - 2 pages to handle *from*
 - 2 pages to handle *to*

The **maximum number of frames** is defined by the amount of **total frames** in the system (i.e., physical memory, RAM)



Allocation of Frames

Two major allocation schemes

1. Fixed allocation

▶ Equal allocation

- Each process is allocated **same number** of frames
- Disadvantage
 - » Space waste for small process

▶ Propositional allocation

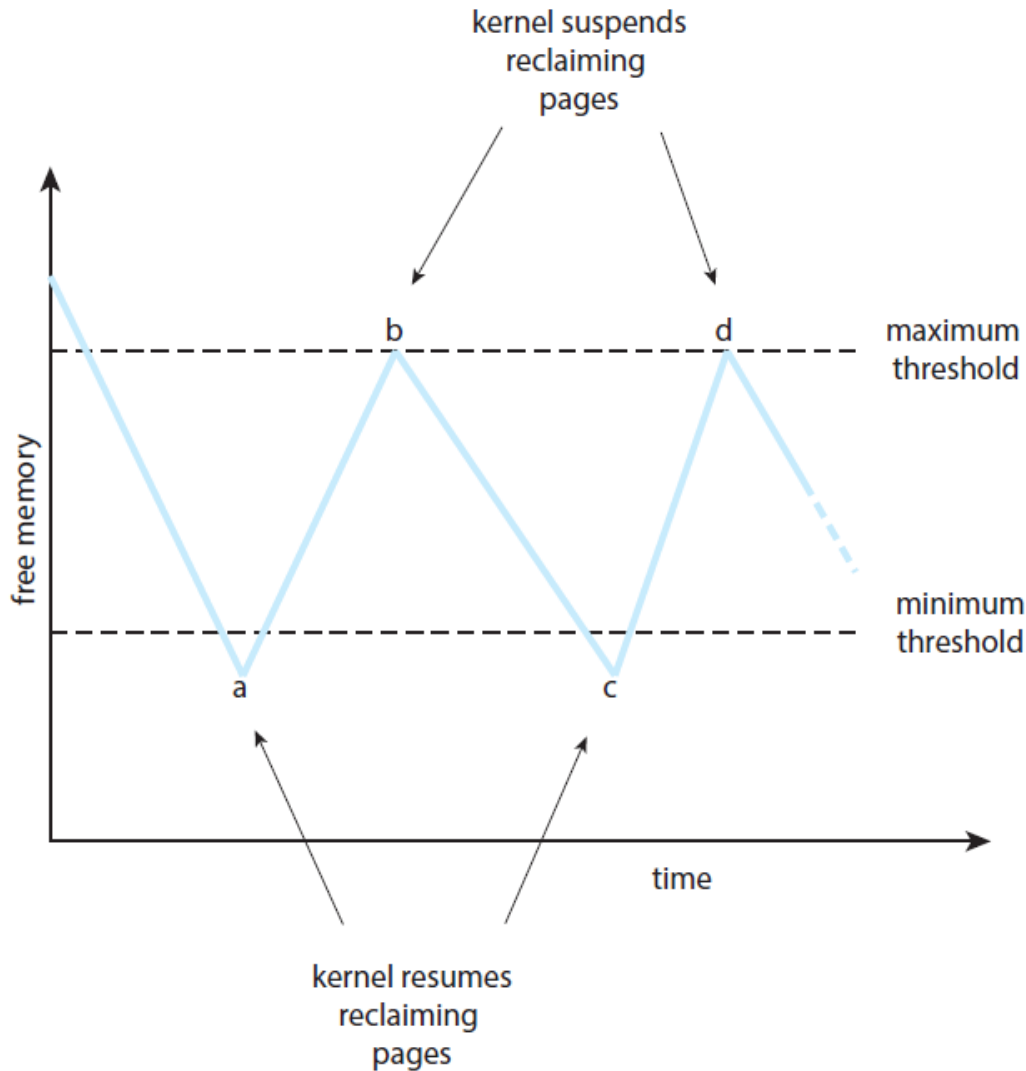
- Allocate frames according to **the size of a process**
- Disadvantage
 - » Process size might be changed during the execution

2. Priority allocation

- ▶ the ratio of frames depends on the **combination of size and priority of a process**
- ▶ Replace the pages of process with lower priority



Allocation of Frames



Page replacement happens before the number of frames falls below a certain threshold



Thrashing(系统颠簸)

If a process does not have “enough” frames in memory, the page-fault rate is very high

Replace page frequently

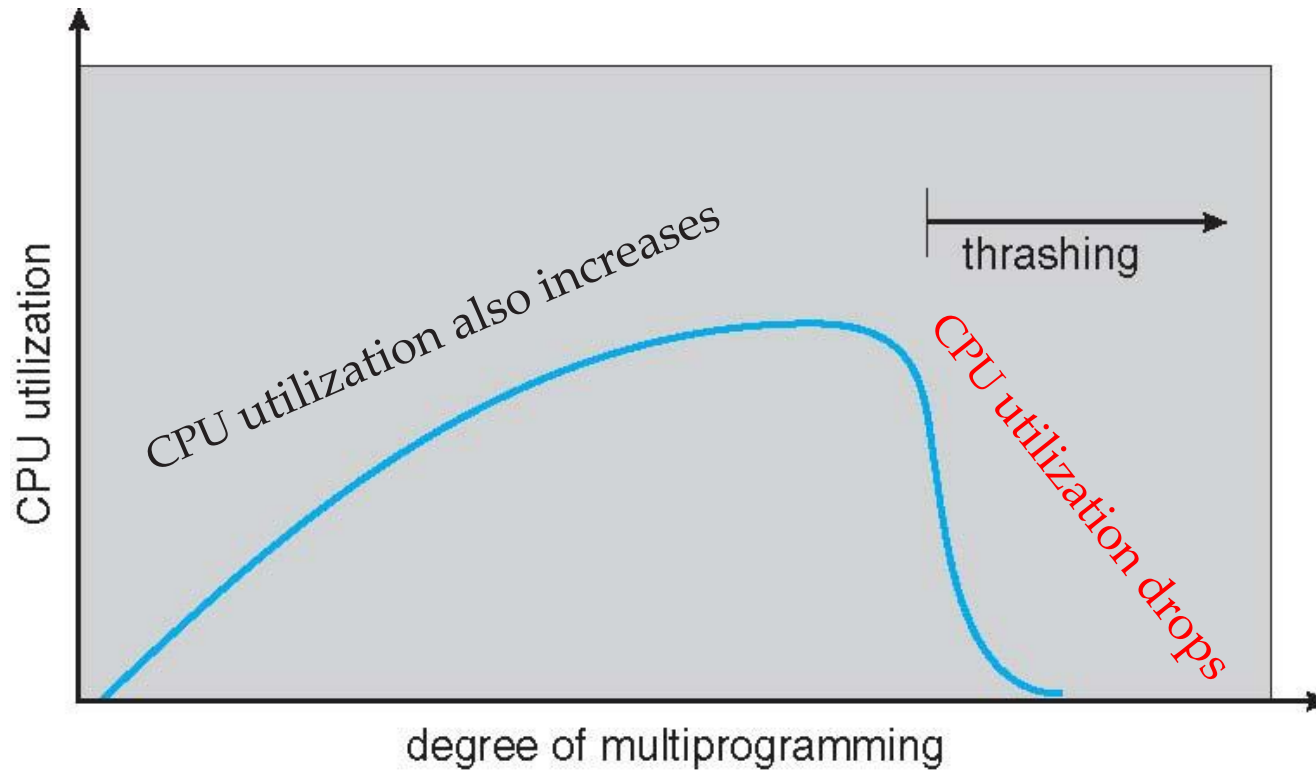
The replaced page might be used again

This leads to:

- ▶ Low CPU utilization
- ▶ Operating system keeps adding new process, trying to increase CPU utilization
- ▶ Things get worse -> entering a bad cycle
- ▶ Thrashing
 - A process is busy swapping pages in and out, instead of doing useful work.



Thrashing (Cont.)



Solutions to thrashing


1. Decrease the degree of multiprogramming
2. Establish “acceptable” [page-fault frequency \(PFF\) rate](#) and use local replacement policy
3. Install enough physical memory (hardware)
4. Install faster hard disk



Page-Fault Frequency

If actual rate too low, process loses frame (we think too many frames assigned, remove some frames)

If actual rate too high, process gains frame





Allocating Kernel Memory

Allocating memory for **Kernel** is **different** from the allocation of memory for **user applications**

Special features of kernel

Kernel requests memory for structures of varying sizes,

- ▶ E.g.,
 - structure for PCB
 - structure for file descriptor

There are multiples instances for each structure,

- ▶ e.g., there is a PCB for each process

Memory needs to be contiguous

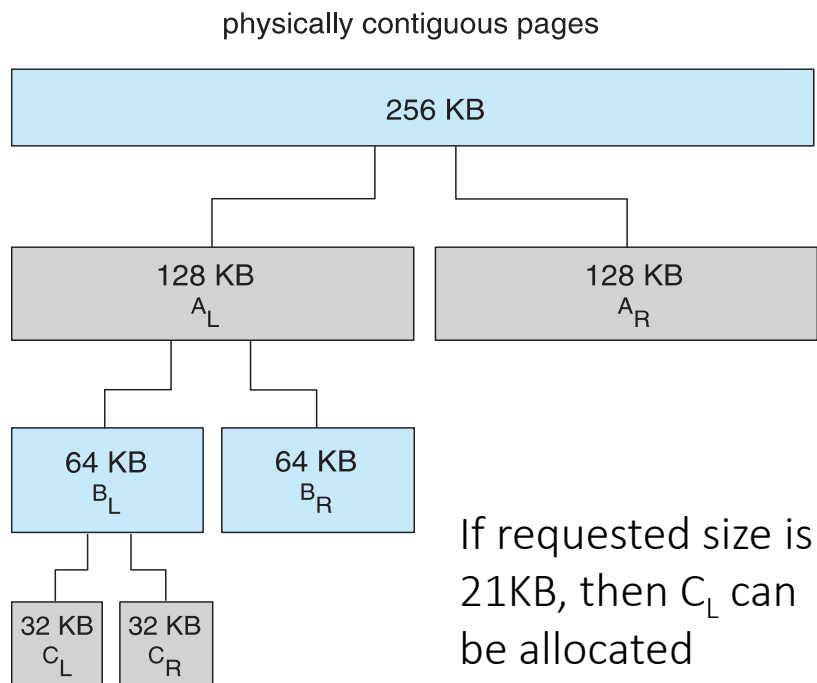


Allocating Kernel Memory: Two Methods

Two allocation methods

1. Buddy system

- ▶ Allocates memory from fixed-size segment consisting of physically-contiguous pages
- ▶ **Power-of-2 allocator**: memory allocated in units is sized as power of 2, one smallest but bigger than requested size



If requested size is 21KB, then C_L can be allocated

Advantage:

Unused chunks can be merged to a bigger one

Thinking:

Disadvantage?



Allocating Kernel Memory: Two Methods

Two allocation methods

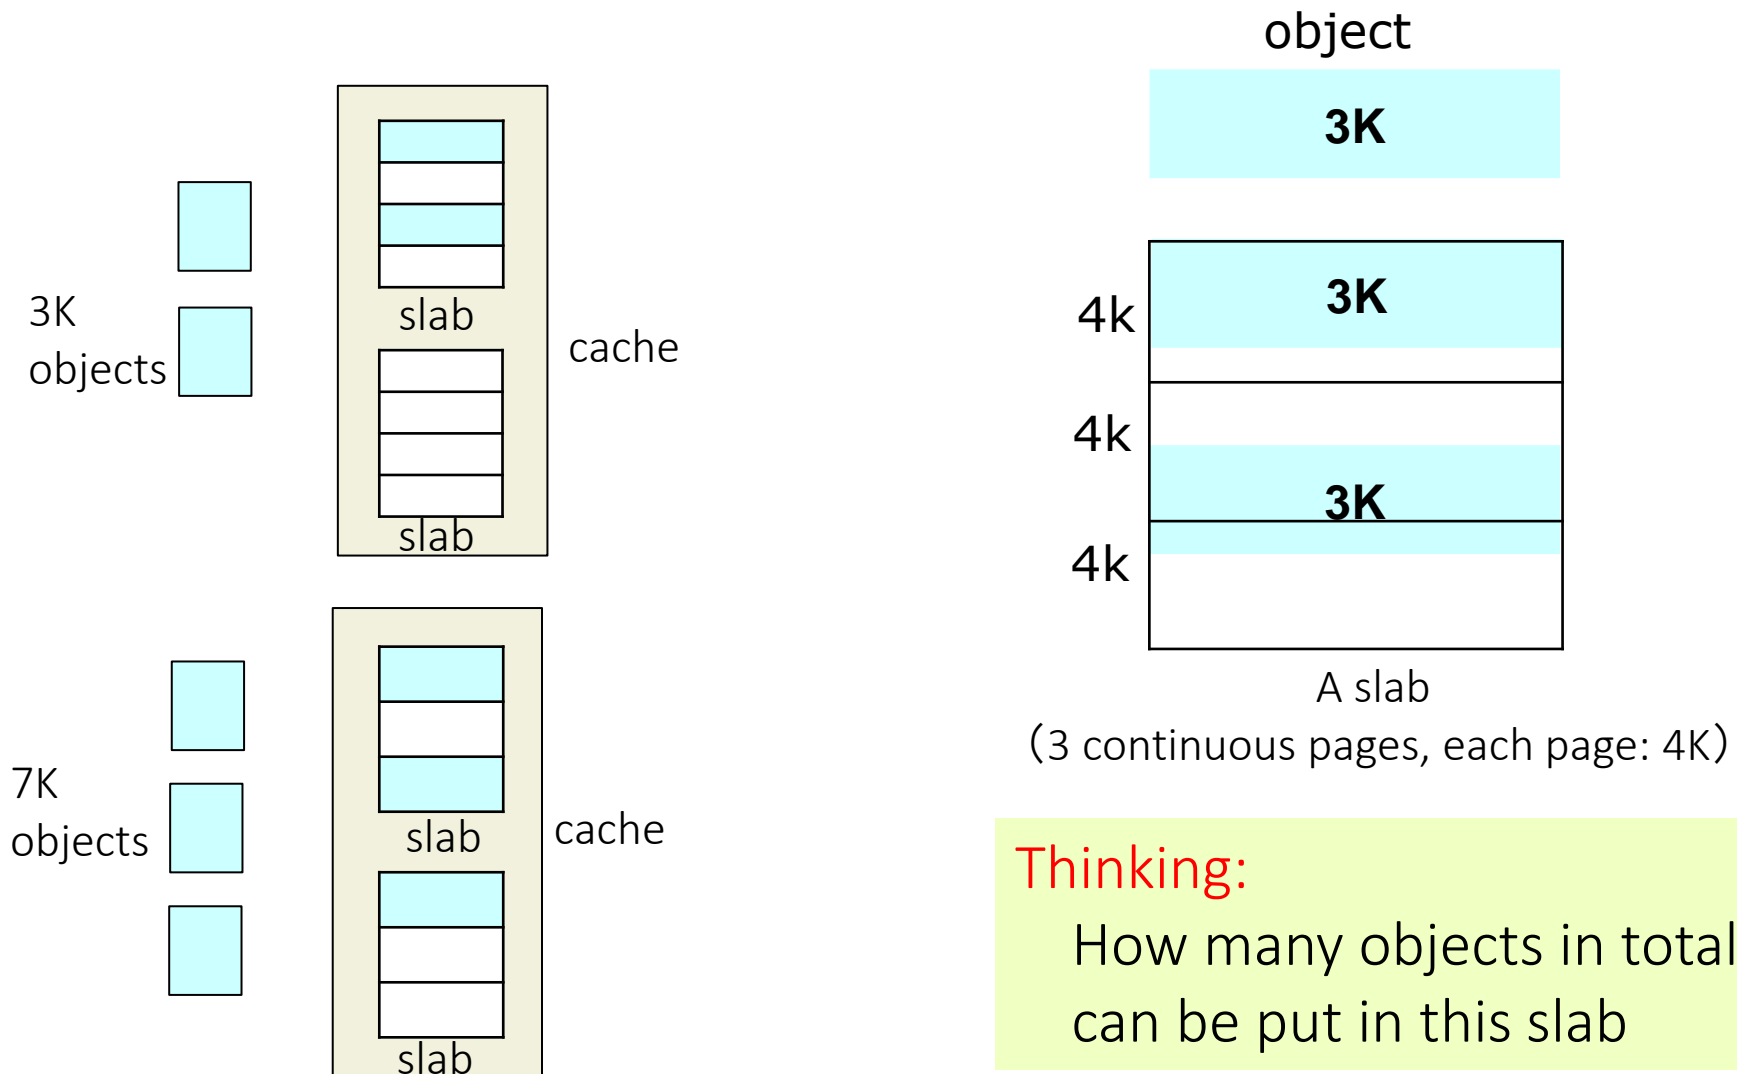
2. Slab allocator

▶ A cache

- is used for unique kernel data structure
- consists of one or more slabs
- A **slab** is one or more physically contiguous pages
 - » filled with same kind of objects – instantiations of the data structure,
 - » Each object has two status: free and used
- If slab is full of used objects, next object is allocated from empty slab
 - » If no empty slabs, new slab allocated



Slab Allocation: An Example





Other Considerations

To reduce page faults

Pre-paging

- ▶ Pre-page all or some of the pages a process will need, before they are referenced
- ▶ If pre-paged pages are unused, I/O and memory was wasted
- ▶ Need to balance

Page size selection

- ▶ Fragmentation
- ▶ Page table size
- ▶ Need to balance

TLB reach

- ▶ Depend on page size and TLB size
- ▶ Need to balance

Program structure

- ▶ What can programmers do?



Other Issues – Program Structure

Two programs (assume data are stored row by row)

Program 1: accessing data column by column

```
for (j = 0; j < 128; j++)    128 x 128 = 16,384 page faults!  
    for (i = 0; i < 128; i++)  
        data[i][j] = 0; //each i is in different page
```

Program 2: accessing data row by row

```
for (i = 0; i < 128; i++)    128 page faults only!  
    for (j = 0; j < 128; j++)  
        data[i][j] = 0; //each j is in different page
```

data[0][0]
data[0][1]
.....
data[1][0]
data[1][1]
.....
data[2][0]
data[2][1]
.....
data[3][0]
data[2][1]
.....



Example: Microsoft Windows

Uses demand paging with **clustering**.

Clustering brings in pages surrounding the faulting page

Processes are assigned **working set minimum** and **working set maximum**

Working set minimum is the minimum number of pages the process is guaranteed to have in memory

A process may be assigned as many pages up to its working set maximum

When the amount of free memory in the system falls below a threshold, **automatic working set trimming** is performed to restore the amount of free memory

- ▶ Working set trimming removes pages from processes that have pages in excess of their working set minimum

End of Chapter 10

